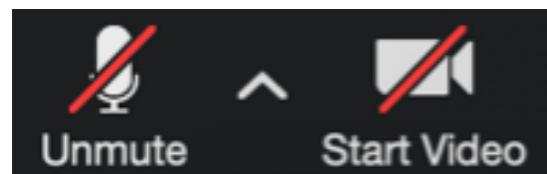




AEROSP 584 - Navigation and Guidance: From Perception to Control



Lectures start at
10:30am EST

Vasileios Tzoumas

Lecture 19



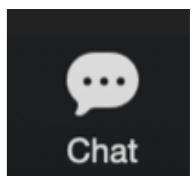
Based on slides made by Luca Carlone @



To ask questions:



or



Raise Hand

Today

- Recap on 2-view
- RANSAC
- 3D-3D correspondences

SRI International

RANDOM SAMPLE CONSENSUS: A PARADIGM FOR MODEL FITTING WITH APPLICATIONS TO IMAGE ANALYSIS AND AUTOMATED CARTOGRAPHY

Technical Note 213

March 1980

By: Martin A. Fischler, Senior Computer Scientist
Robert C. Bolles, Computer Scientist

Artificial Intelligence Center
SRI International
Menlo Park, California 94025

SRI Projects 5300 and 1009

The work reported herein was supported by the Defense Advanced Research Projects Agency under Contract Nos. DAAG29-76-C-0057 and DAAG29-79-C-0588.

333 Ravenswood Avenue • Menlo Park, CA 94025
415-326-62 a460585.pdf 34-486

IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. PAMI-9, NO. 5, SEPTEMBER 1987

698

V. Cappellini and A. G. Constantinides, Eds. Amsterdam, The Netherlands: Elsevier, pp. 730-735.

[3] J. O'Rourke, "Motion detection using Hough techniques," in *Proceedings of Pattern Recognition and Image Processing*, Dallas, TX, 1981, p. 737.

[4] T. M. Sangerberg, L. Davis, and D. Headwood, "An iterative Hough transform for 3-D object recognition," in *Pattern Recognition*, vol. 17, no. 6, pp. 621-629, 1984.

[5] R. C. Bolles, M. A. Fischler, and J. L. M. Lavine, "Fast Hough Transform," in *Proceedings of the 1984 Conference on Pattern Recognition*, Bellair, MI, 1984, pp. 75-83.

[6] R. C. Bolles, "A Hough Transform based on B-splines," in *Proc. Conf. Computer Vision and Pattern Recognition*, Miami Beach, FL, 1986, pp. 640-642.

[7] R. C. Bolles, "Shape matching using new connected components algorithm for virtual memory computers," *Computer Graphics and Image Processing*, vol. 22, pp. 380-390, 1983.

[8] A. Bowyer and J. Woodcock, *A practical guide to geometry*. London: Butterworth, 1983.

[9] R. C. Bolles and T. Tousignant, "On the detection of structures in noisy pictures," *Pattern Recognition*, vol. 9, pp. 95-98, 1977.

[10] M. Van Veen and P. C. A. Groot, "Detection of structures in the presence of noise," *Pattern Recognition*, vol. 14, pp. 141-145, 1981.

[11] R. Lini, "A new three-dimensional connected components algorithm," *Pattern, Vision, Graphics, Image Processing*, vol. 23, pp. 207-217, 1983.

[12] Least-Squares Fitting of Two 3-D Point Sets
K. S. ARUN, T. S. HUANG, AND S. D. BLOSTEIN

Abstract—Two point sets $\{p_i\}$ and $\{p'_i\}$, $i = 1, 2, \dots, N$, are related by $p'_i = Rp_i + T + N_r$, where R is a rotation matrix, T is a translation vector, and N_r is a noise vector. Given $\{p_i\}$ and $\{p'_i\}$, we present an algorithm to find the least-squares solution of R and T , which is based on the singular value decomposition (SVD) of a 3×3 matrix. This new algorithm is compared to two earlier algorithms with respect to computer time requirements.

Index Terms—Computer vision, least-squares, motion estimation, quaternions, singular value decomposition.

I. INTRODUCTION

In many computer vision applications, notably the estimation of motion parameters of a rigid object using 3-D point correspondences [1], the motion estimation is the relative attitude of a rigid object with respect to a reference [2]. We consider the rigid mathematical problem. We are given two 3-D point sets $\{p_i\}$, $i = 1, 2, \dots, N$ (here, p_i and p'_i are considered as 3×1 column matrices)

$$p'_i = Rp_i + T + N_r \quad (1)$$

where R is a 3×3 rotation matrix, T is a translation vector (3×1 column matrix), and N_r is a noise vector. (We assume that the rotation is around an axis passing through the origin). We want to find R and T to minimize

$$\Sigma^2 = \sum_{i=1}^N \|p'_i - (Rp_i + T)\|^2 \quad (2)$$

Manuscript received July 1, 1986; revised April 9, 1987. Recommended for acceptance by S. W. Zucker. This work was supported by the National Science Foundation under Grant RI-8605400.

The authors are with the Coordinated Science Laboratory, University of Illinois, Urbana, IL 61801.

IEEE Log Number 871589.

0162-8828/87/0900-0698\$01.00 © 1987 IEEE

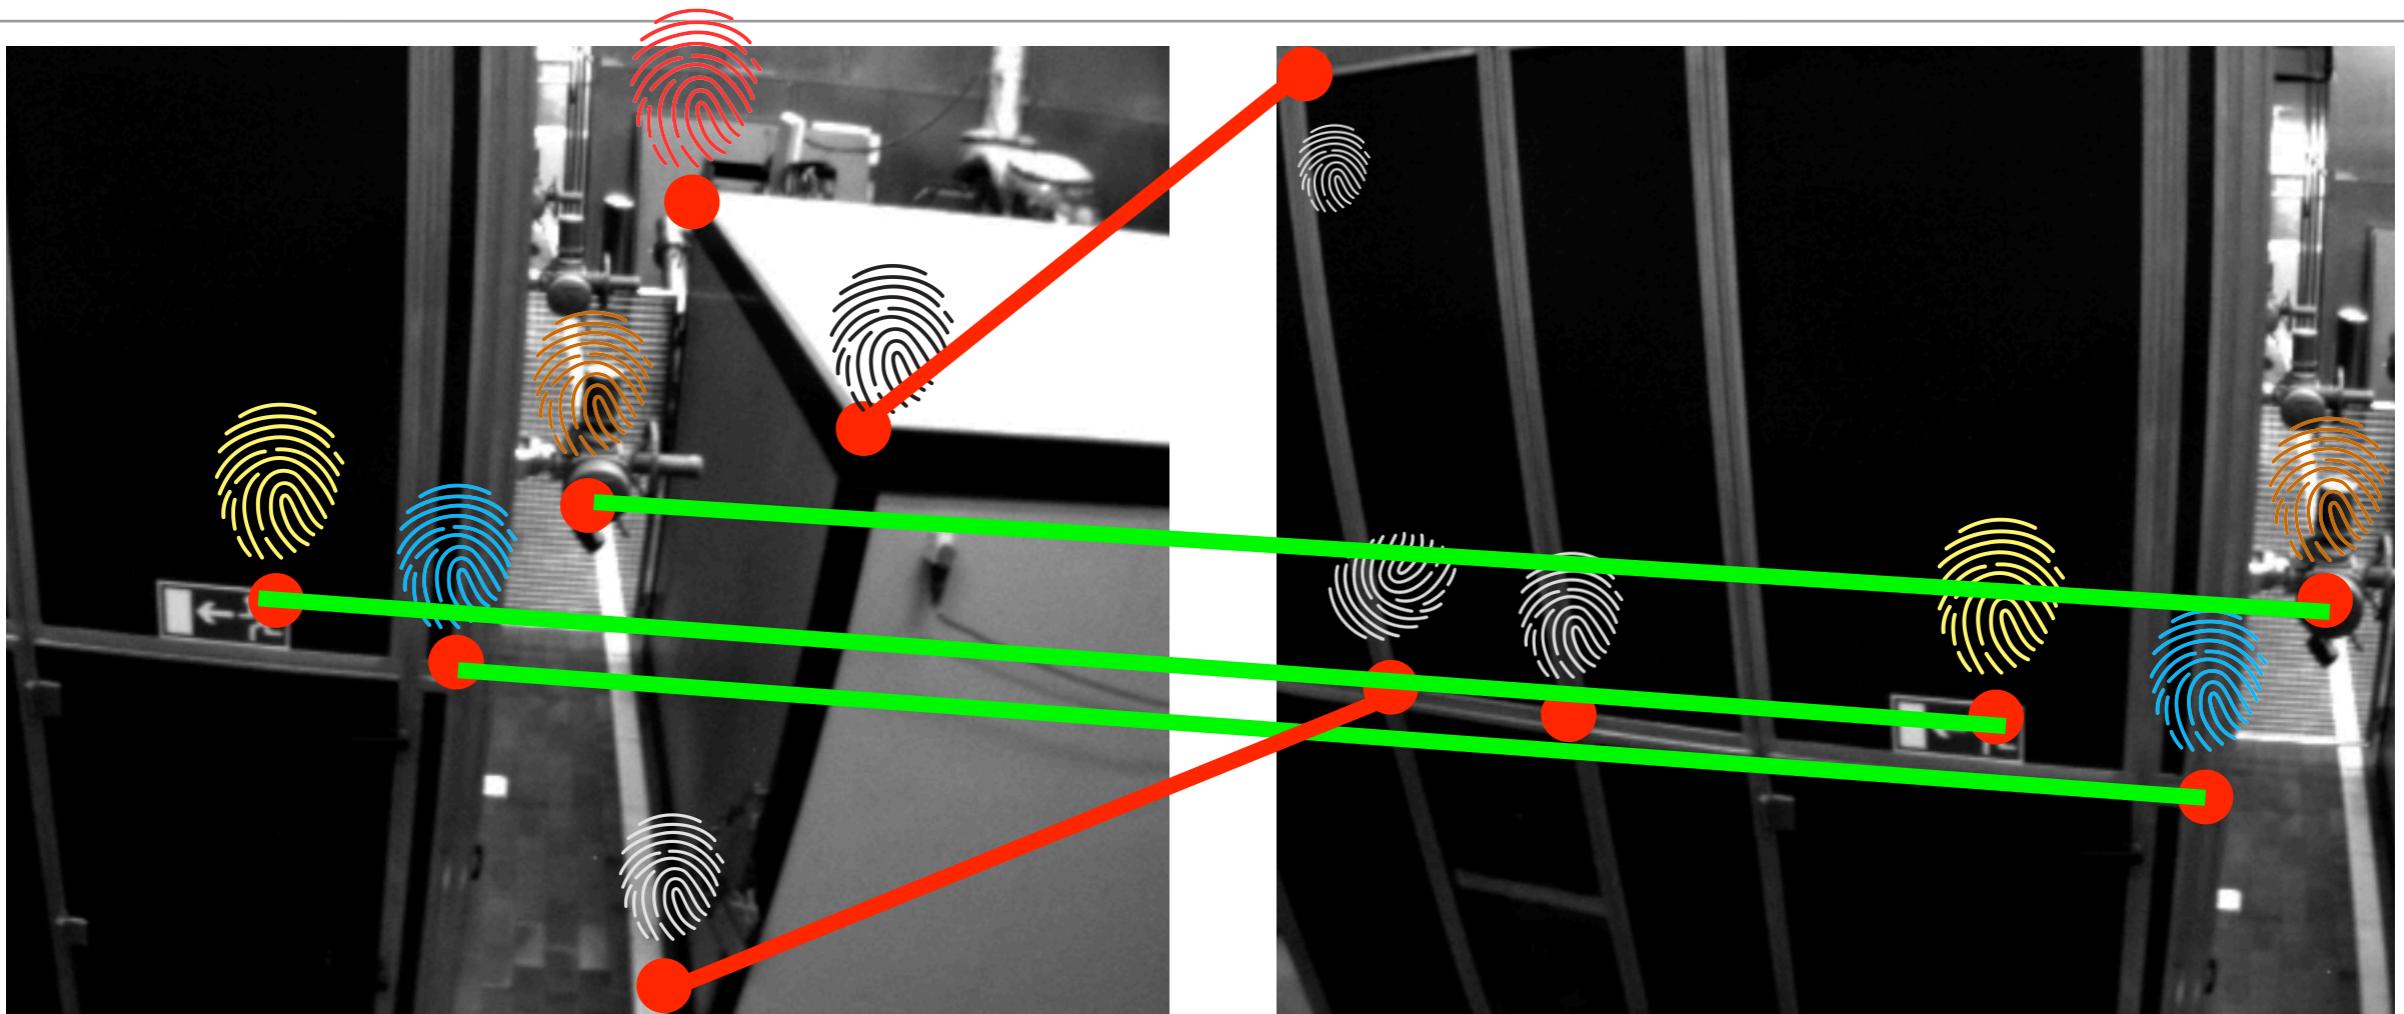
[1] M .Fisher, R. Bollets, "Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography", SRI Technical Note, 1980.

[2] K.S. Arun, T.S. Huang, S.D. Blostein, "Least-Squares Fitting of Two 3-D Point Sets", IEEE Transactions on Pattern Analysis and Machine Intelligence (PAMI), 9(5), 698-700, 1987.

[1]

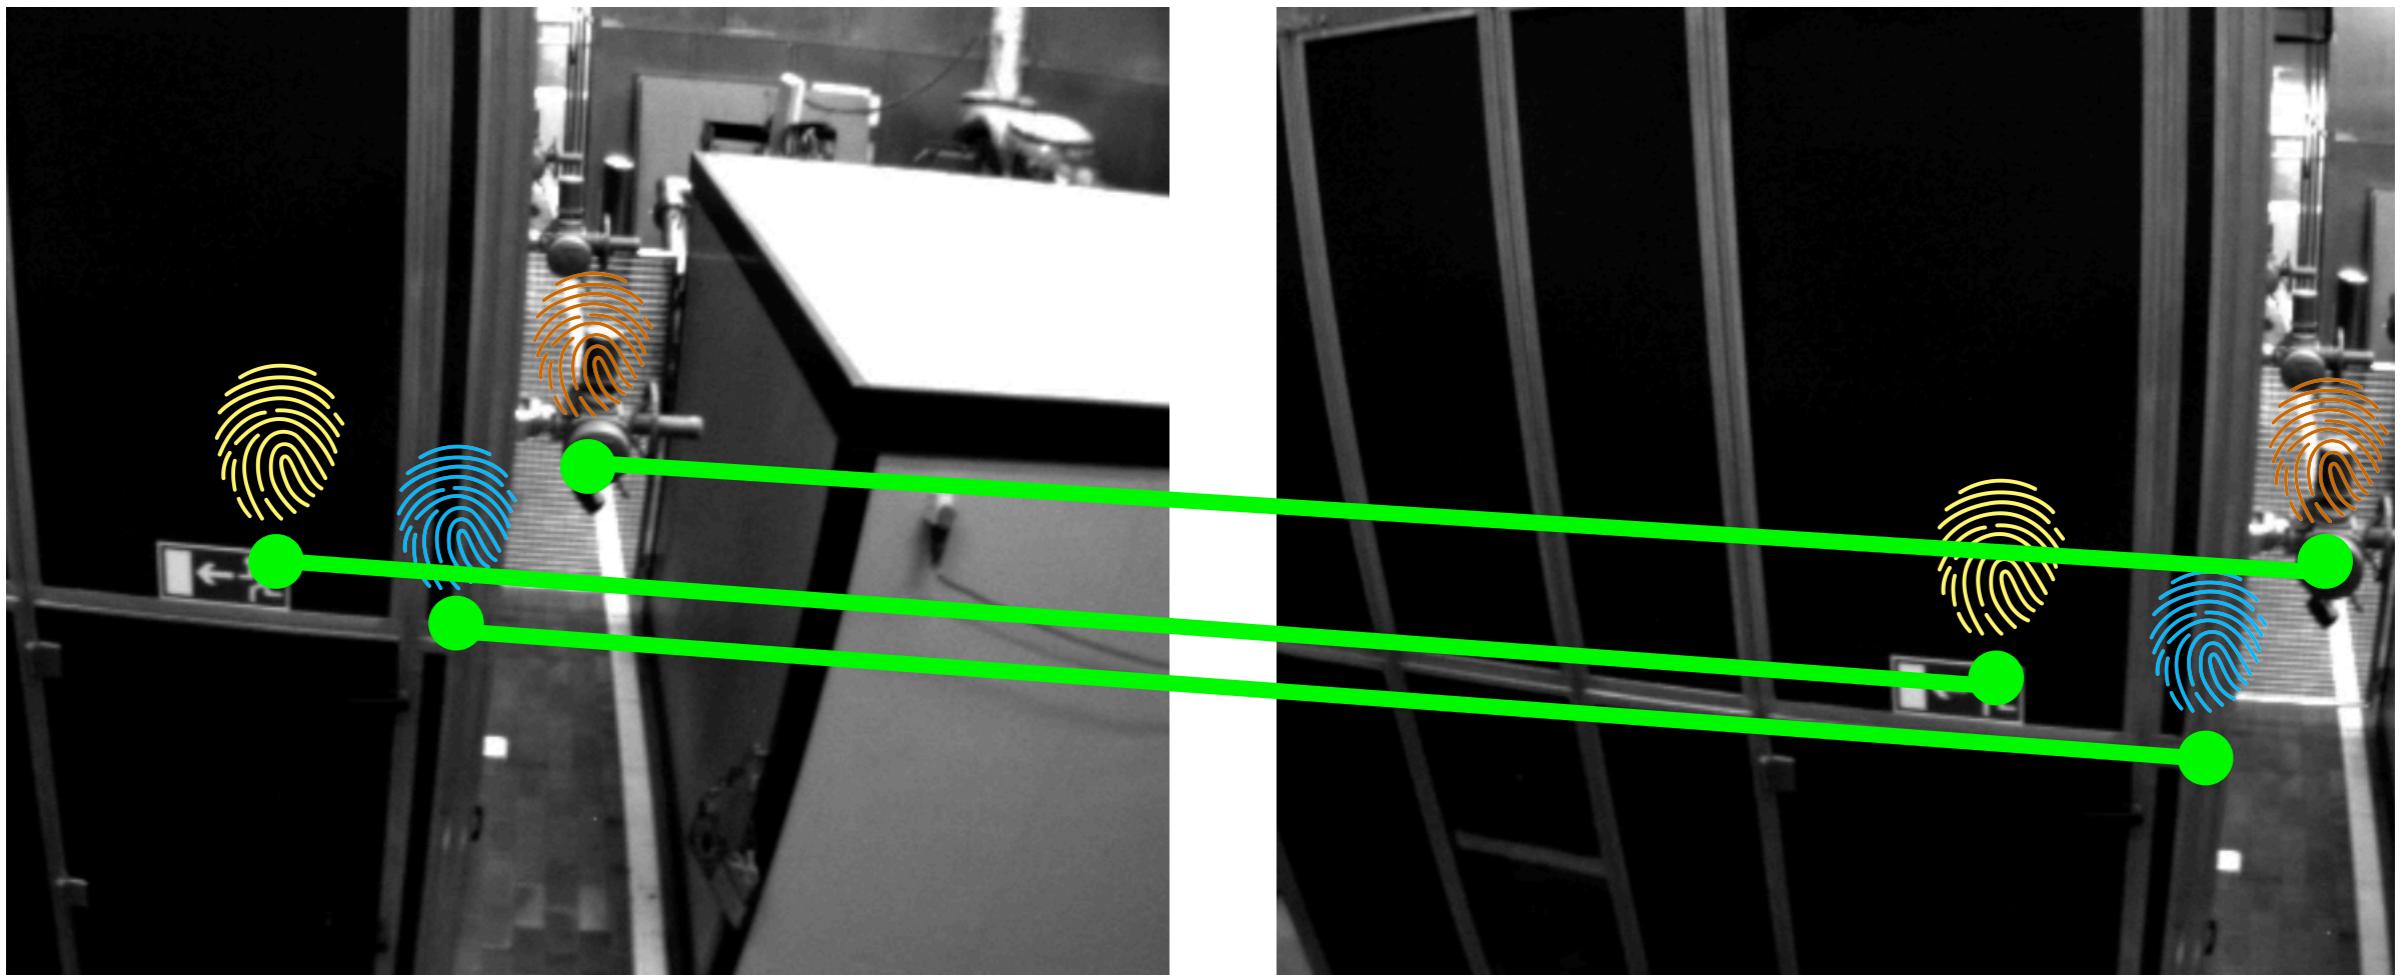
[2]

2-view Geometry



Question: can we estimate the motion of the camera between I_1 and I_2 using pixel correspondences?

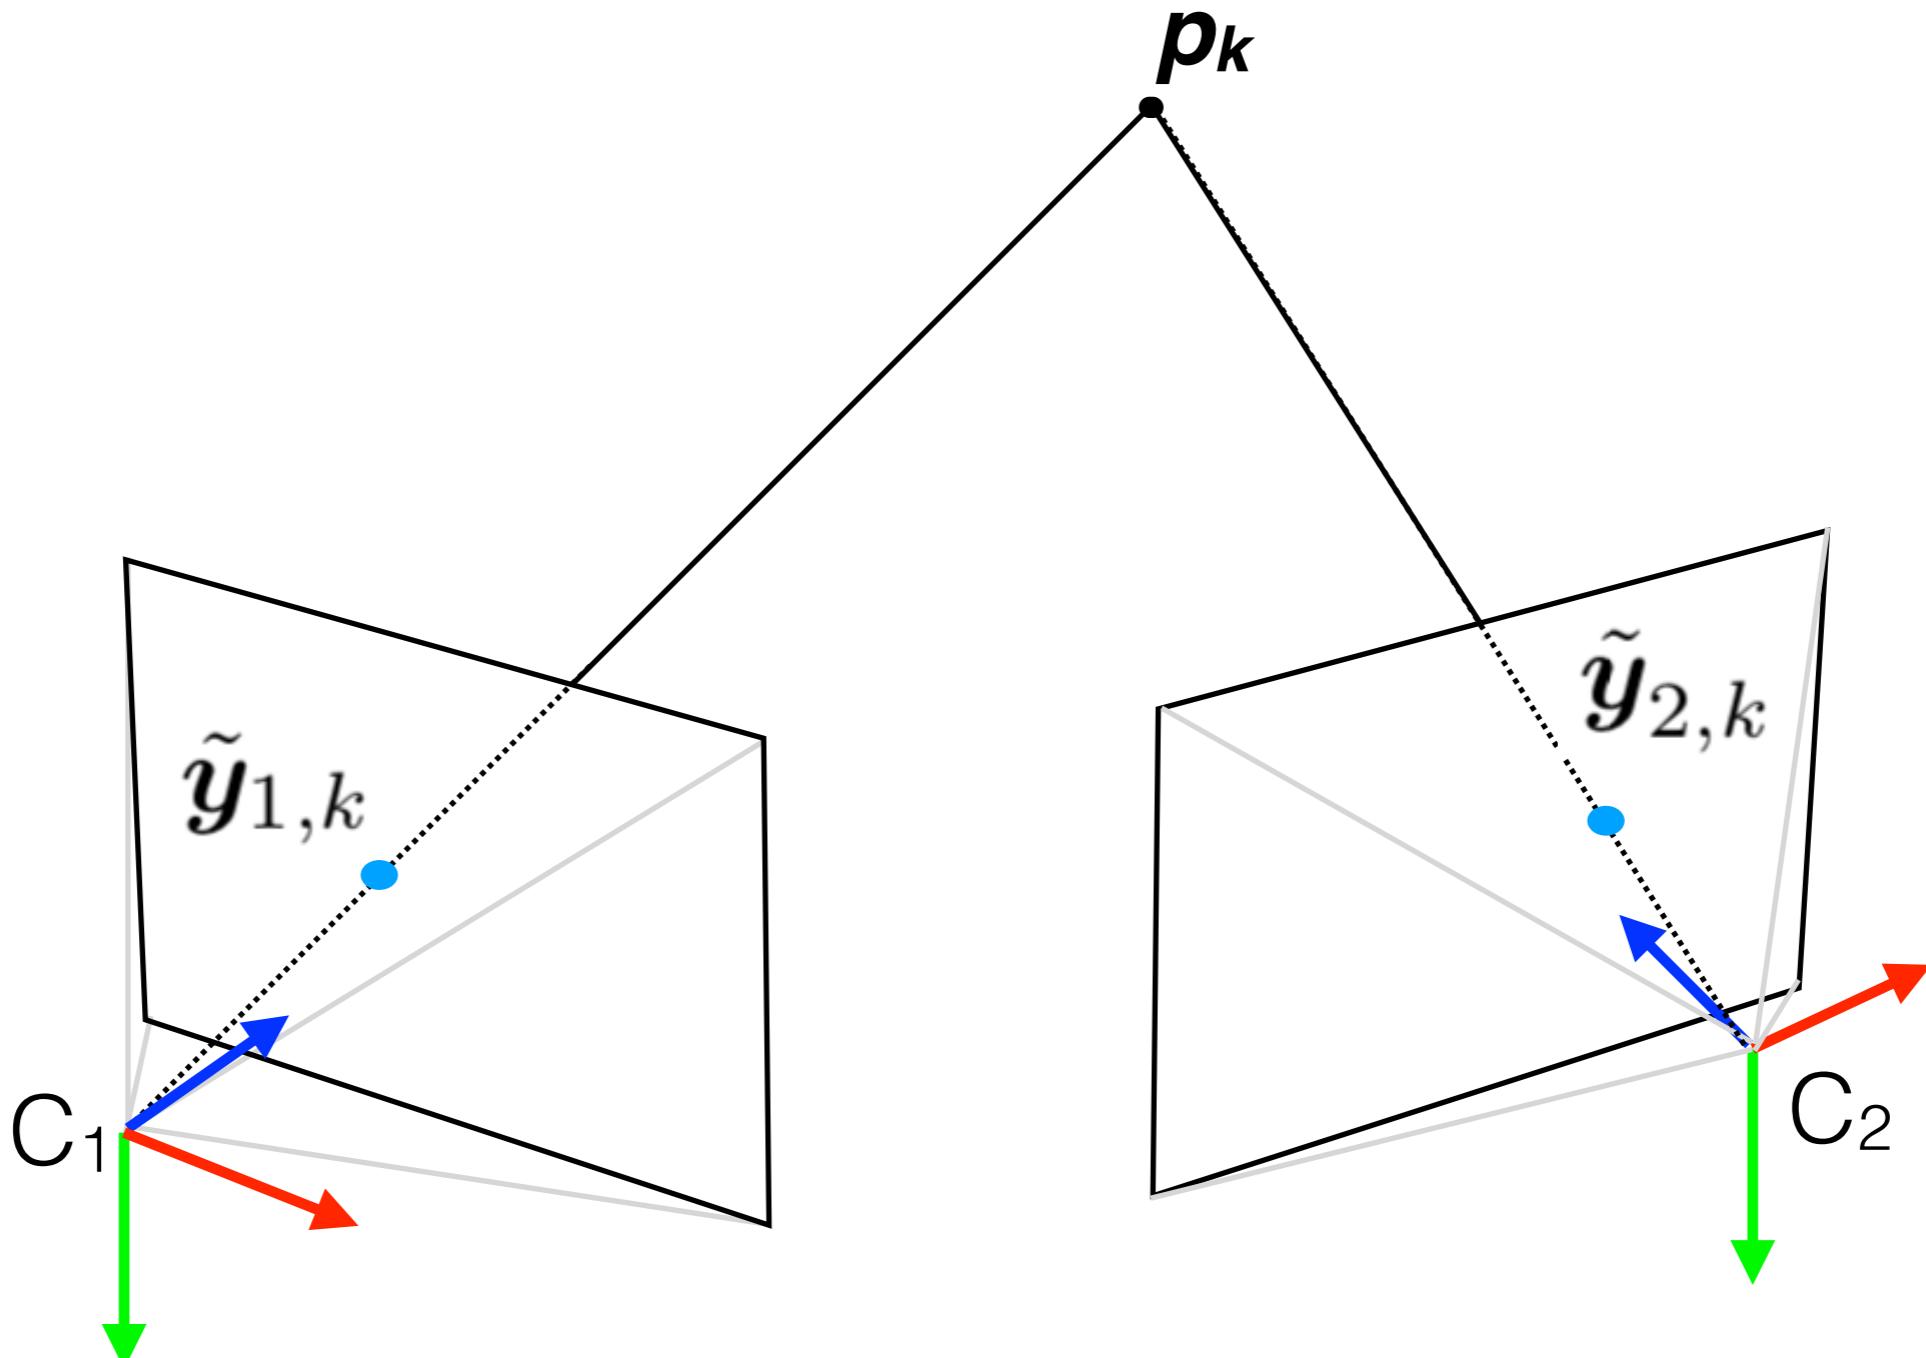
2-view Geometry



**Monday's
assumptions:**

- no wrong correspondences (outliers)
- 3D point is not moving
- camera calibration is known

2-view Geometry



Essential matrix encodes relative pose
(up to scale) between C_1 and C_2

Estimating Poses from Correspondences

Given N calibrated pixel correspondences:

$$(\tilde{\mathbf{y}}_{1,k}, \tilde{\mathbf{y}}_{2,k}) \text{ for } k = 1, \dots, N$$

1. leverage the epipolar constraints to estimate the essential matrix \mathbf{E}

$$\tilde{\mathbf{y}}_{2,k}^T \mathbf{E} \tilde{\mathbf{y}}_{1,k} = 0$$

For 8 points: $\mathbf{A}\mathbf{e} = 0$ $N > 8$ points: $\arg \min_{\|\mathbf{e}\|=1} \|\mathbf{A}\mathbf{e}\|^2$

2. Retrieve the rotation and translation (up to scale) from the \mathbf{E}

$$\mathbf{E} = [\mathbf{t}]_\times \mathbf{R}$$

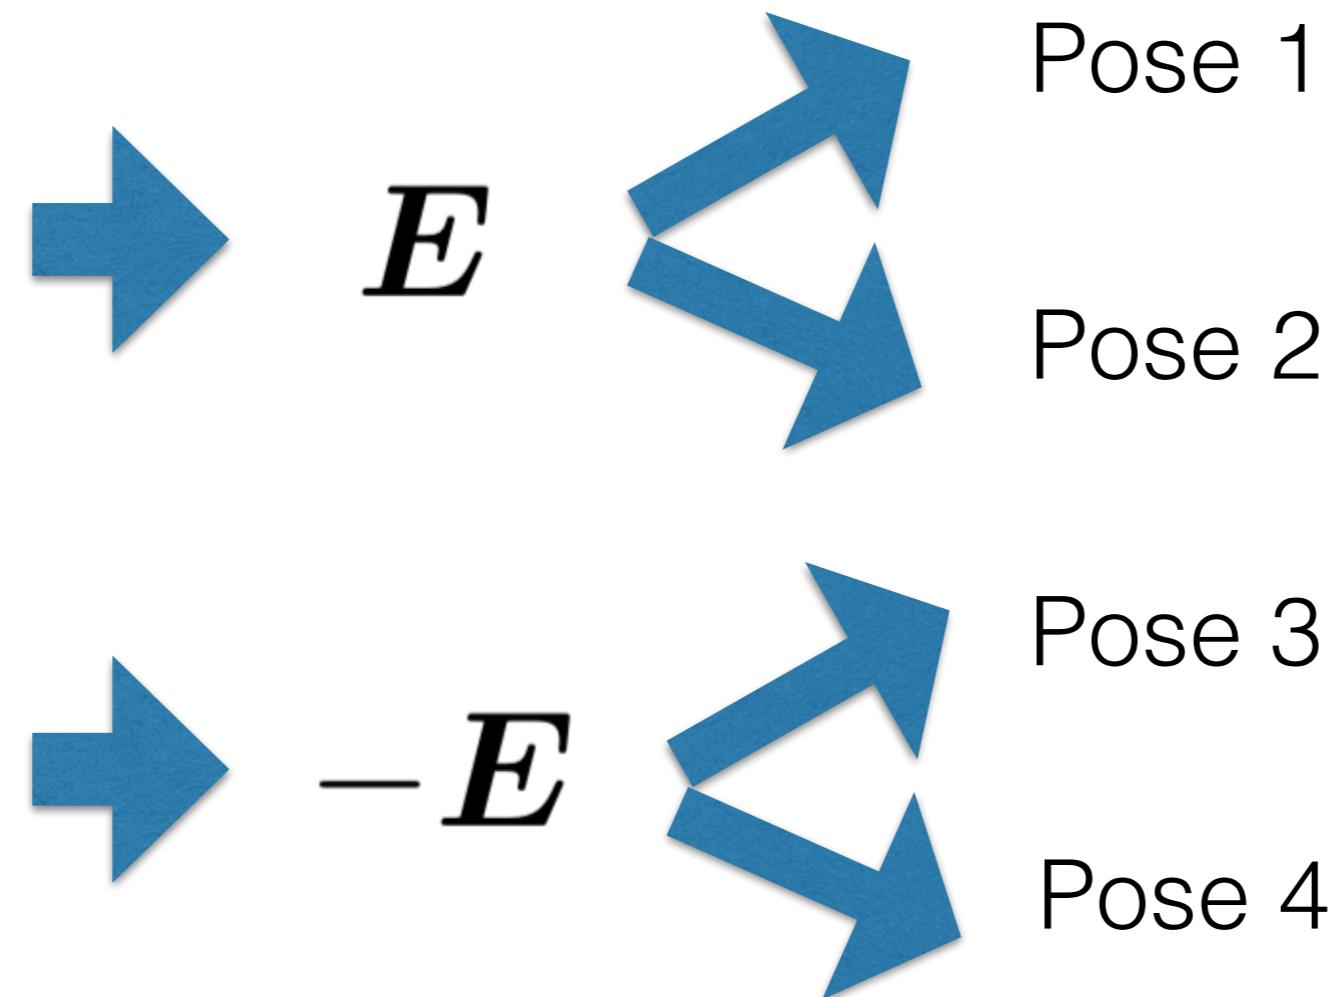
Retrieving Pose from Essential Matrix

Theorem 1 (Pose recovery from essential matrix, Thm 5.7 in [1]). *There exist exactly two relative poses (\mathbf{R}, \mathbf{t}) with $\mathbf{R} \in \text{SO}(3)$ and $\mathbf{t} \in \mathbb{R}^3$ corresponding to a nonzero essential matrix \mathbf{E} (i.e., such that $\mathbf{E} = [\mathbf{t}]_\times \mathbf{R}$):*

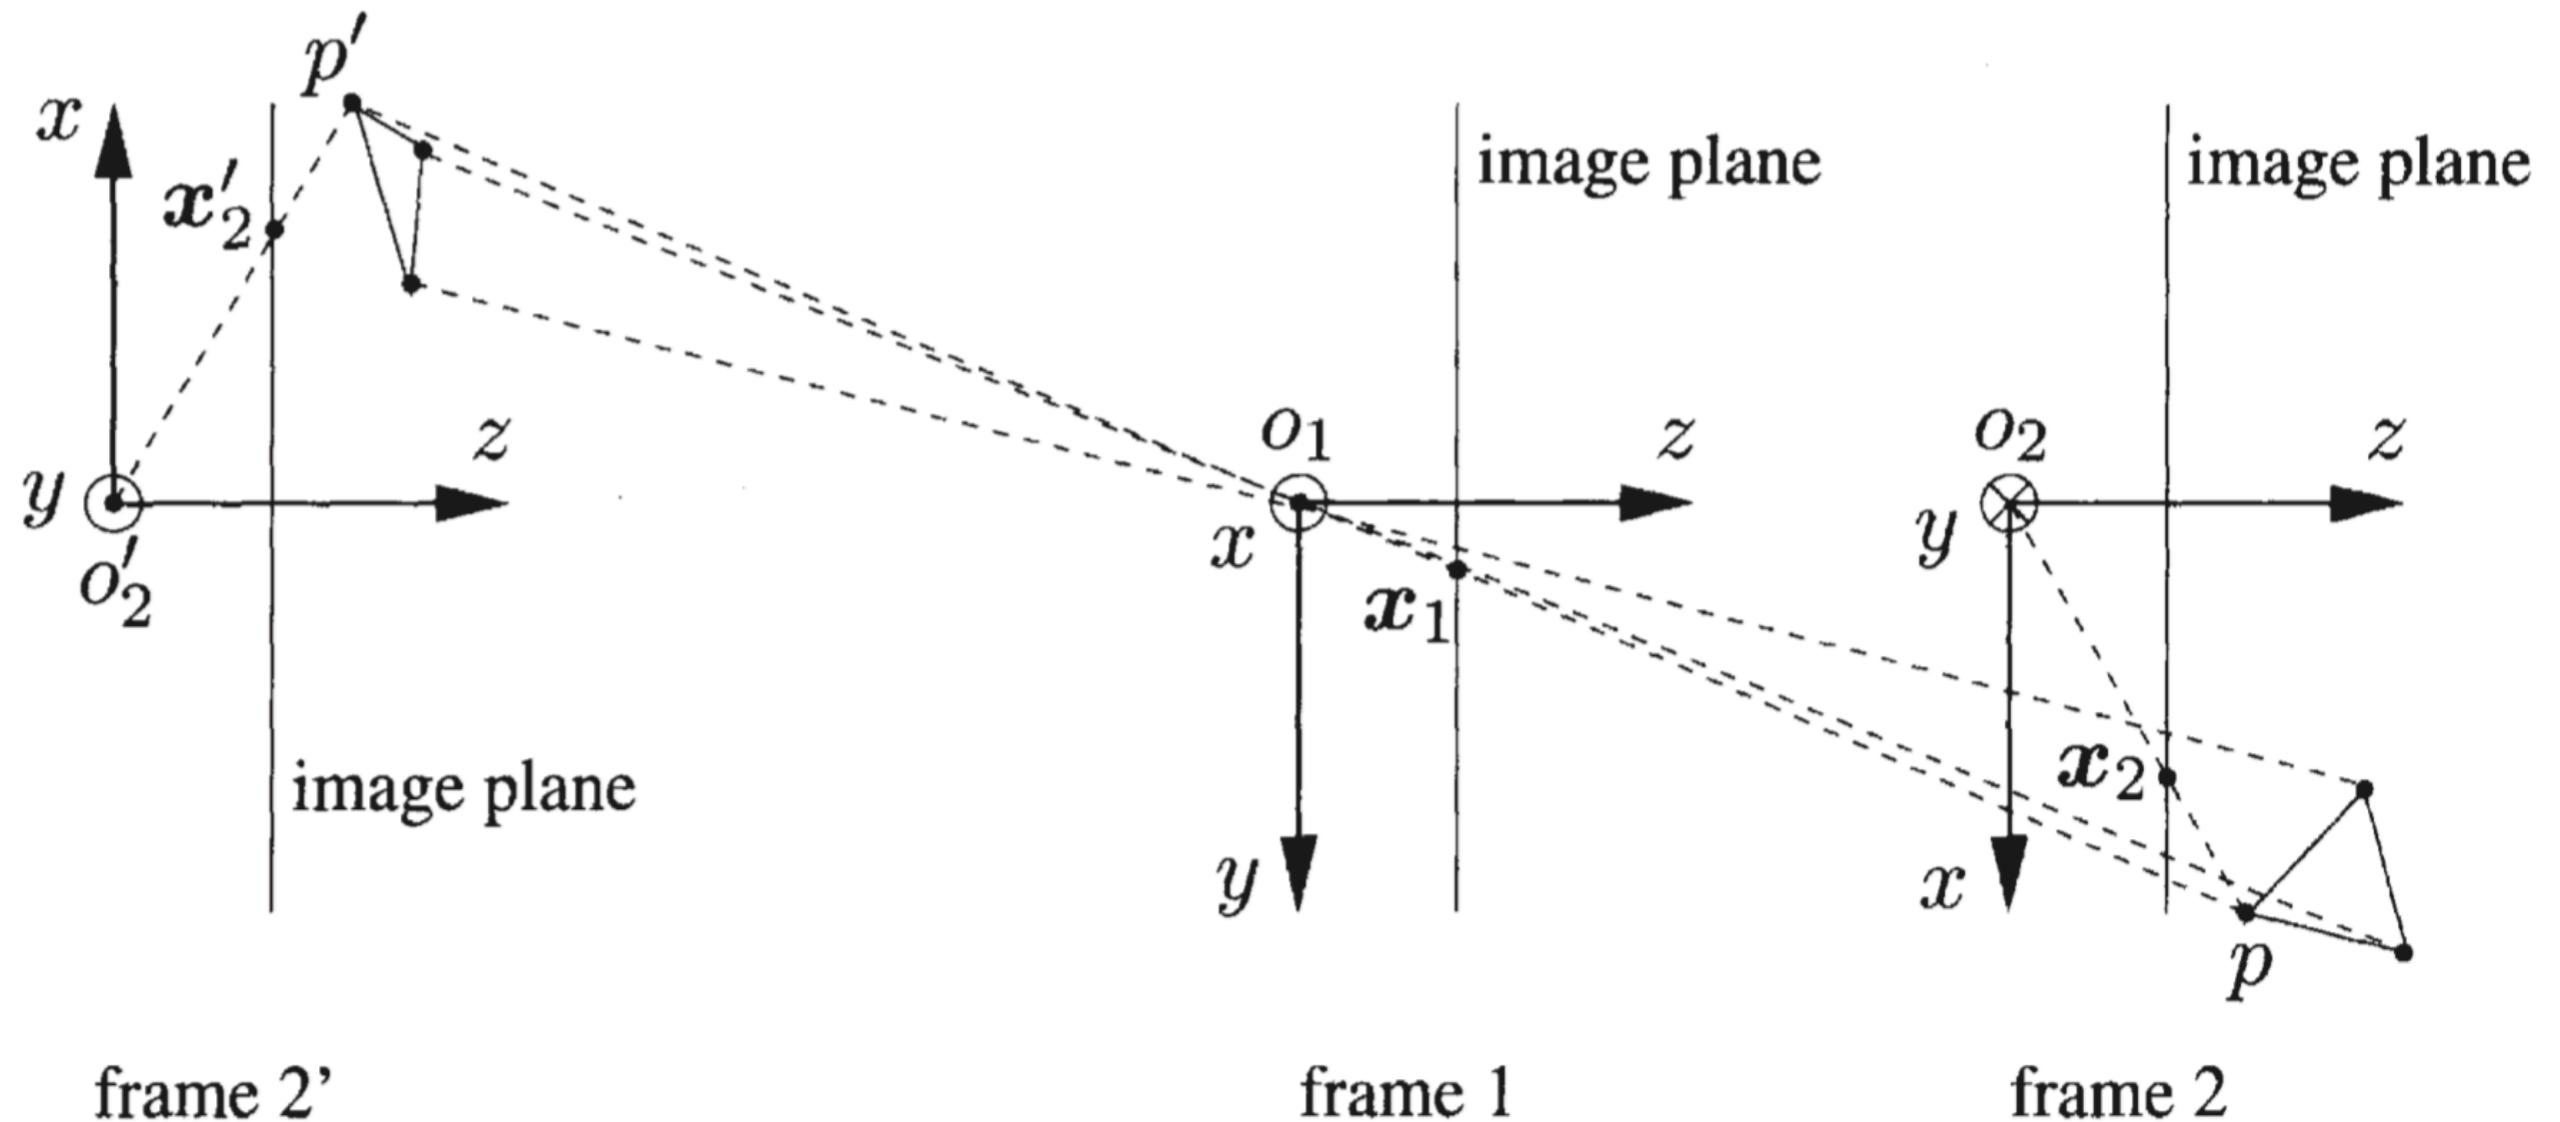
$$\mathbf{t}_1 = \mathbf{U} \mathbf{R}_z(+\pi/2) \mathbf{\Sigma} \mathbf{U}^\top \quad \mathbf{R}_1 = \mathbf{U} \mathbf{R}_z(+\pi/2) \mathbf{V}^\top \quad (13.19)$$

$$\mathbf{t}_2 = \mathbf{U} \mathbf{R}_z(-\pi/2) \mathbf{\Sigma} \mathbf{U}^\top \quad \mathbf{R}_2 = \mathbf{U} \mathbf{R}_z(-\pi/2) \mathbf{V}^\top \quad (13.20)$$

where $\mathbf{E} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^\top$ is the singular value decomposition of the matrix \mathbf{E} , and $\mathbf{R}_z(+\pi/2)$ is an elementary rotation around the z -axis of an angle $\pi/2$.

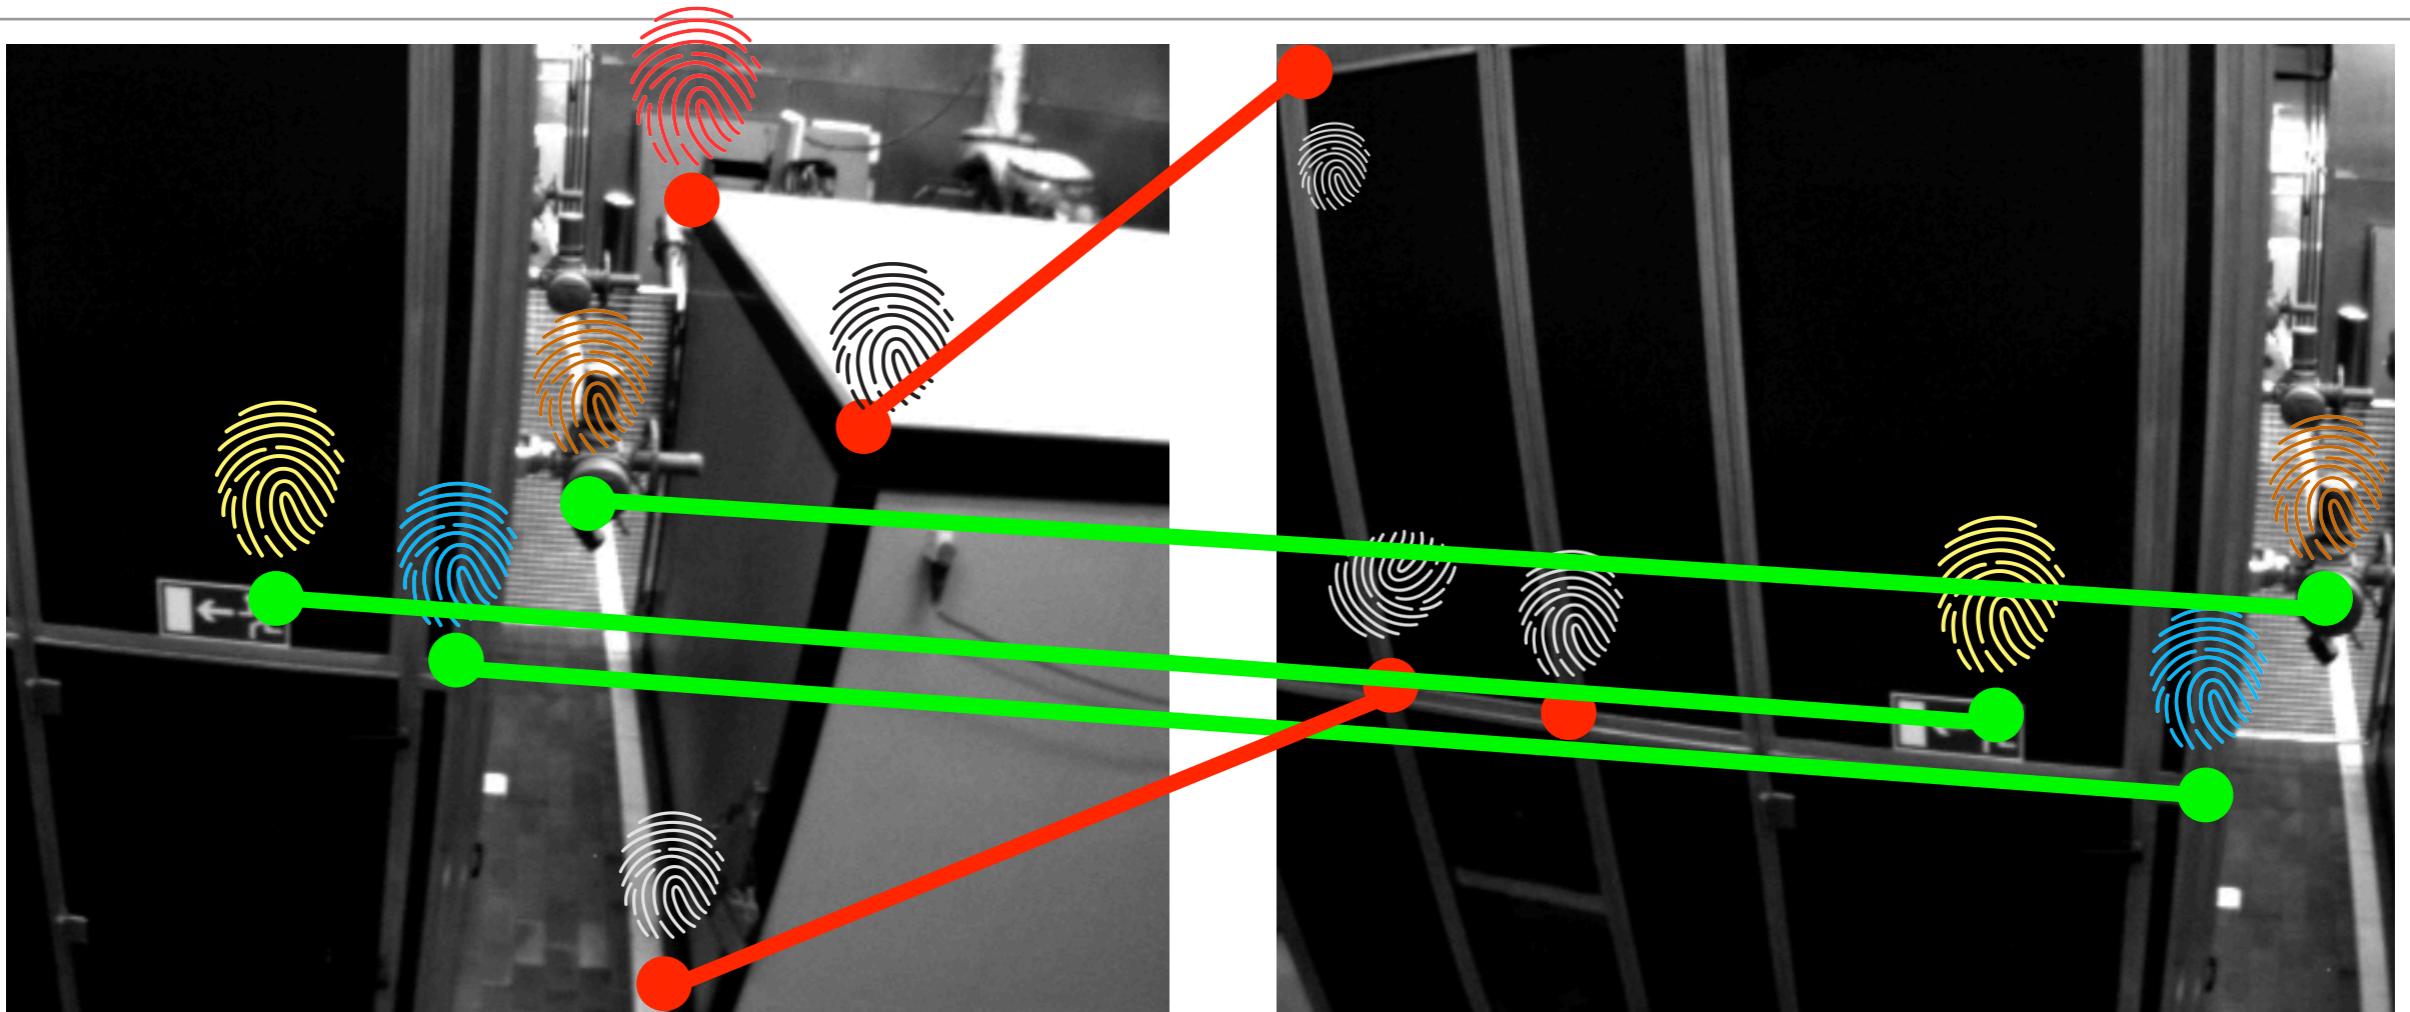


Chirality constraints



Points must be in front of the cameras!

2-view Geometry



In practice:

- Many wrong correspondences (outliers)
- Some 3D points might be moving

RANSAC

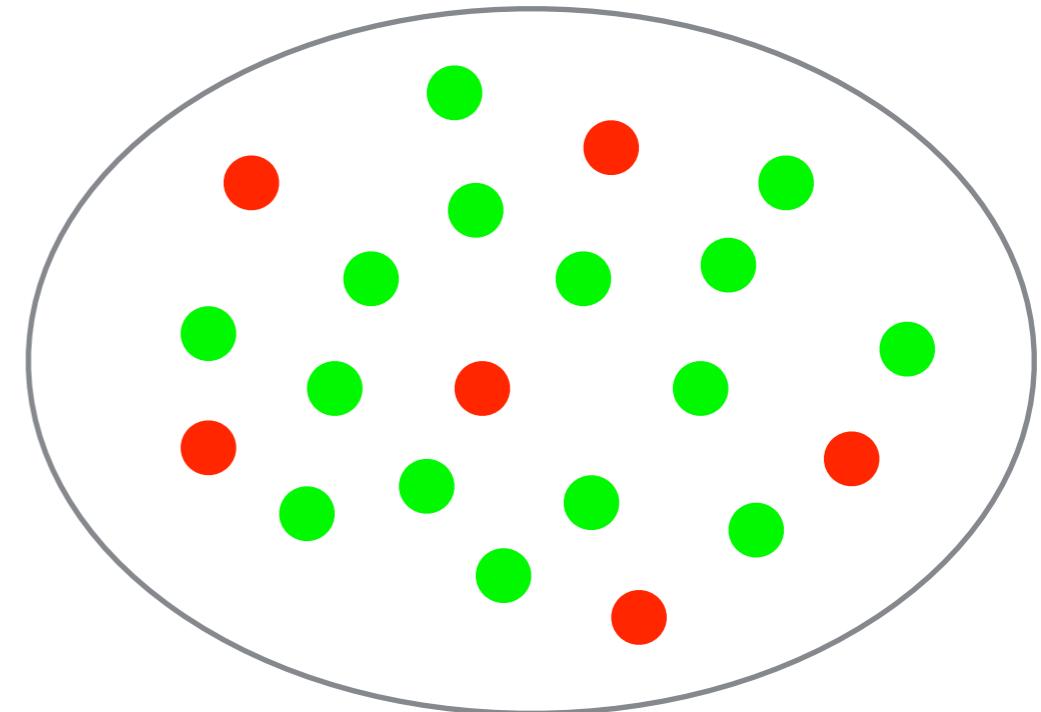
RANdom SAmple Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points



RANSAC

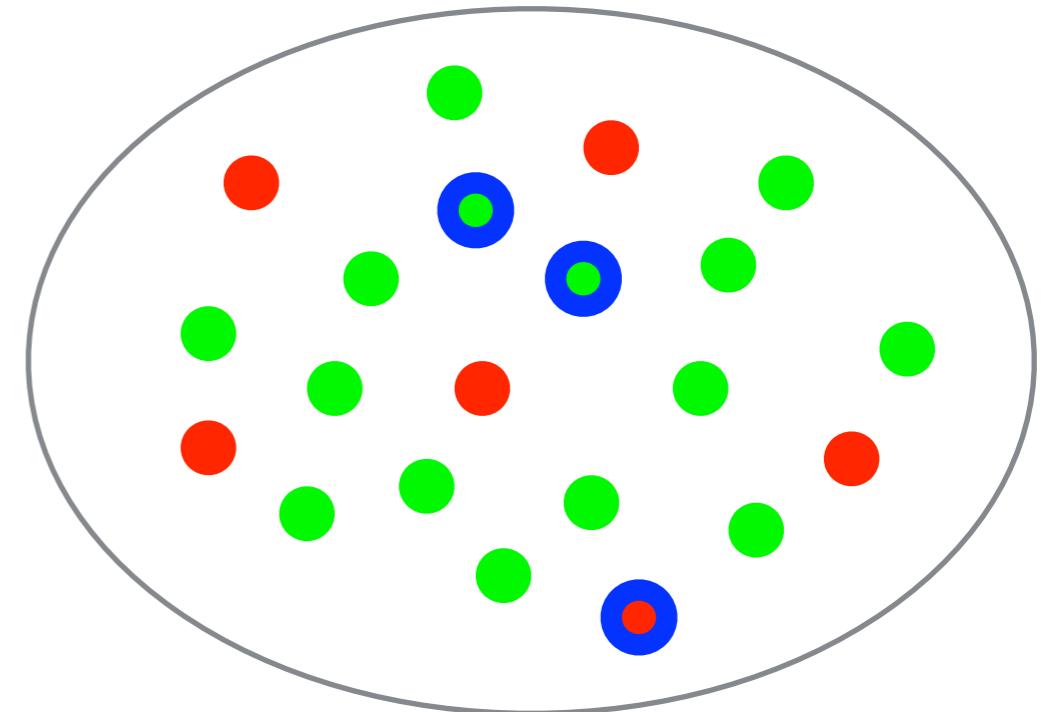
RANdom SAmple Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points

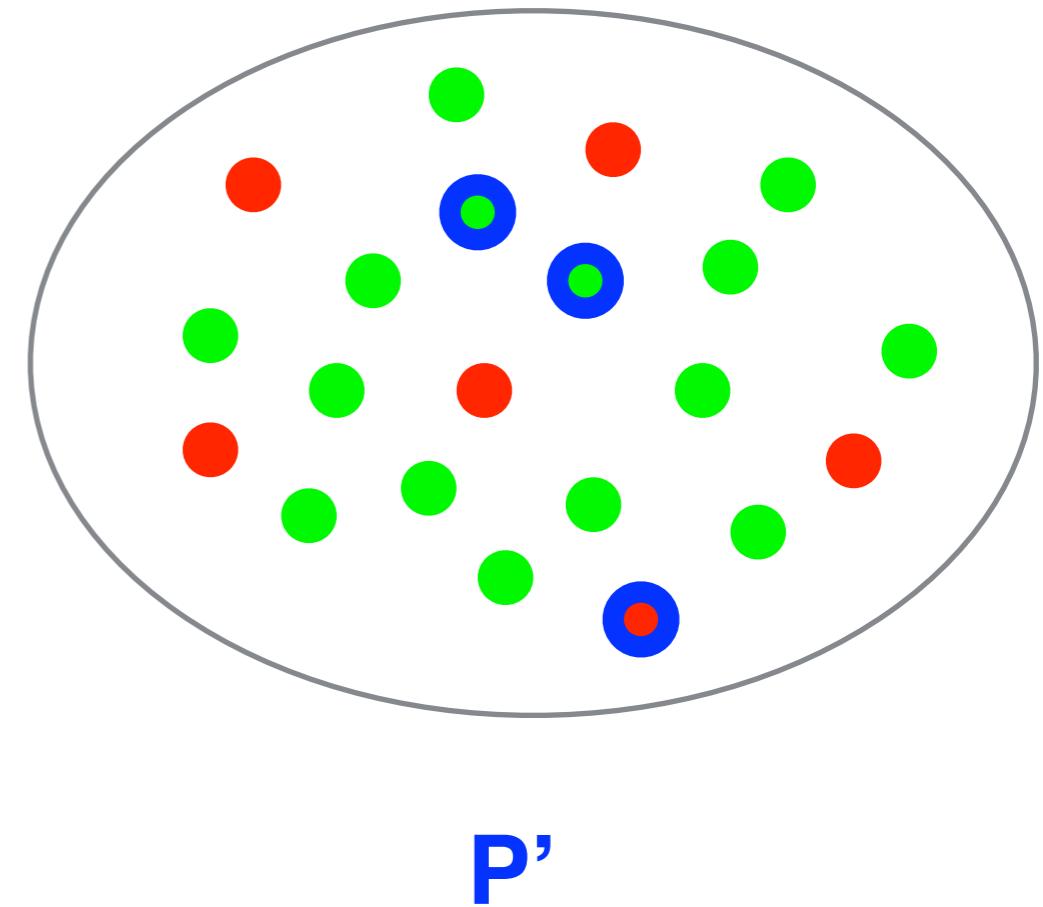


RANSAC

RANdom SAmples Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)



Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points

RANSAC

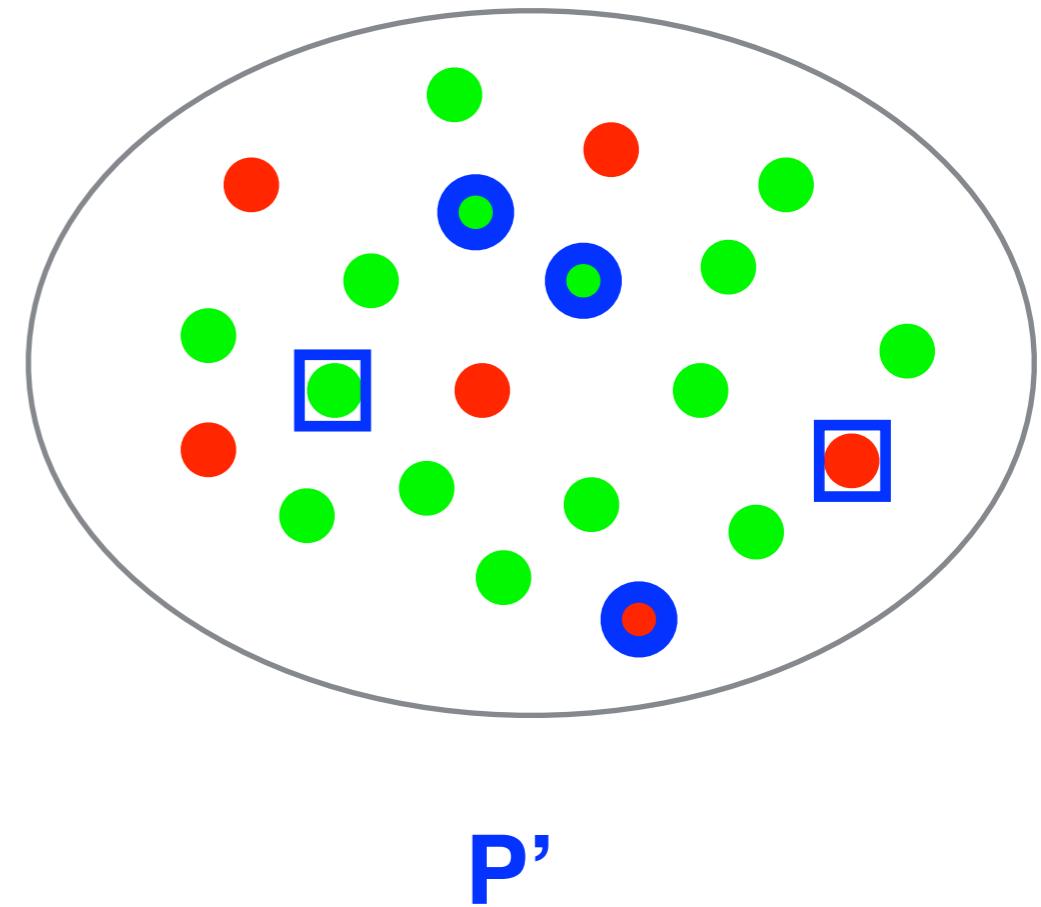
RANdom SAmples Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points



Consensus Set

RANSAC

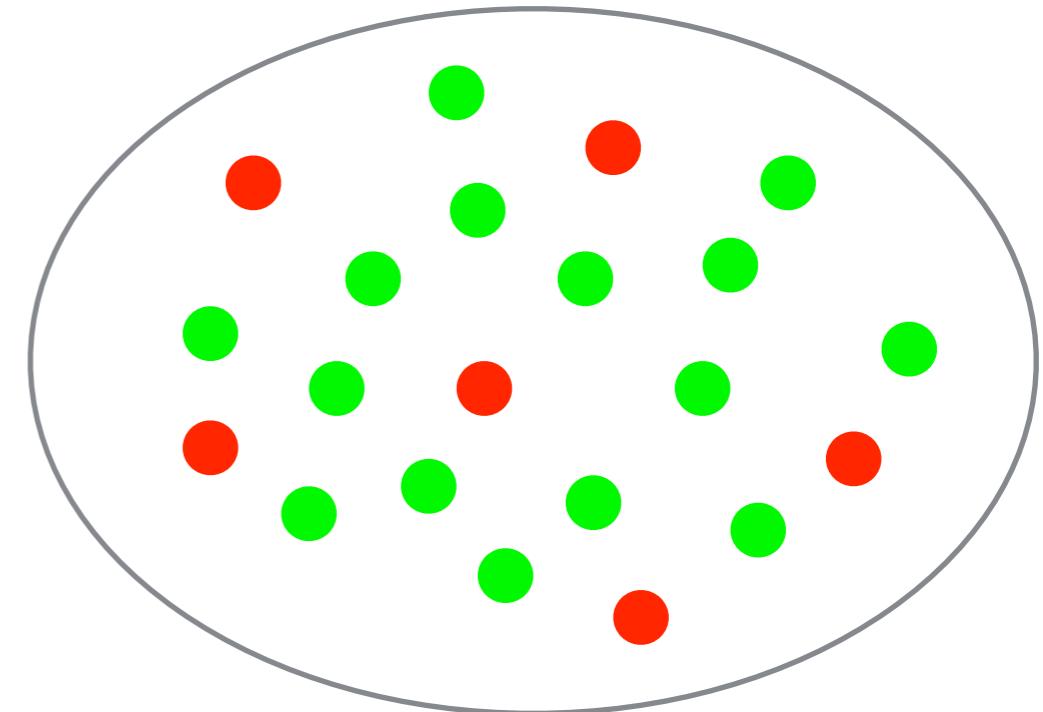
RANdom SAmples Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points



RANSAC

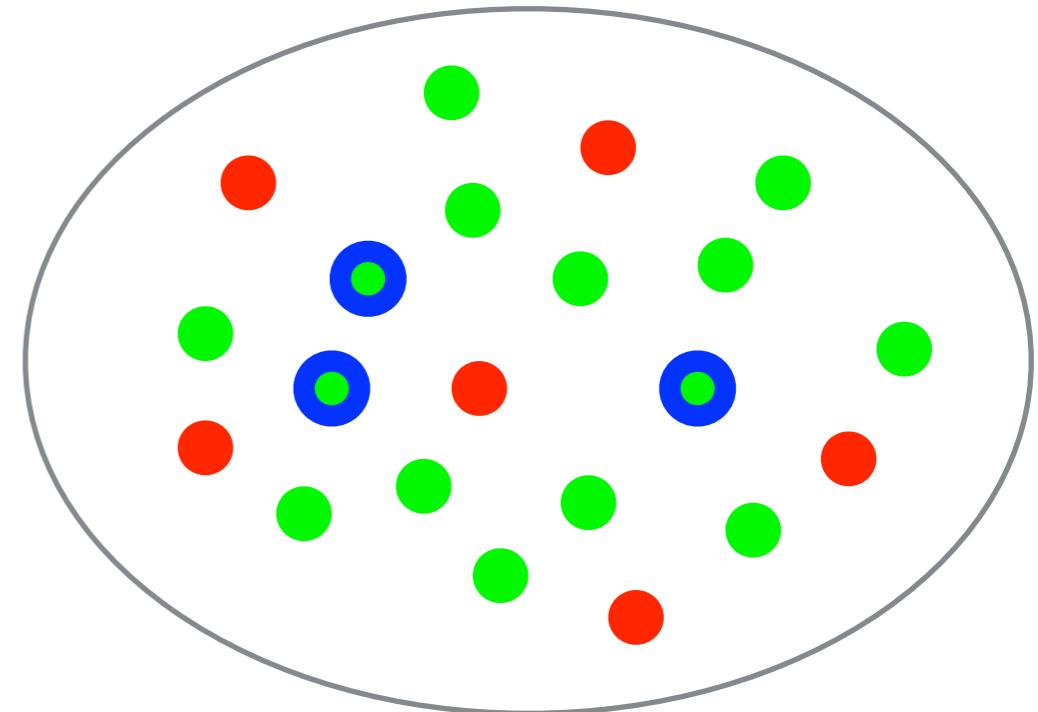
RANdom SAmples Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points

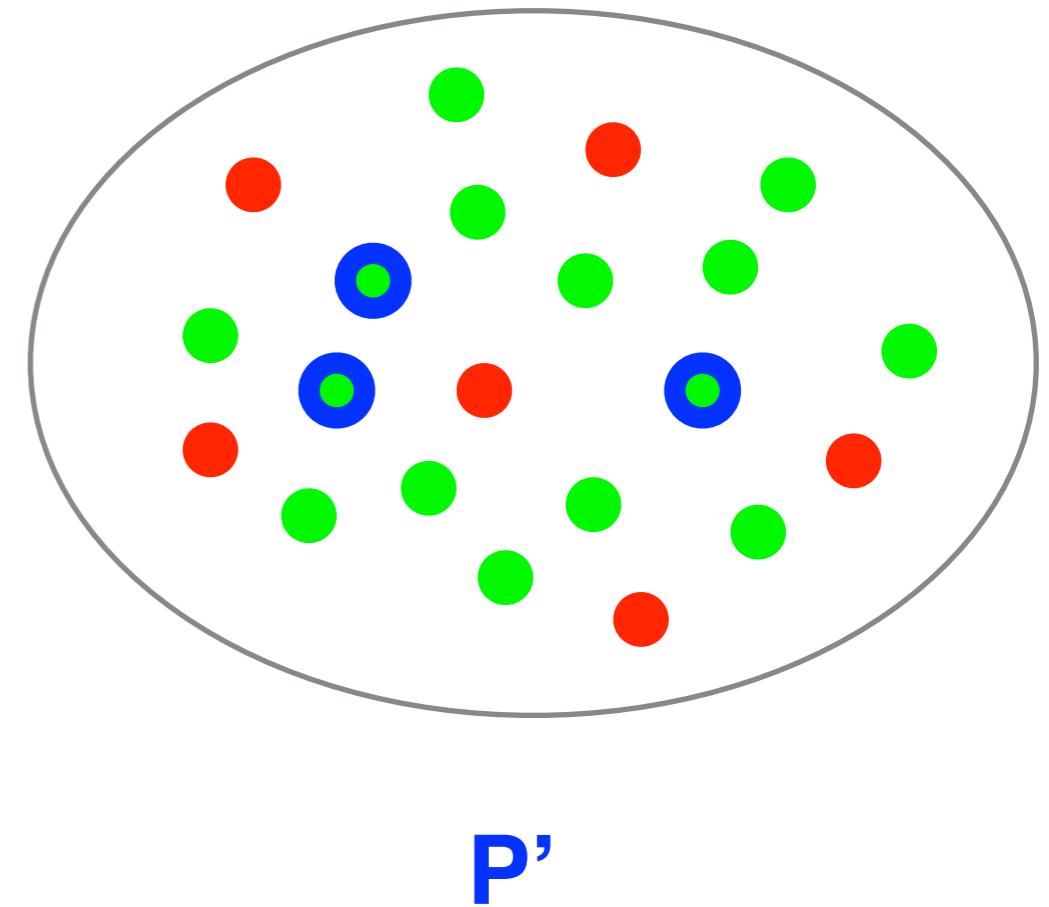


RANSAC

RANdom SAmples Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)



Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points

RANSAC

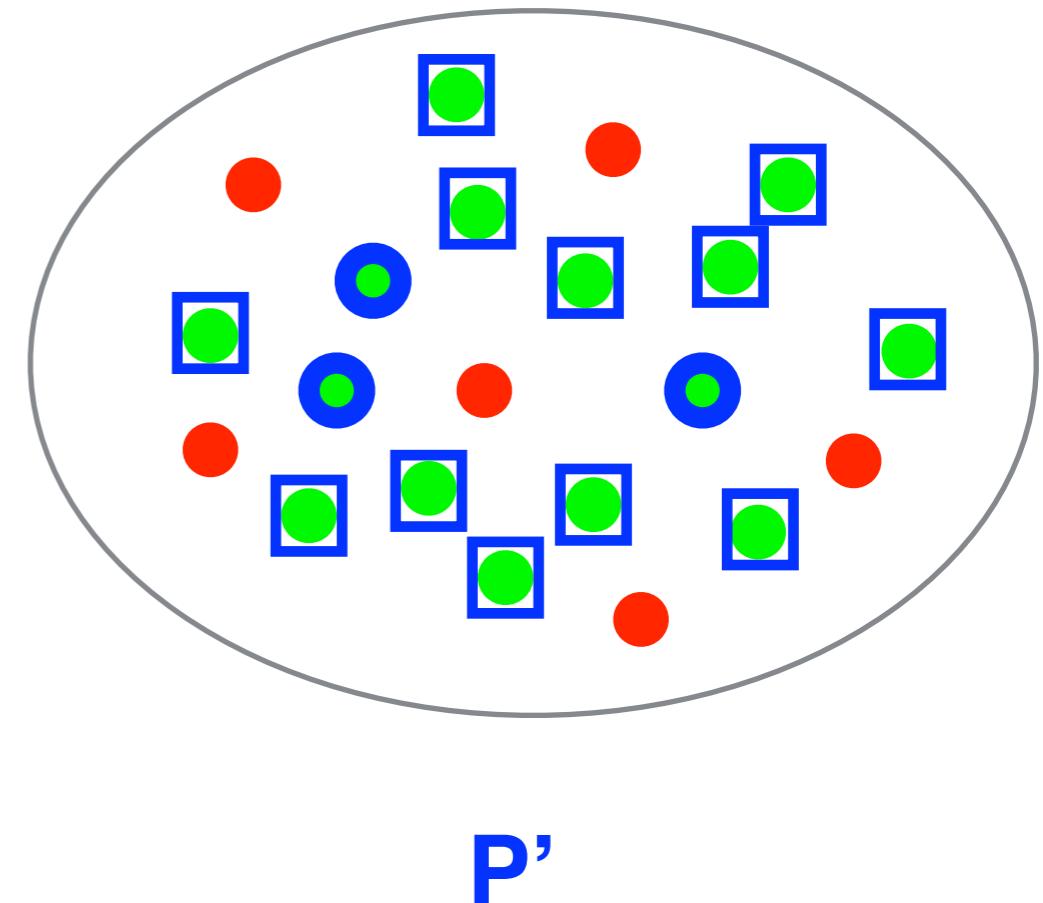
RANdom SAmple Consensus

Problem: estimate model P from N data points, possibly corrupted with outliers.

Assume: we have an algorithm to estimate P from n data points ($n \ll N$)

Basic idea:

1. sample n points
2. compute an estimate P' of P
3. count how many other points agree with P'
4. repeat until you get a P' that agrees with many points

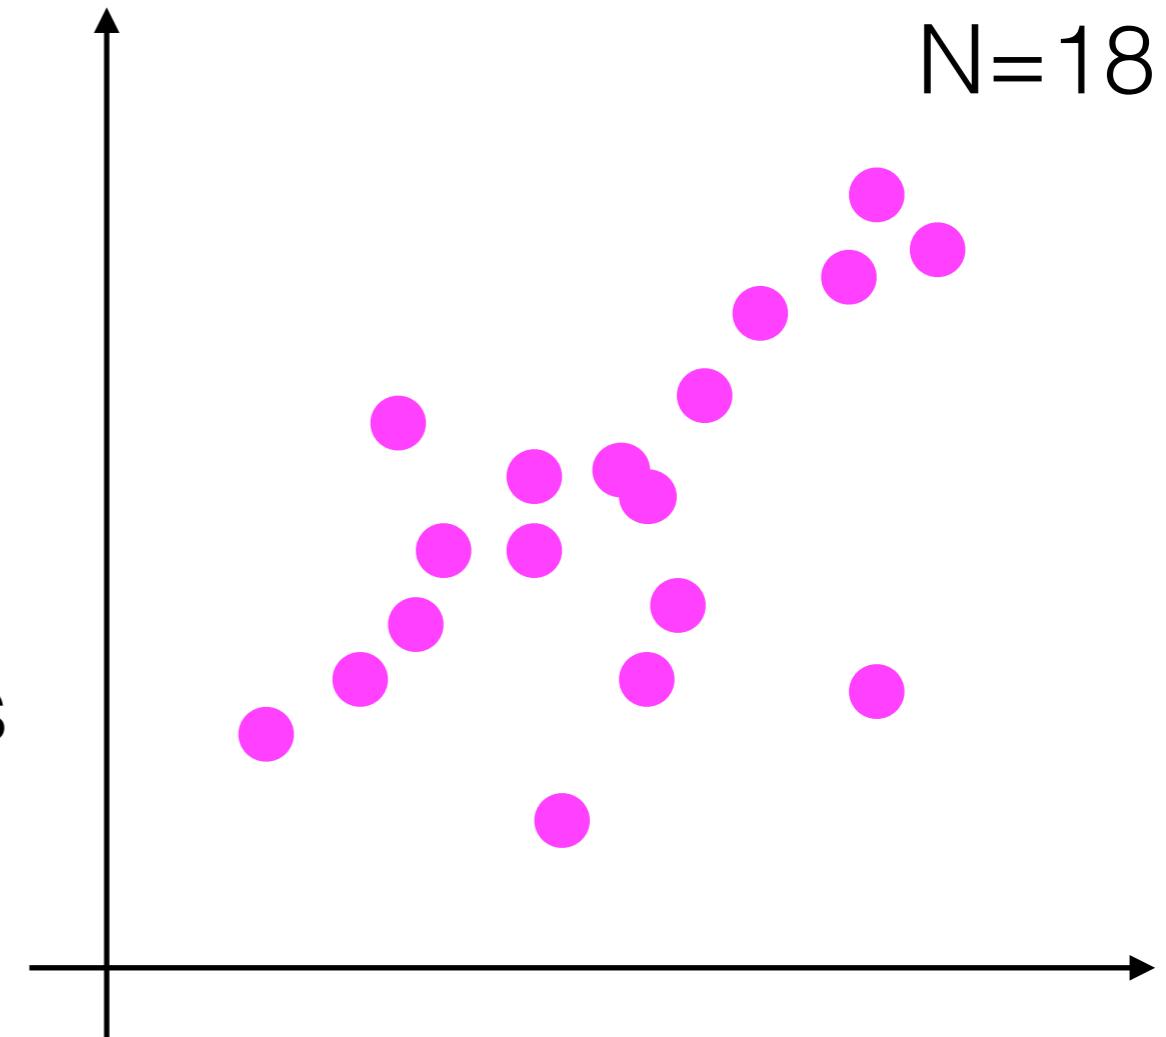


Consensus Set

Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



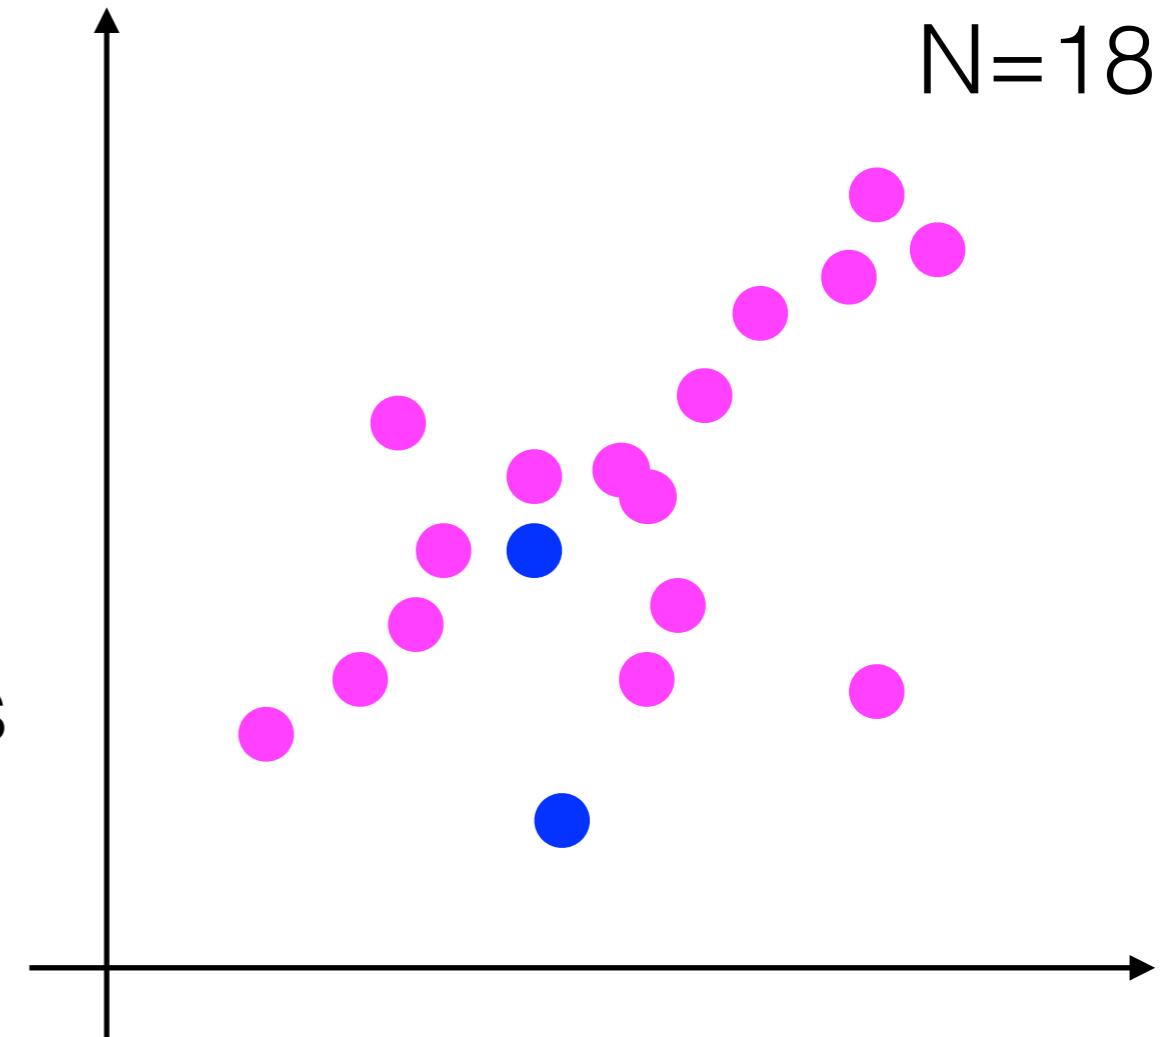
RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points

Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points

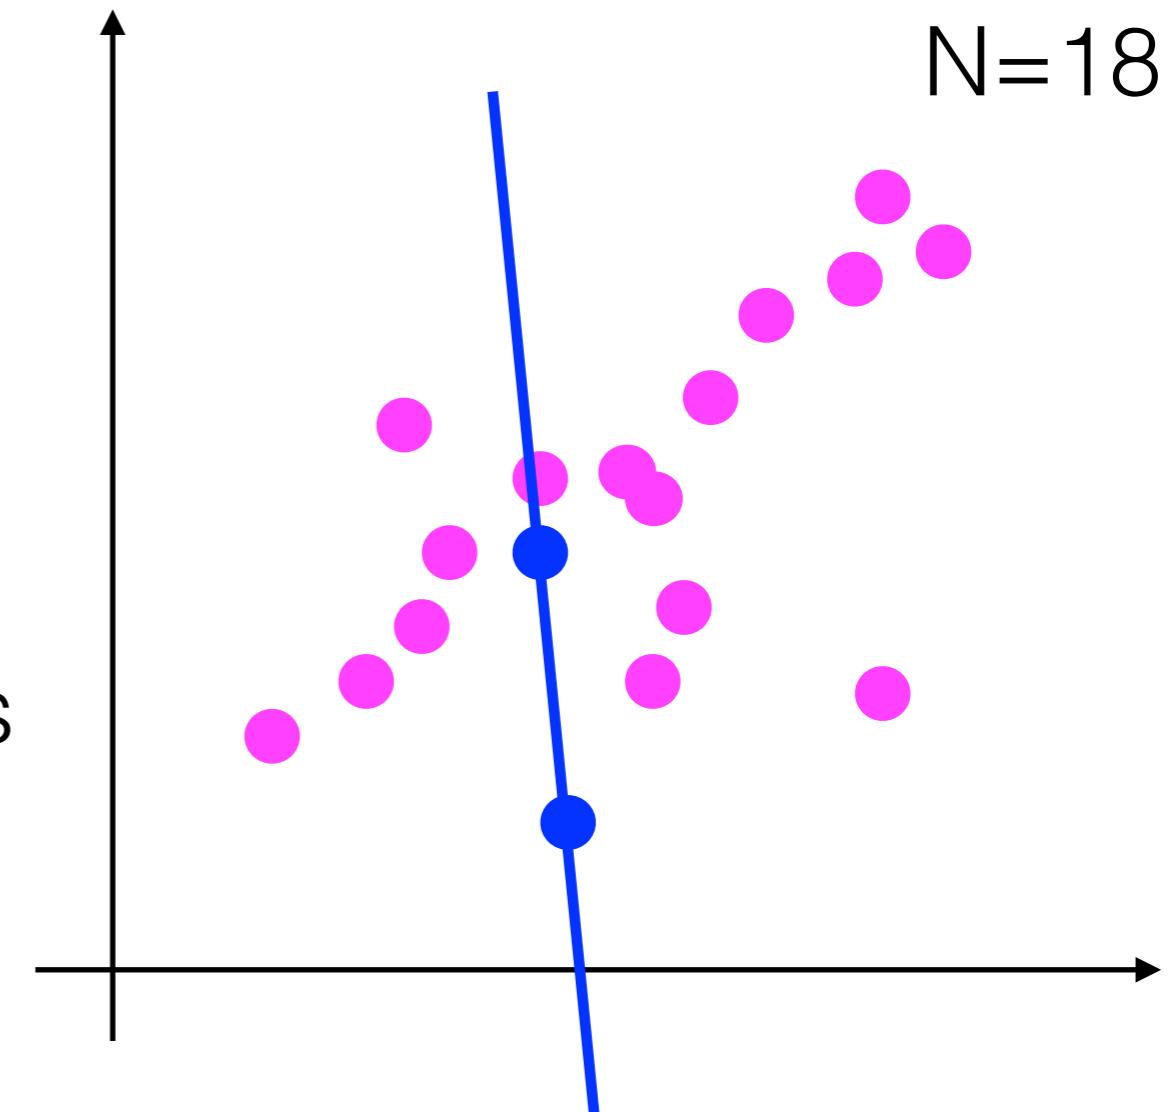
Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points

RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points



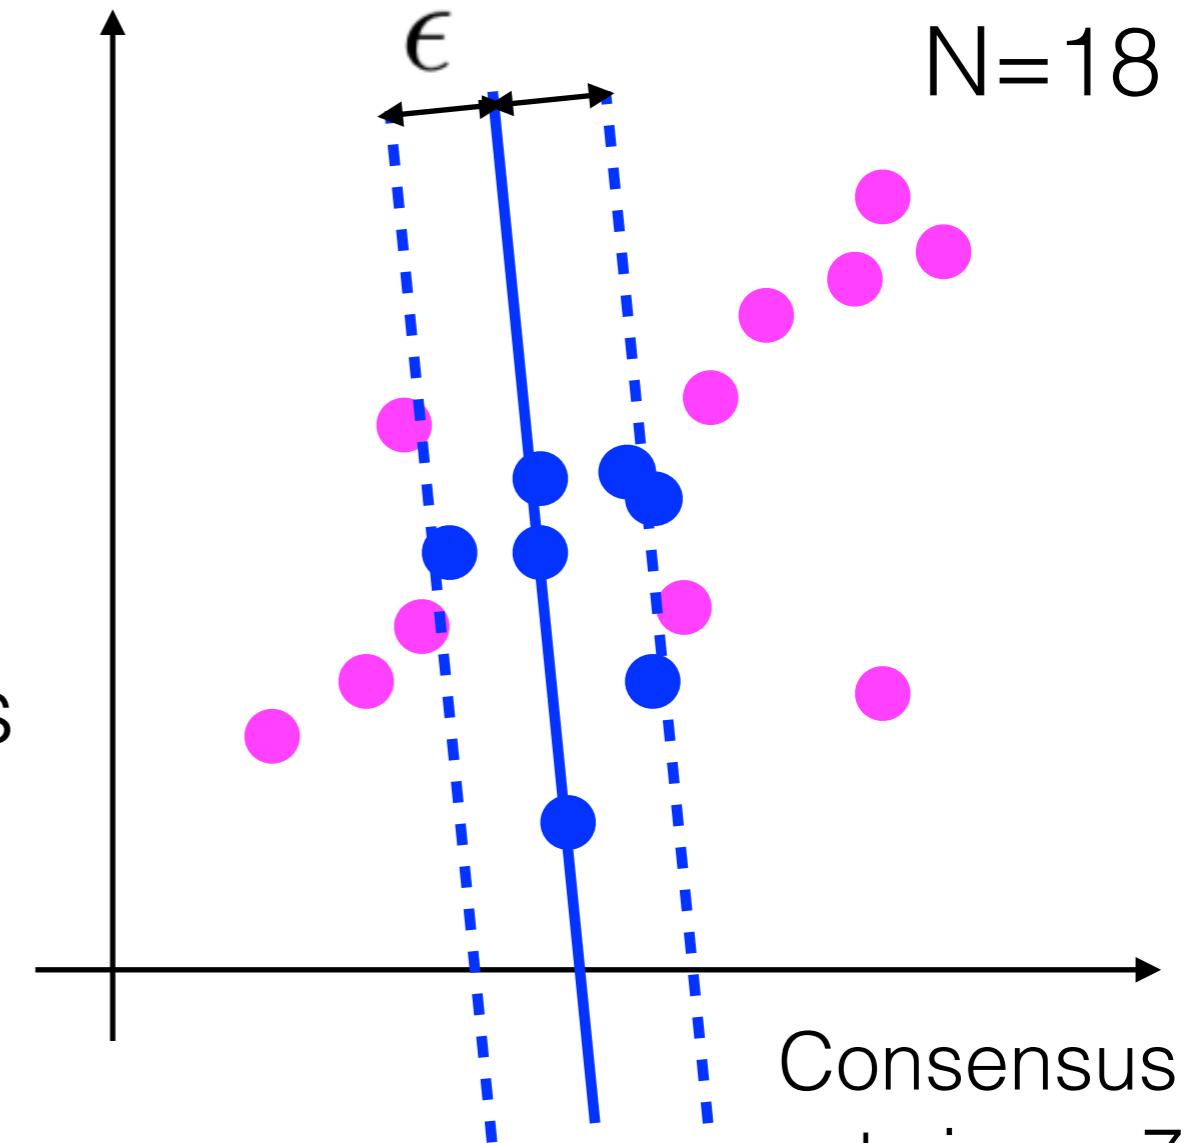
Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points

RANSAC:

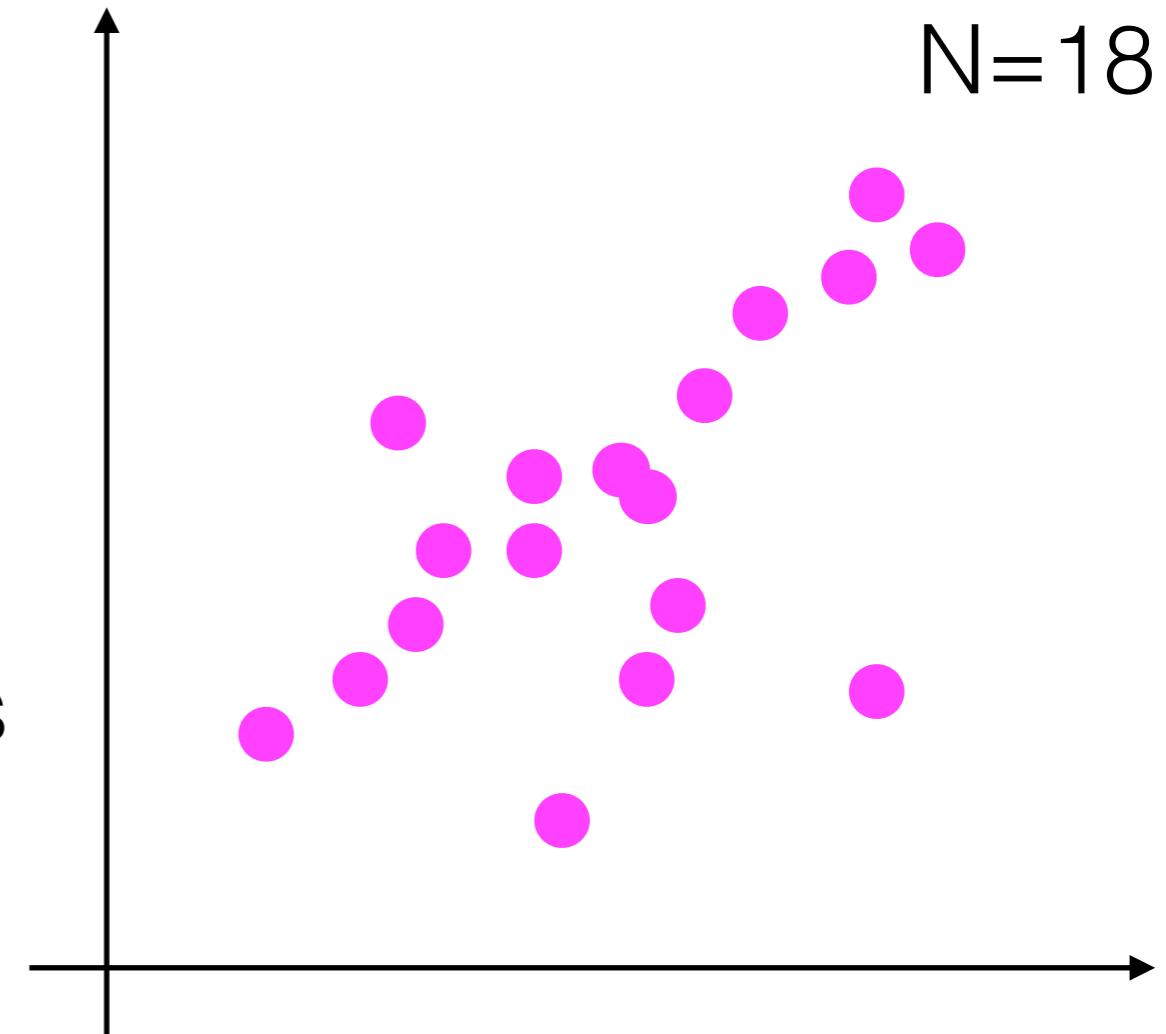
1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points



Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



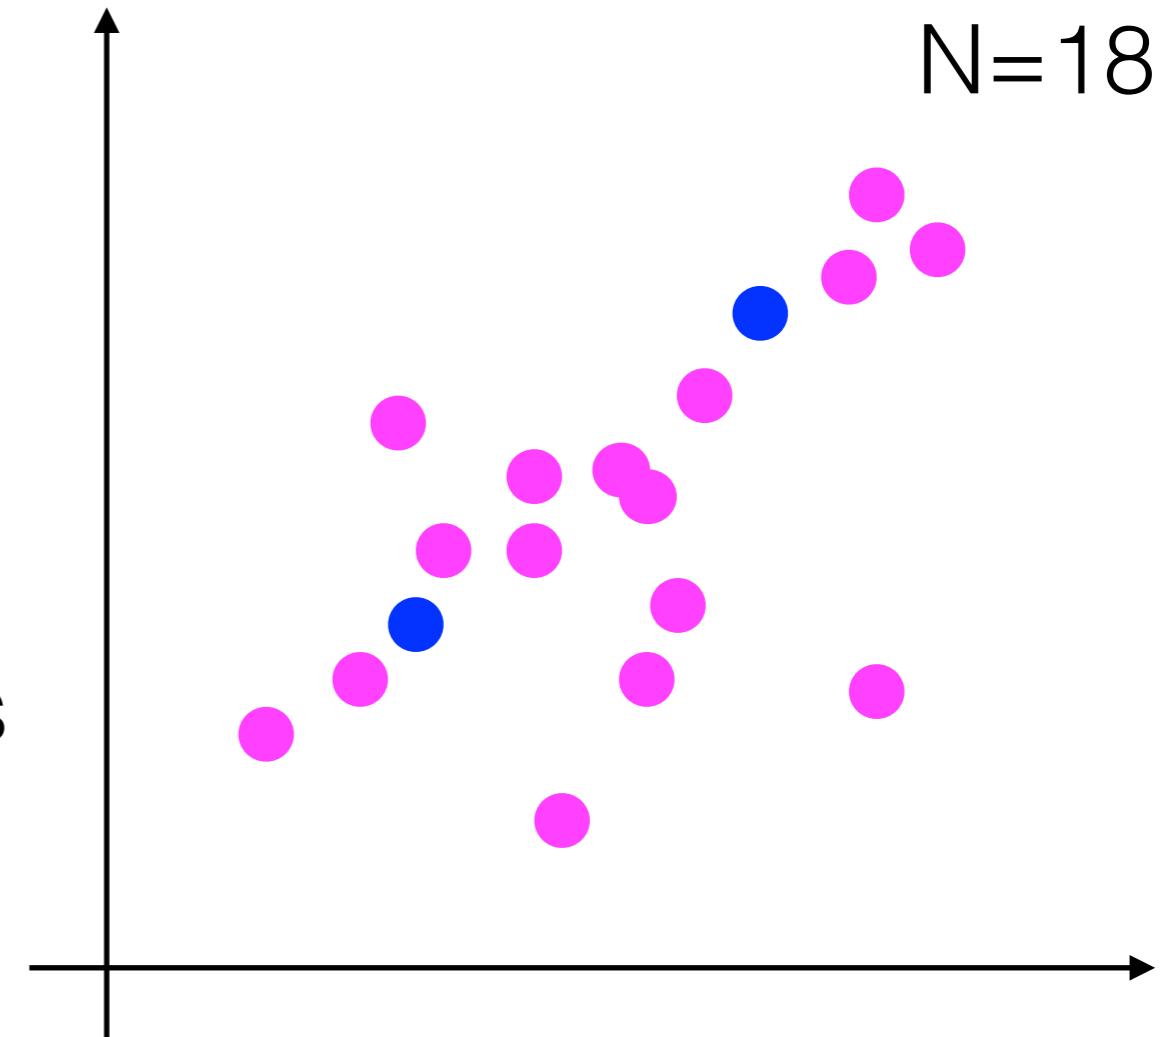
RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points

Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



RANSAC:

1. sample 2 points

2. compute a line estimate P' of P

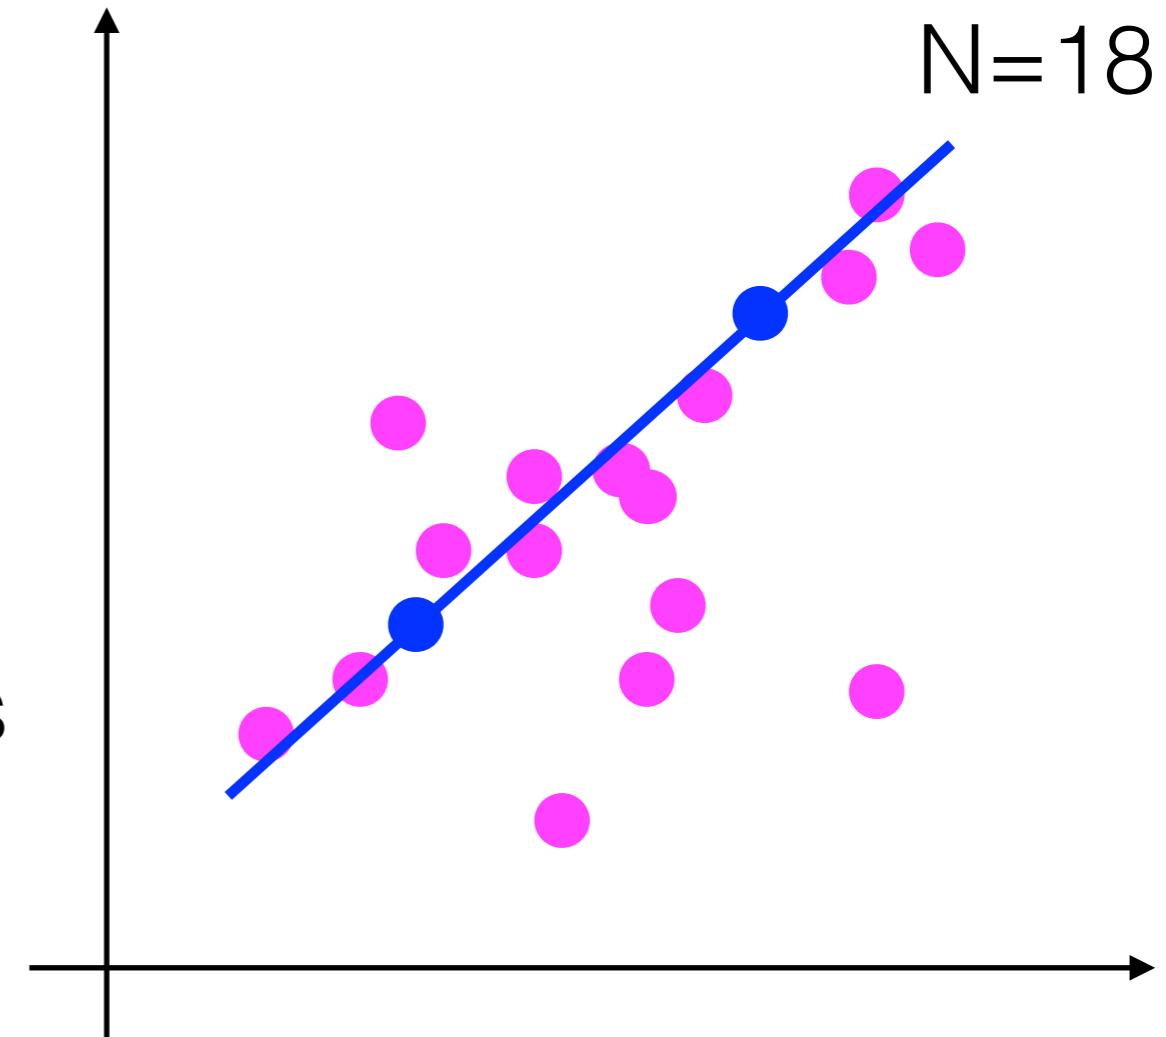
3. count how many points are within a **tolerance** from P'

4. repeat until you get a P' that agrees with many points

Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



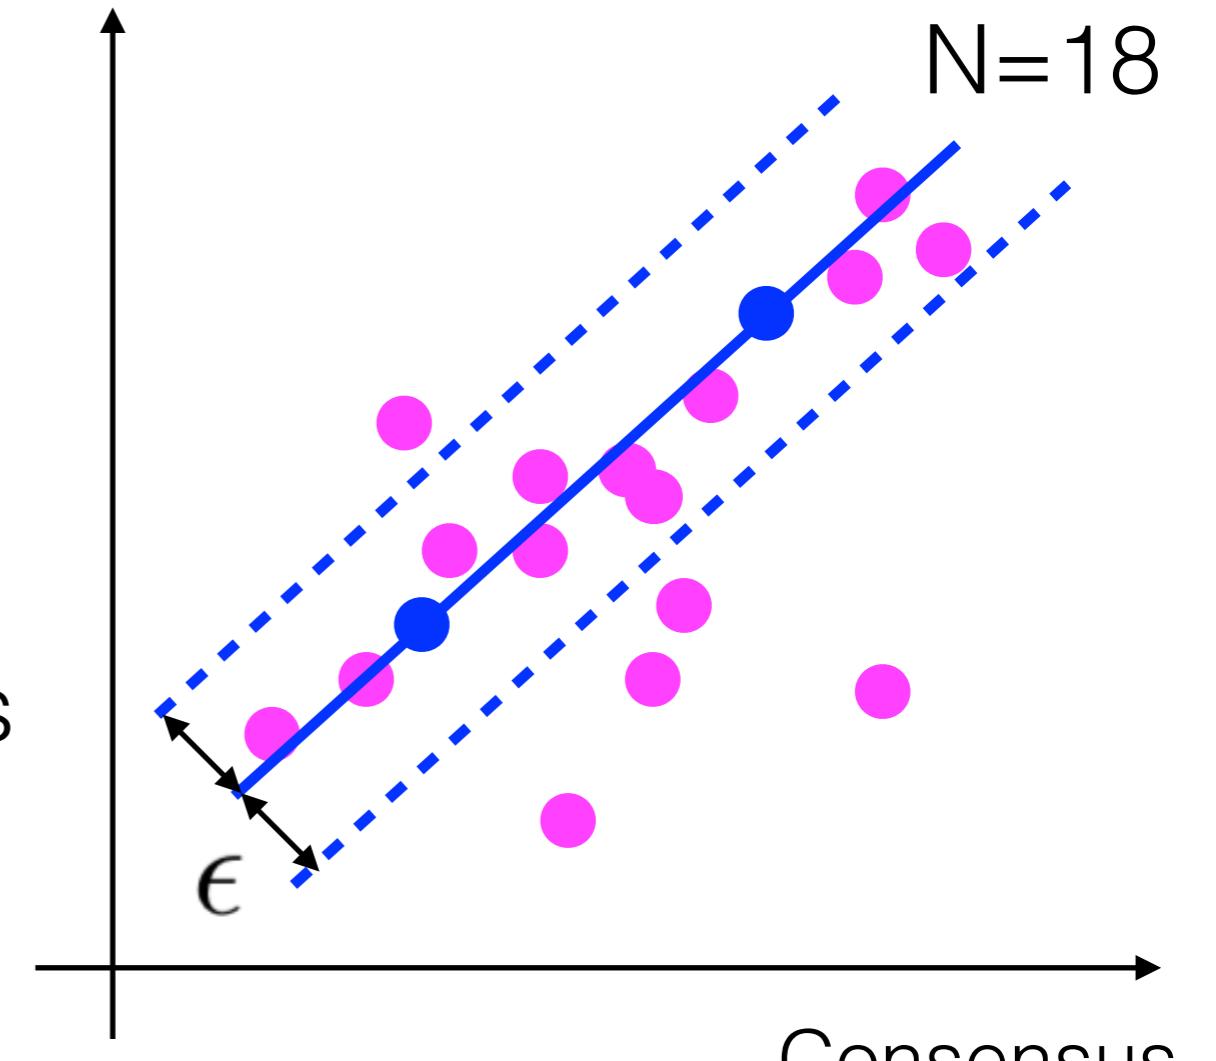
RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points

Example: Linear Regression

Fit a line through
N 2D points, possibly
corrupted with outliers.

Note: we have an algorithm
to estimate a line from $n=2$ points



RANSAC:

1. sample 2 points
2. compute a line estimate P' of P
3. count how many points are within a **tolerance** from P'
4. repeat until you get a P' that agrees with many points

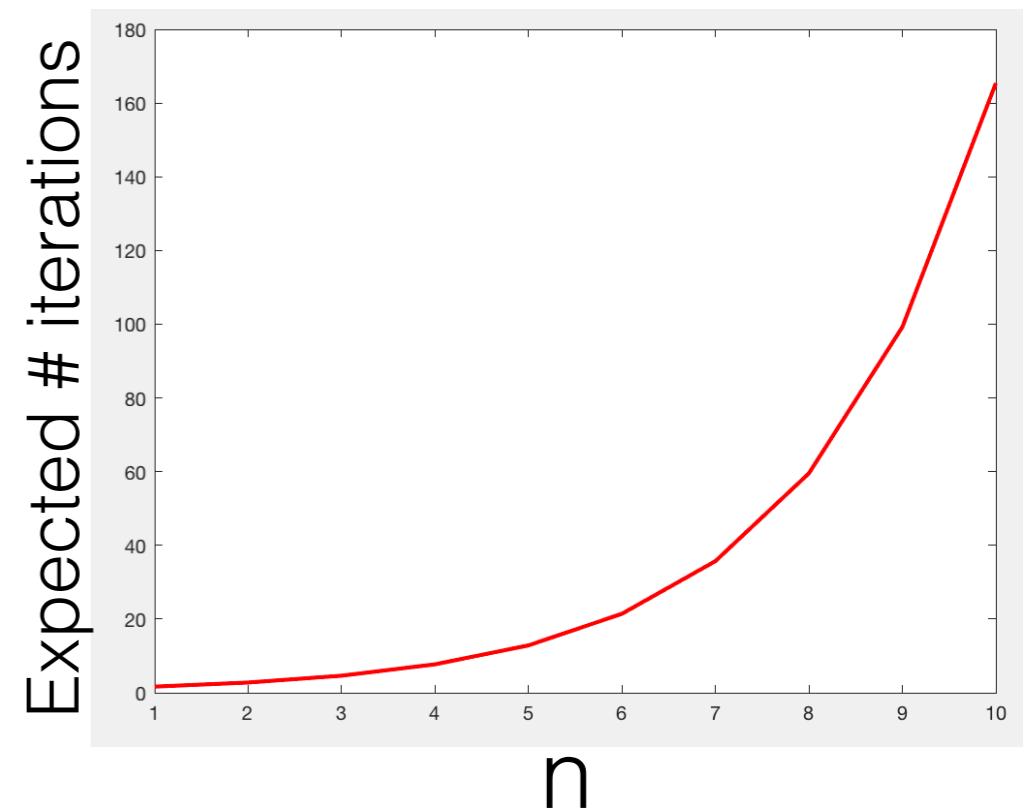
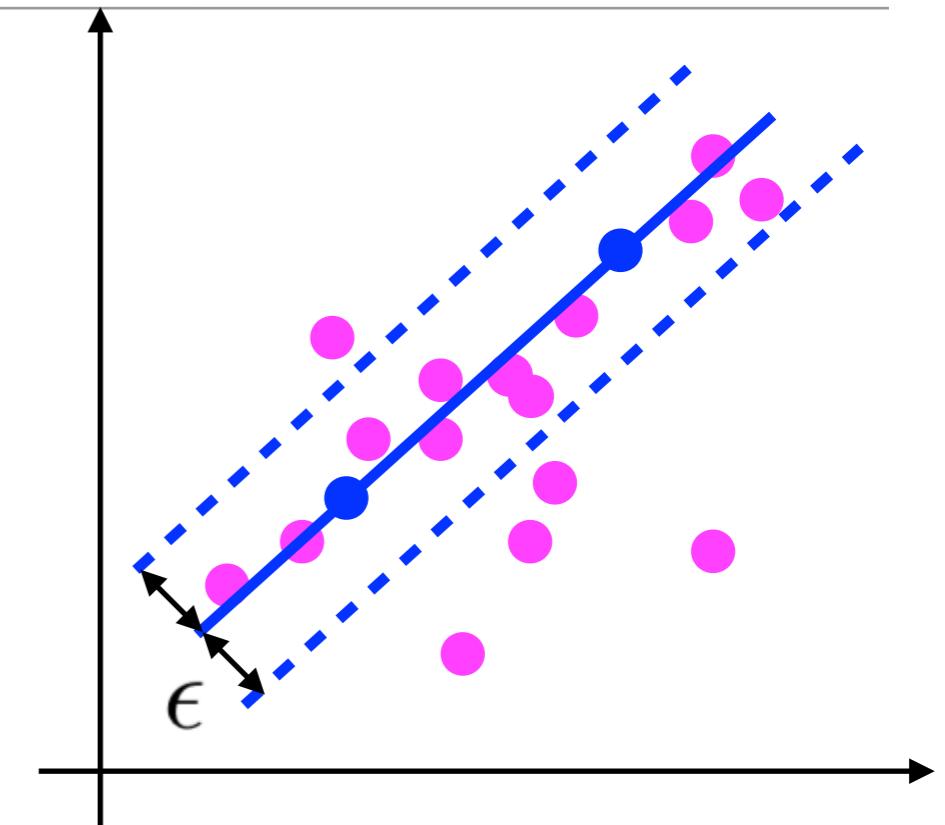
RANSAC: Parameter Tuning

1. Error Tolerance ϵ :
depends on the noise

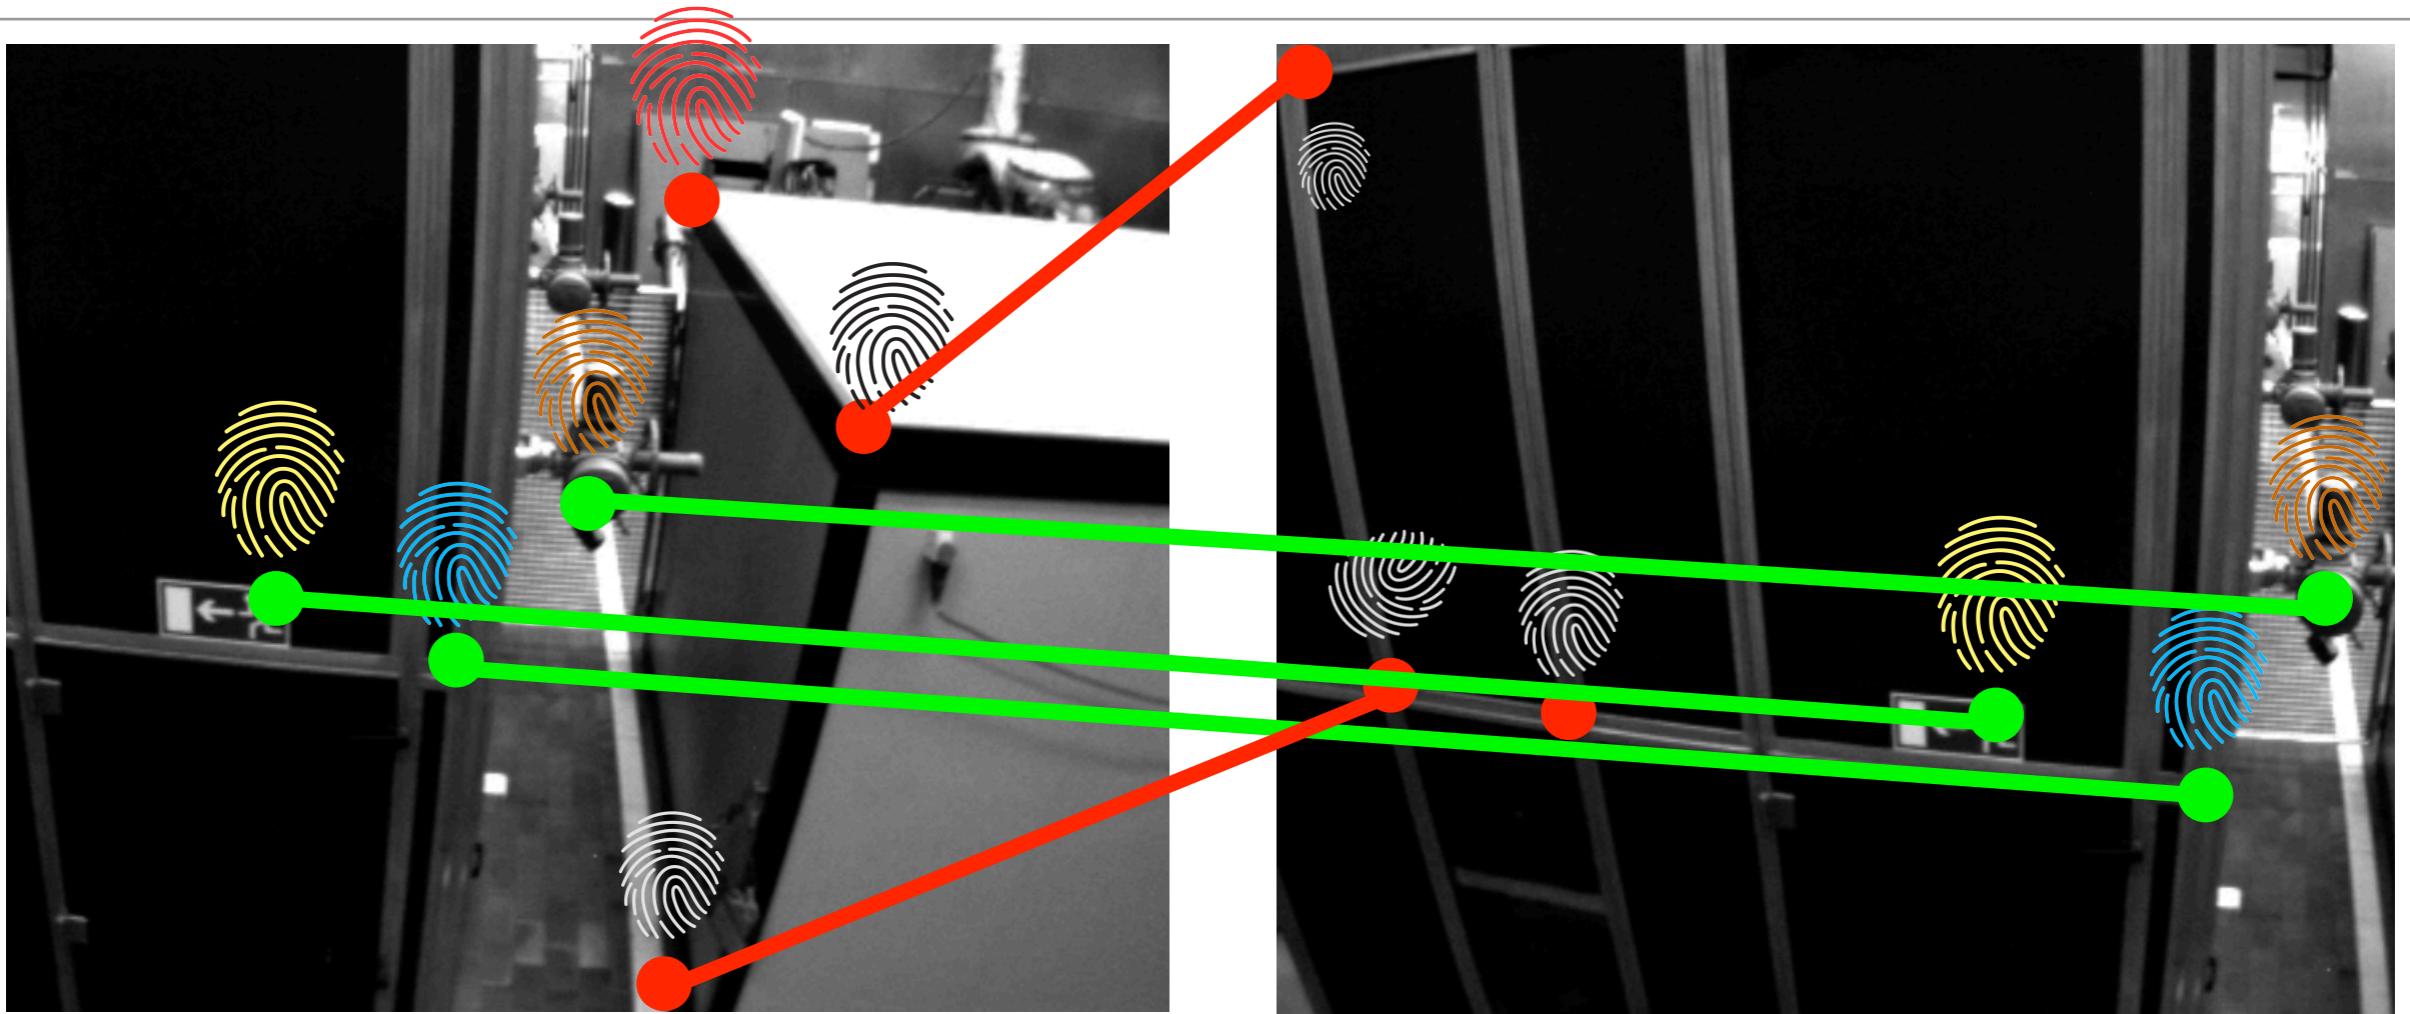
2. Acceptable consensus set:

- from the paper: $n+5$
- rule of thumb: $>50\%$ of points

3. Maximum number of iterations



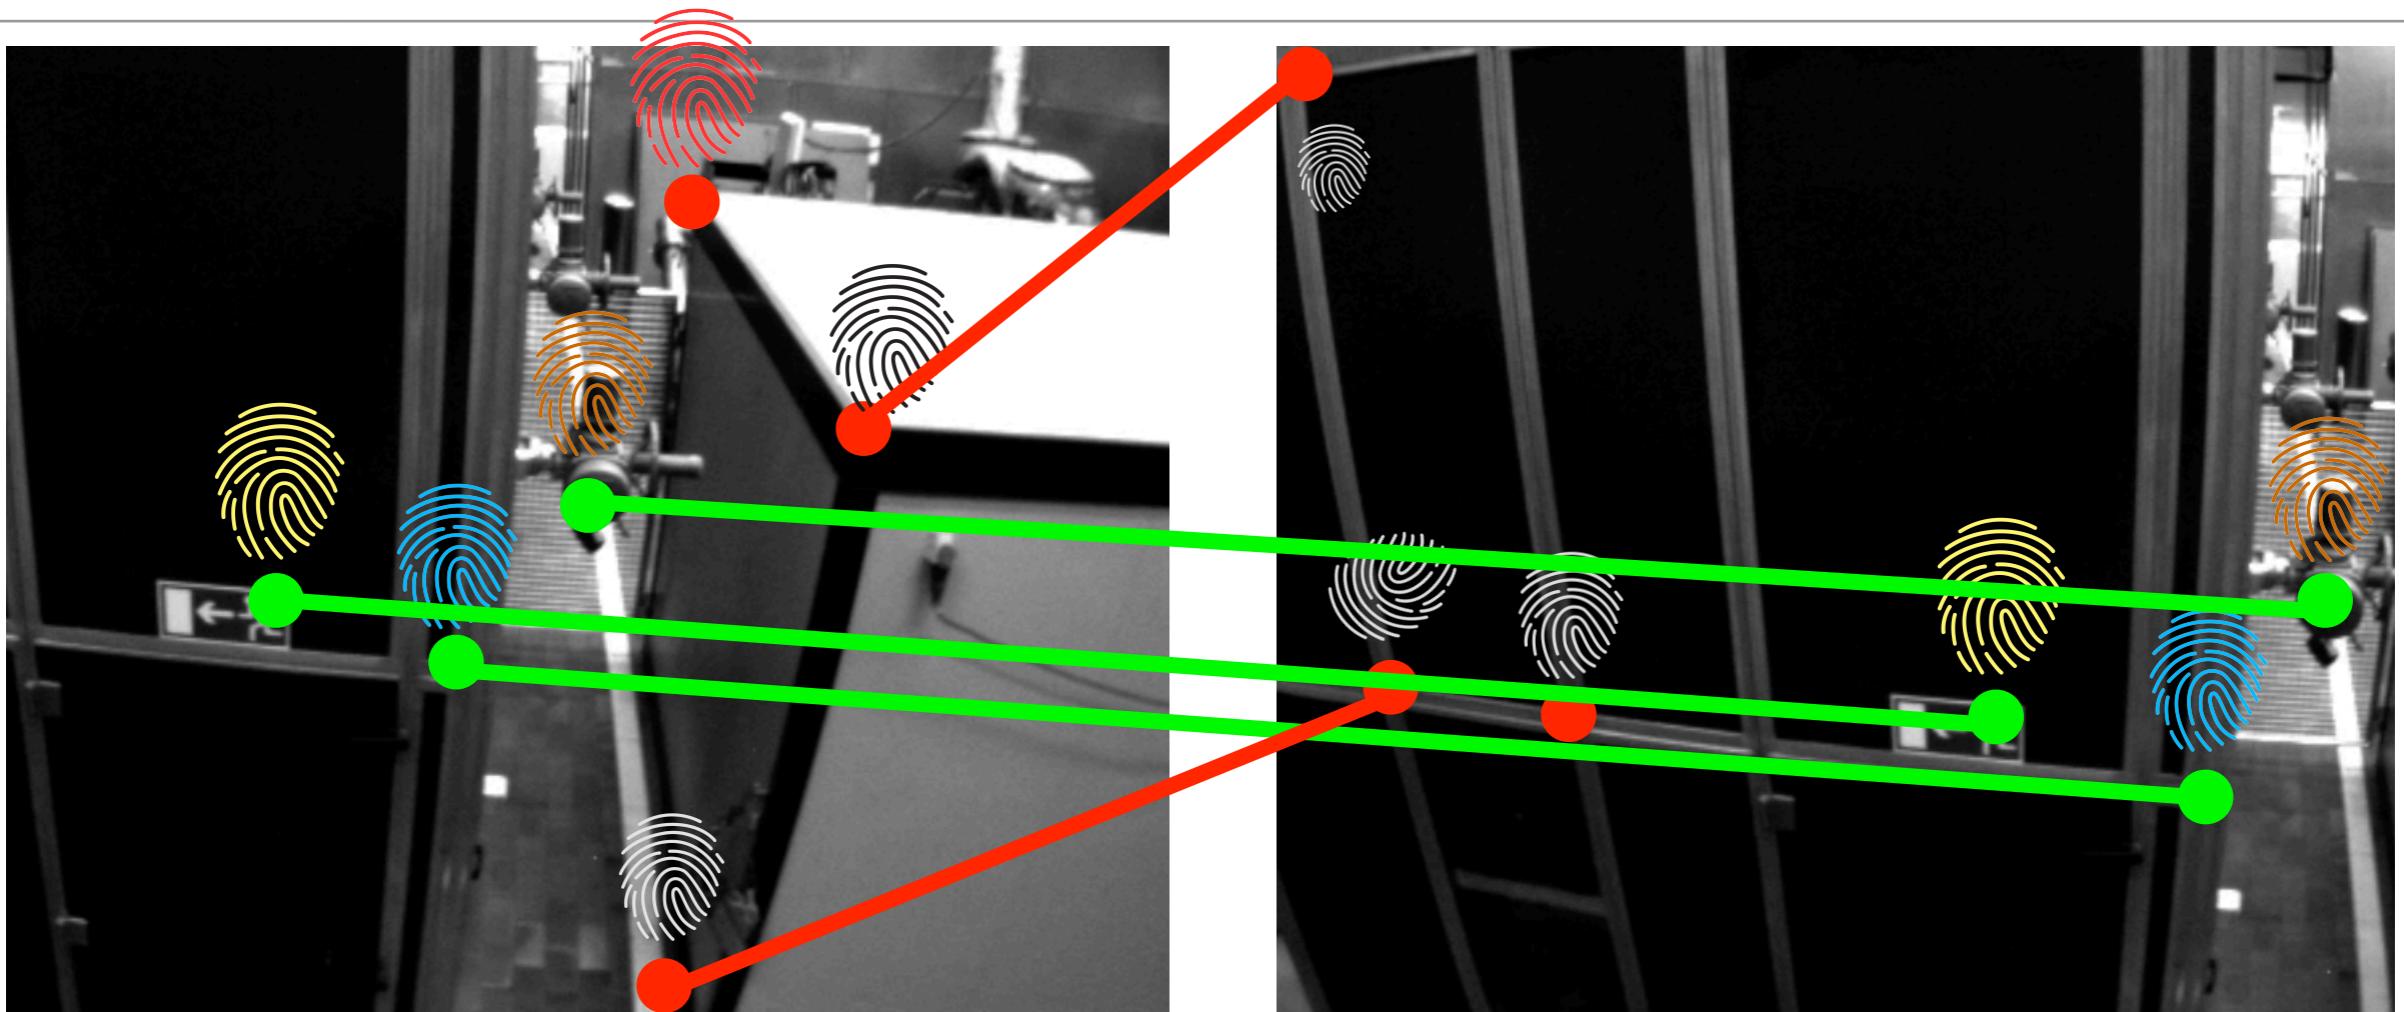
Example: RANSAC for Essential Matrix estimation



RANSAC:

1. sample n point correspondences
2. compute an estimate E' of the essential matrix E
3. count how many points are within a **tolerance** from E'
4. repeat until you get a E' that agrees with many points

Example: RANSAC for Essential Matrix estimation

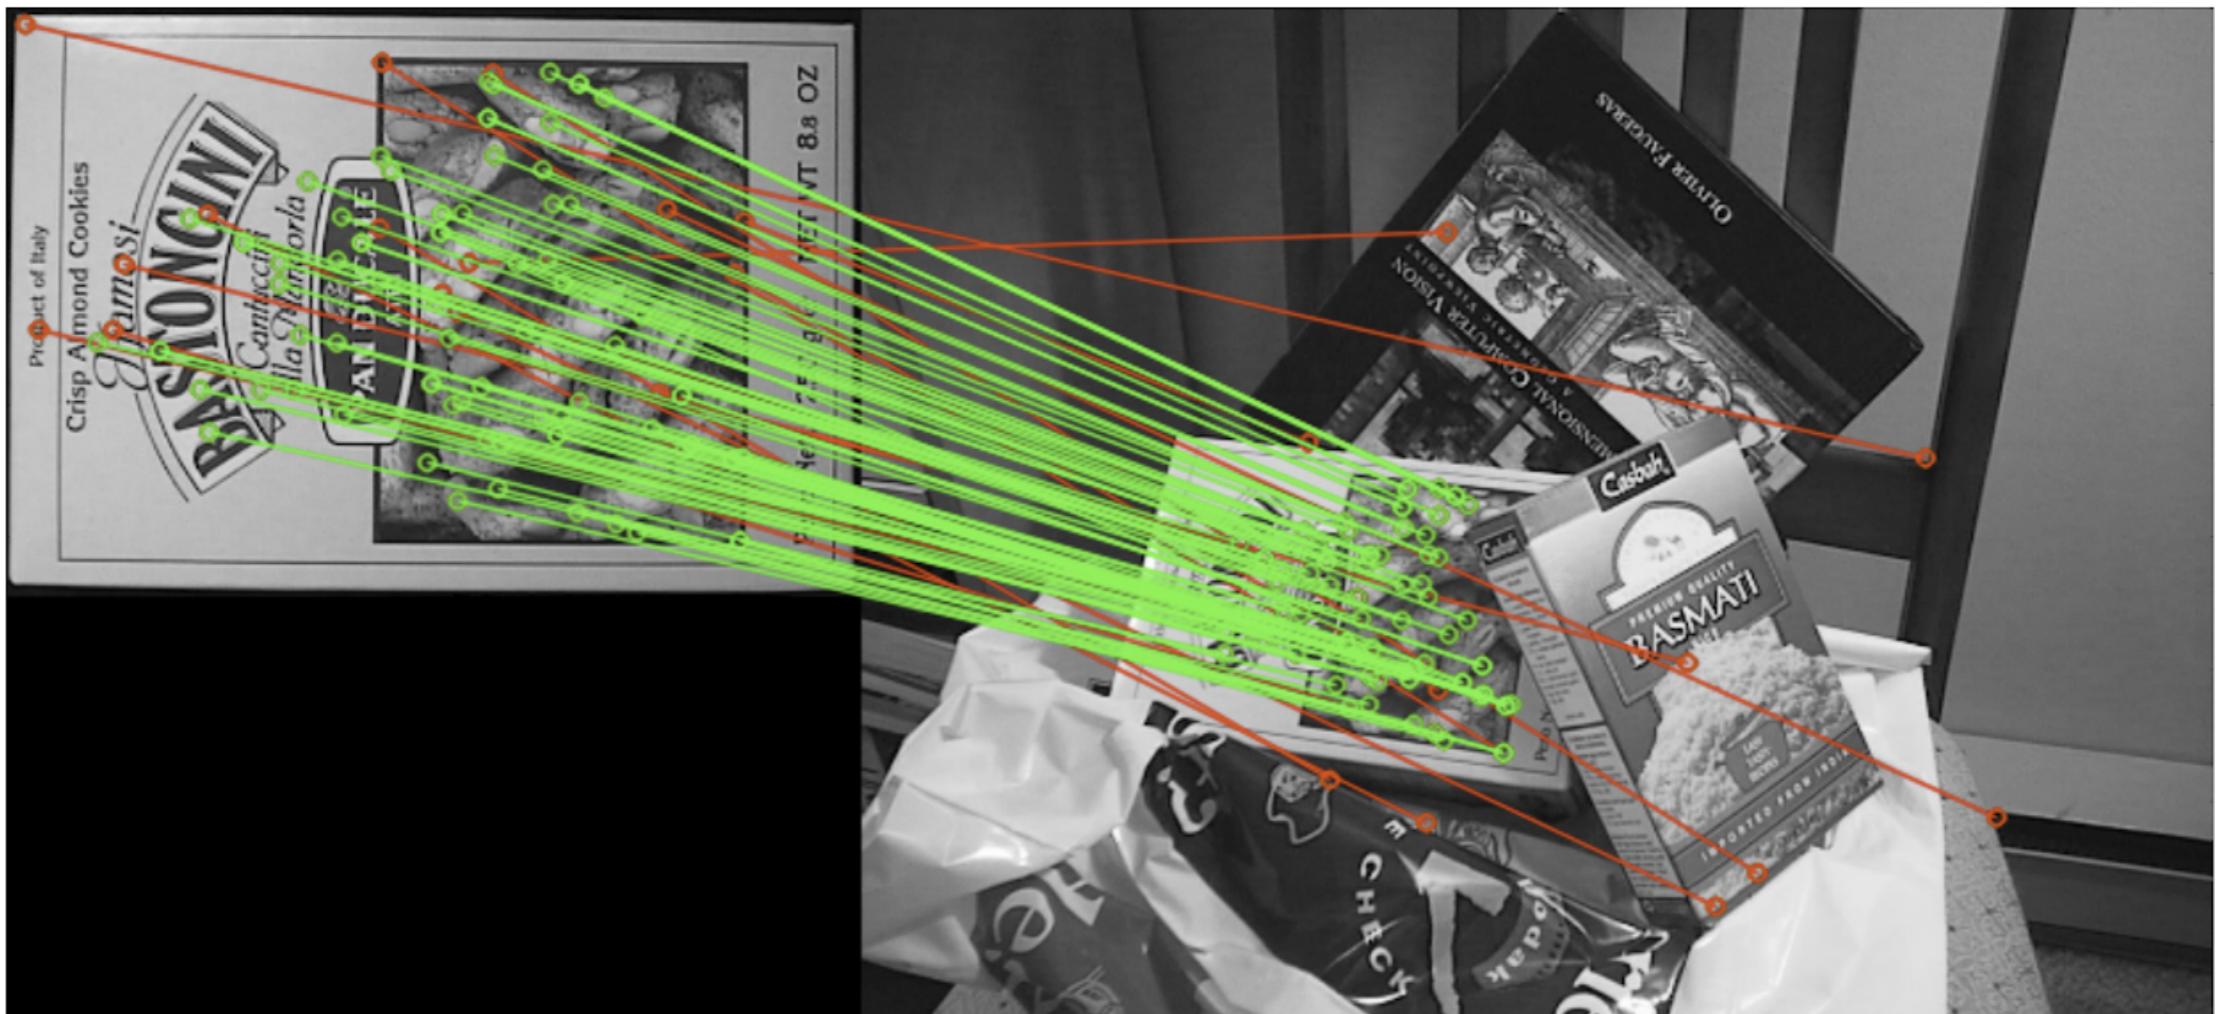


RANSAC

- essentially selects the set of inliers
- provides **geometric verification** for the correspondences

Beyond Motion Estimation

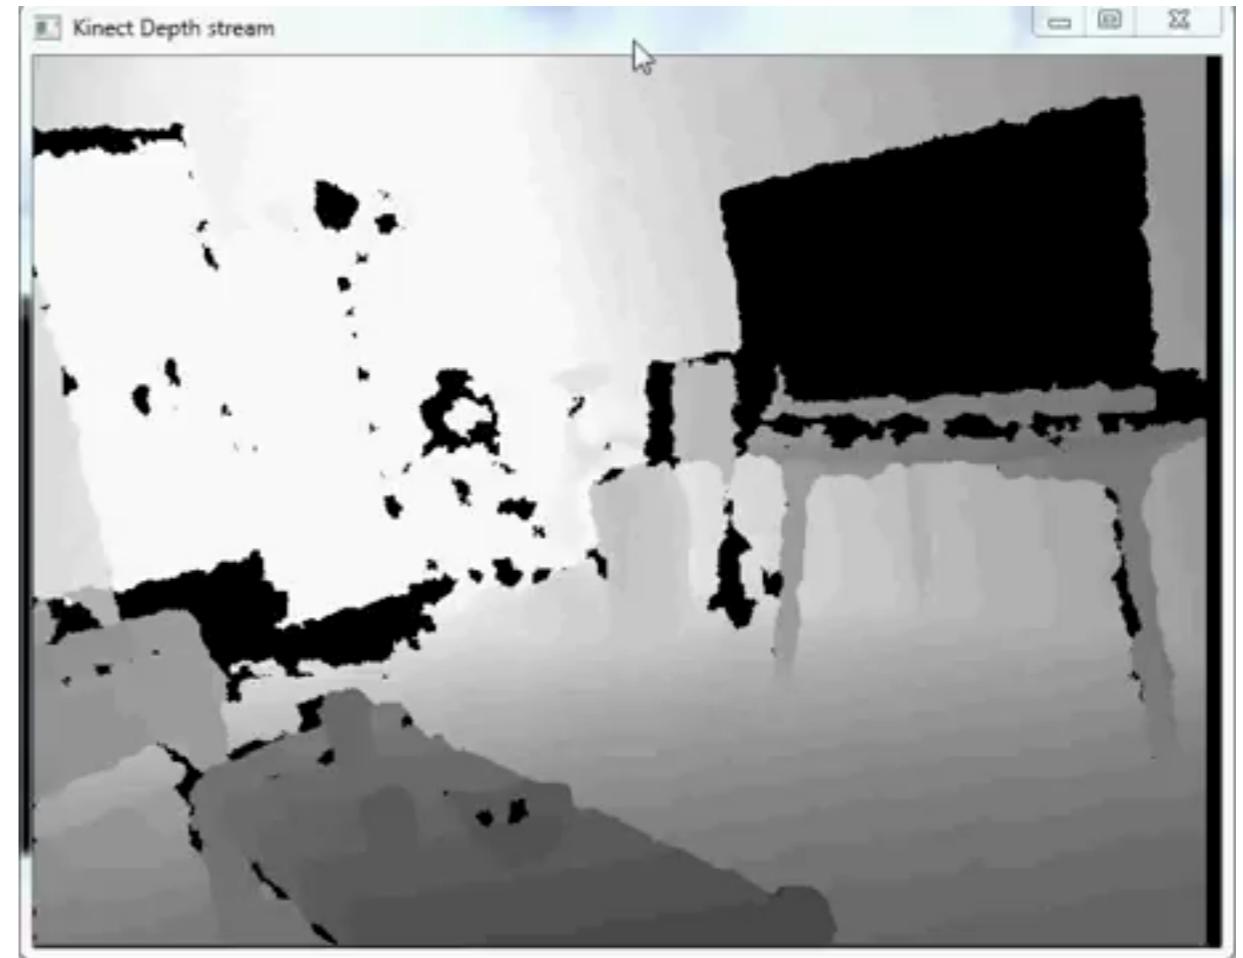
The tools we discussed (feature matching, essential matrix estimation, RANSAC) can be used also for **object detection and localization**



So far: pixel correspondences,
a.k.a., **2D-2D correspondences**

3D-3D Point Correspondences

Structured Light Cameras



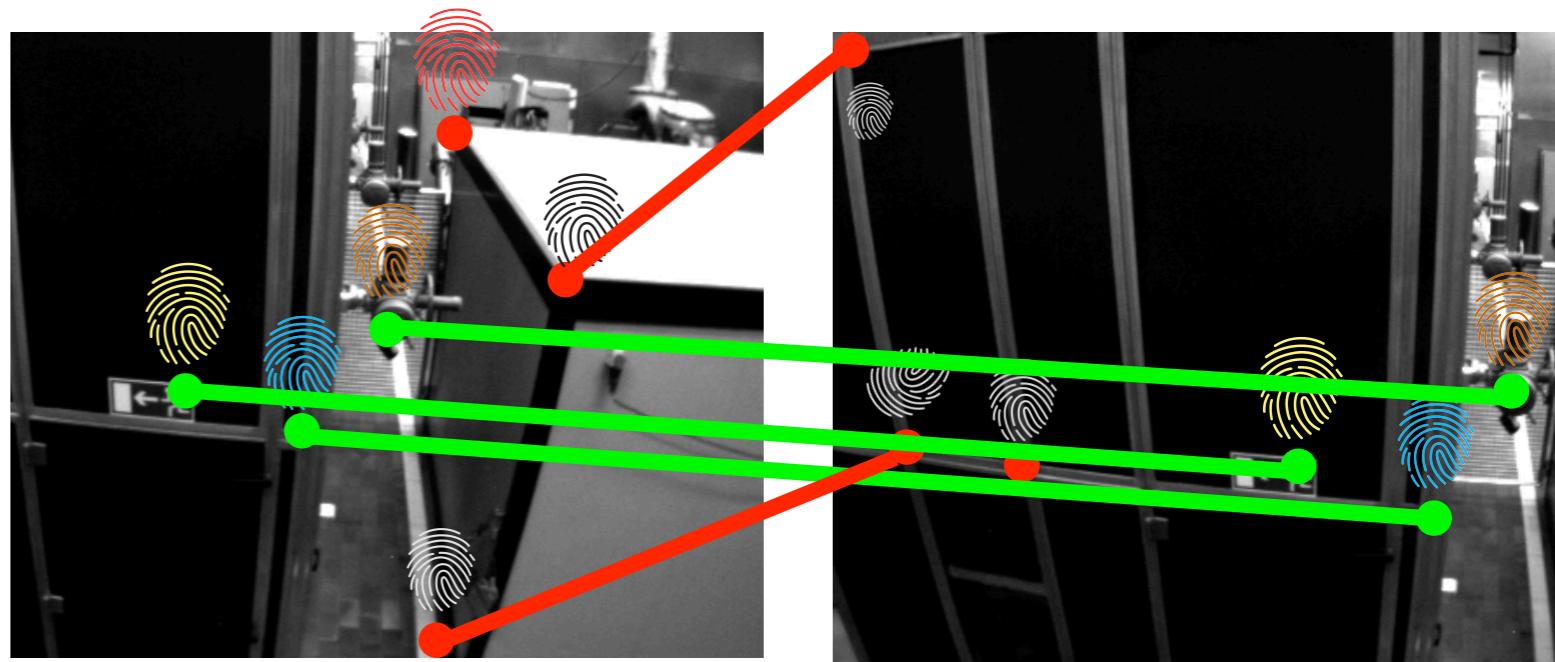
RGB-D cameras can measure depth (D) and image (RGB)

How can we use the depth information to estimate the relative pose between two RGB-D cameras observing the same scene?

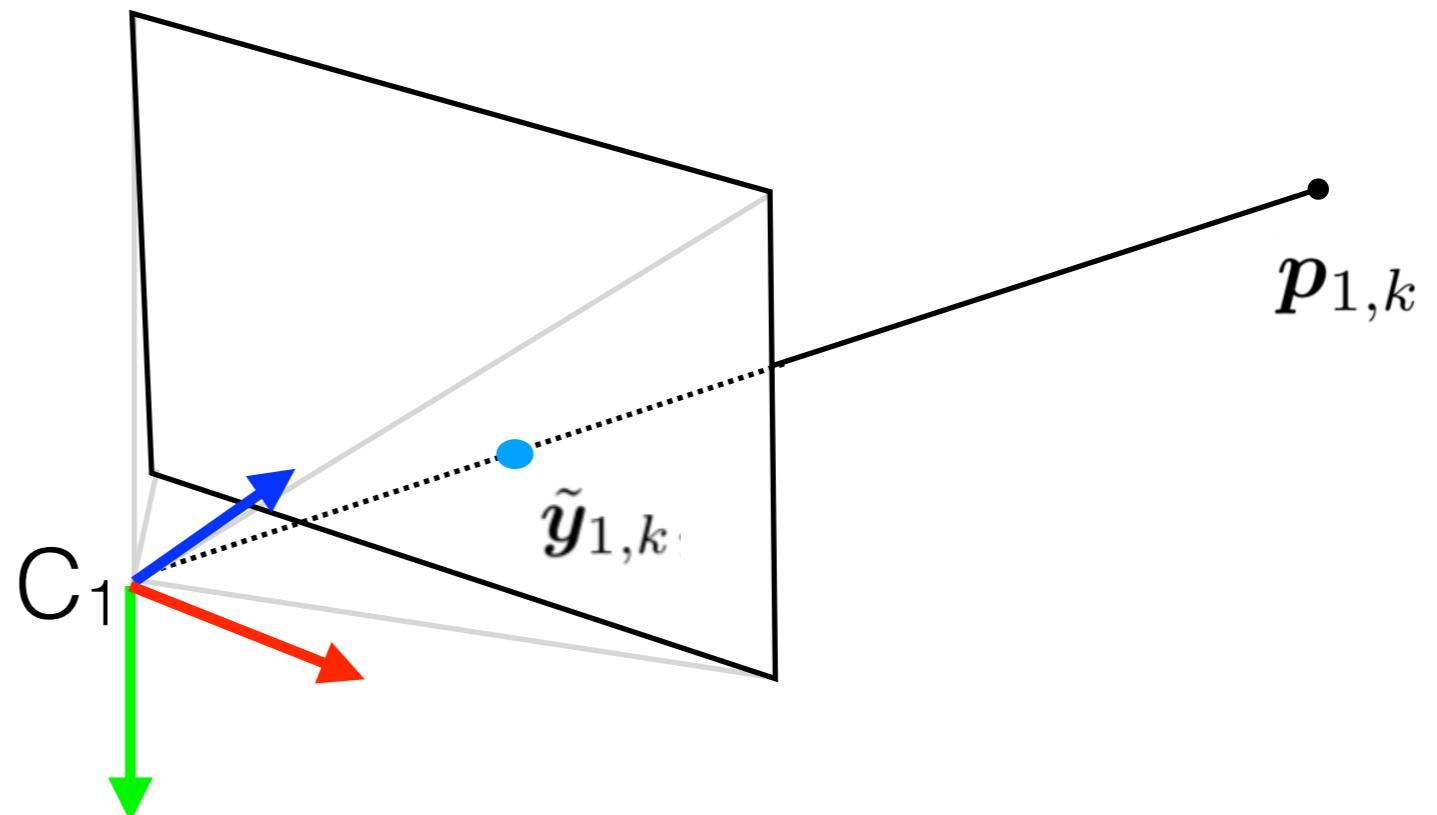
3D-3D Point Correspondences

1. We can use camera images to establish 2D-2D correspondences:

$$(\tilde{y}_{1,k}, \tilde{y}_{2,k}) \text{ for } k = 1, \dots, N$$

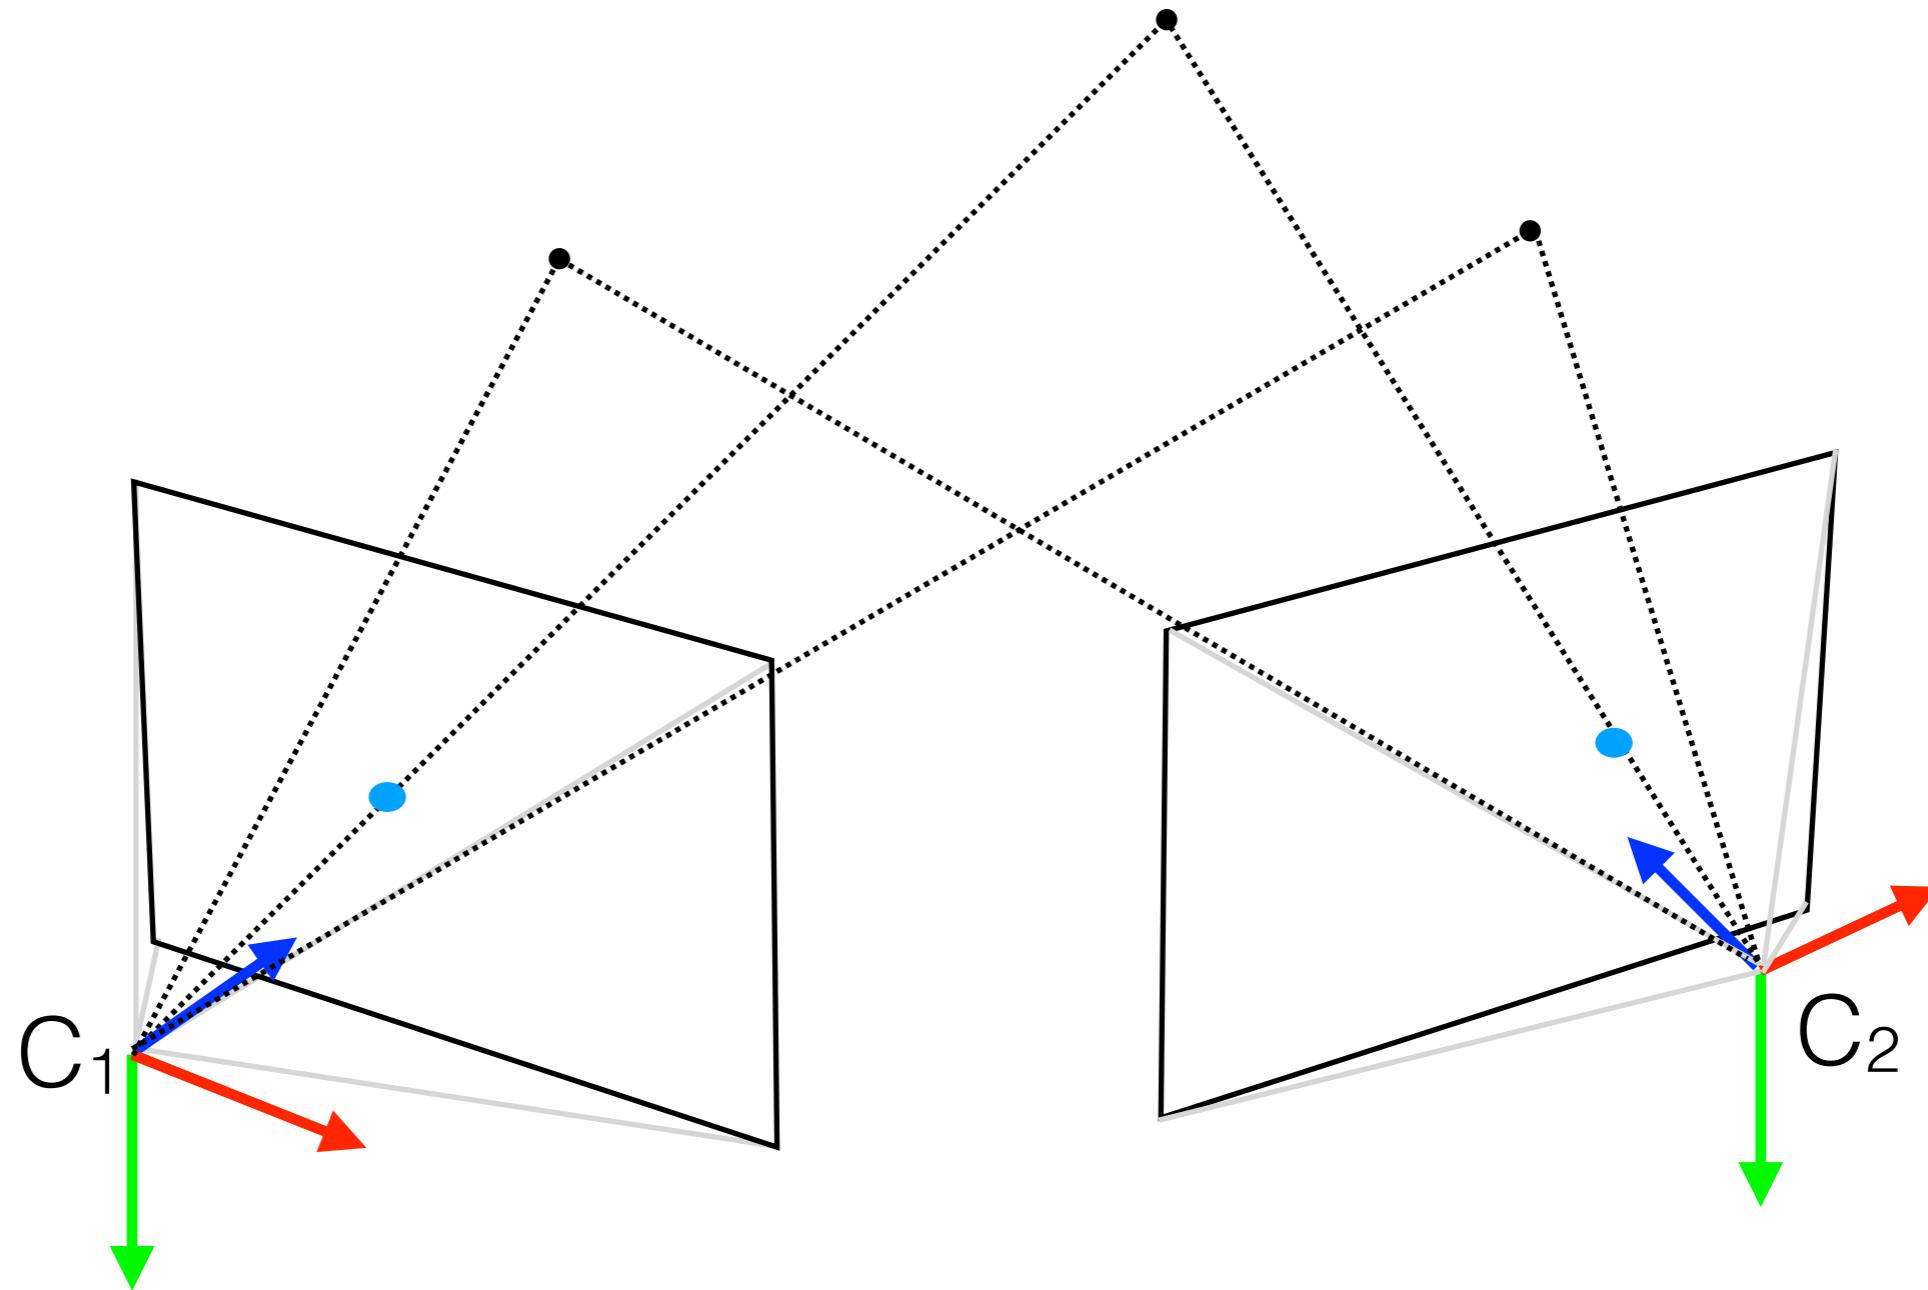


2. For each camera we can compute the set of 3D points corresponding to pixels



We obtain 3D-3D correspondences: $(p_{1,k}, p_{2,k})$ $k = 1, \dots, N$

2-view Geometry from 3D-3D Correspondences



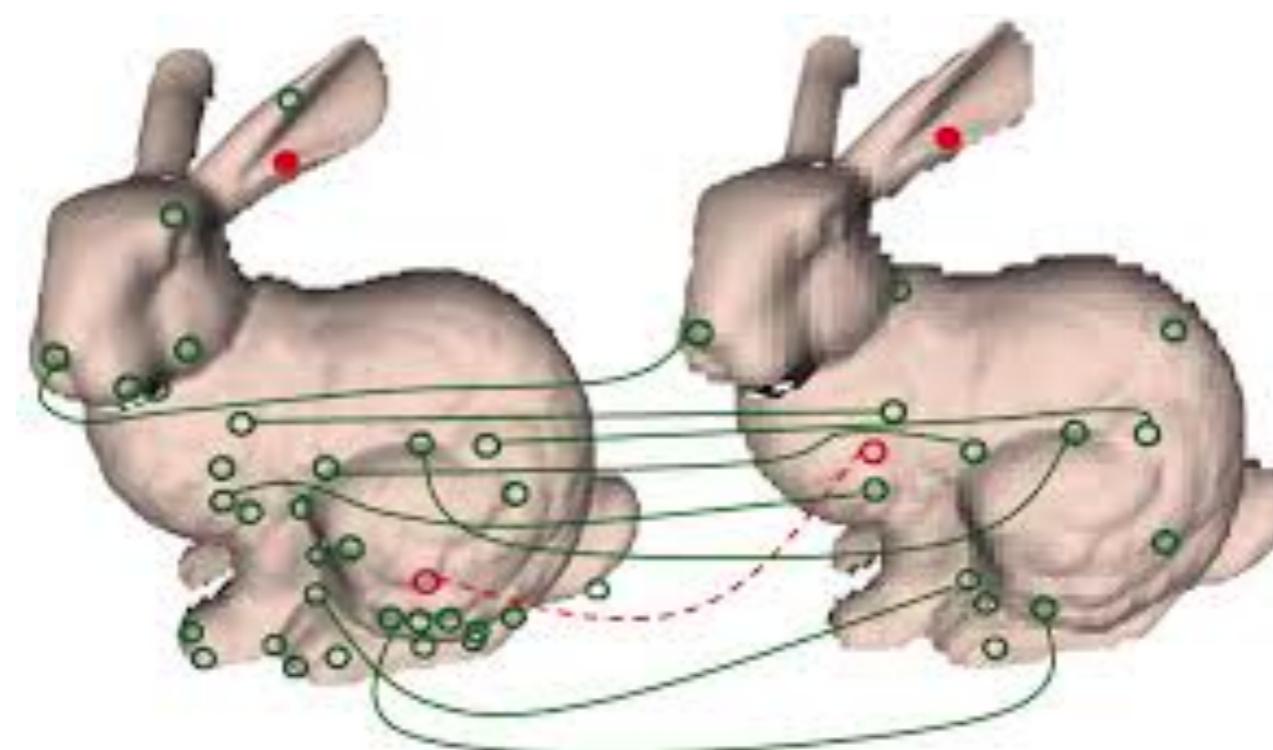
How to estimate the relative pose between the cameras from 3D-3D correspondences $(p_{1,k}, p_{2,k})$ with $k = 1, \dots, N$?

Few More Comments:

3 points are sufficient to compute the relative pose from 3D-3D correspondences

We can use the solver seen today as a 3-point minimal solver within a **RANSAC** method

Also useful for 3D objects localization:



Other names: vector registration, point cloud alignment, ..