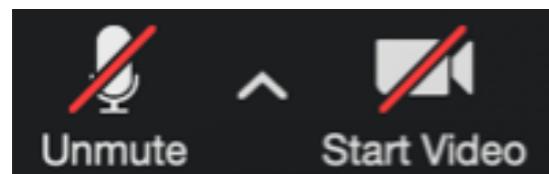




AEROSP 584 - Navigation and Guidance: From Perception to Control



Lectures start at
10:30am EST

Vasileios Tzoumas

Lecture 20

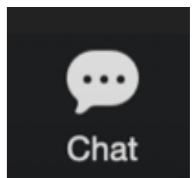


To ask questions:



[Raise Hand](#)

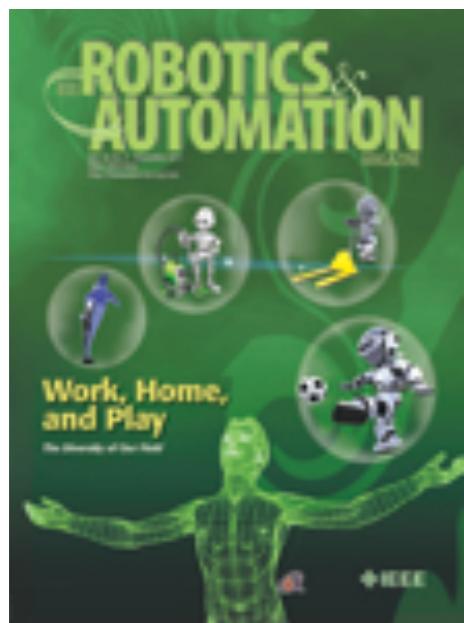
or



[Chat](#)

Today

- VO: Visual Odometry
- VIO: Visual-Inertial Odometry
- (Beyond vision)



Visual Odometry

Part I: The First 30 Years and Fundamentals

Part II: Matching, Robustness, Optimization, and Applications

By Friedrich Fraundorfer and Davide Scaramuzza

On-Manifold Preintegration for Real-Time Visual-Inertial Odometry

Christian Forster, Luca Carlone, Frank Dellaert, Davide Scaramuzza

Abstract—Current approaches for visual-inertial odometry (VIO) are able to attain highly accurate state estimation via nonlinear optimization. However, real-time optimization quickly becomes infeasible as the trajectory grows over time; this problem is further emphasized by the fact that inertial measurements come at high rate, hence leading to fast growth of the number of variables in the optimization. In this paper, we address this issue by preintegrating inertial measurements between selected keyframes into single relative motion constraints. Our first contribution is a *preintegration theory* that properly addresses the manifold structure of the rotation group. We formally discuss the generative measurement model as well as the nature of the rotation noise and derive the expression for the *maximum a posteriori* state estimator. Our theoretical development enables the computation of all necessary Jacobians for the optimization and a-posteriori bias correction in analytic form. The second contribution is to show that the preintegrated IMU model can be seamlessly integrated into a visual-inertial pipeline under the unifying framework of factor graphs. This enables the application of incremental-smoothing algorithms and the use of a *structureless* model for visual measurements, which avoids optimizing over the 3D points, further accelerating the computation. We perform an extensive evaluation of our monocular VIO pipeline on real and simulated datasets. The results confirm that our modelling effort leads to accurate state estimation in real-time, outperforming state-of-the-art approaches.

of monocular vision and gravity observable [1] and provides robust and accurate inter-frame motion estimates. Applications of VIO range from autonomous navigation in GPS-denied environments, to 3D reconstruction, and augmented reality.

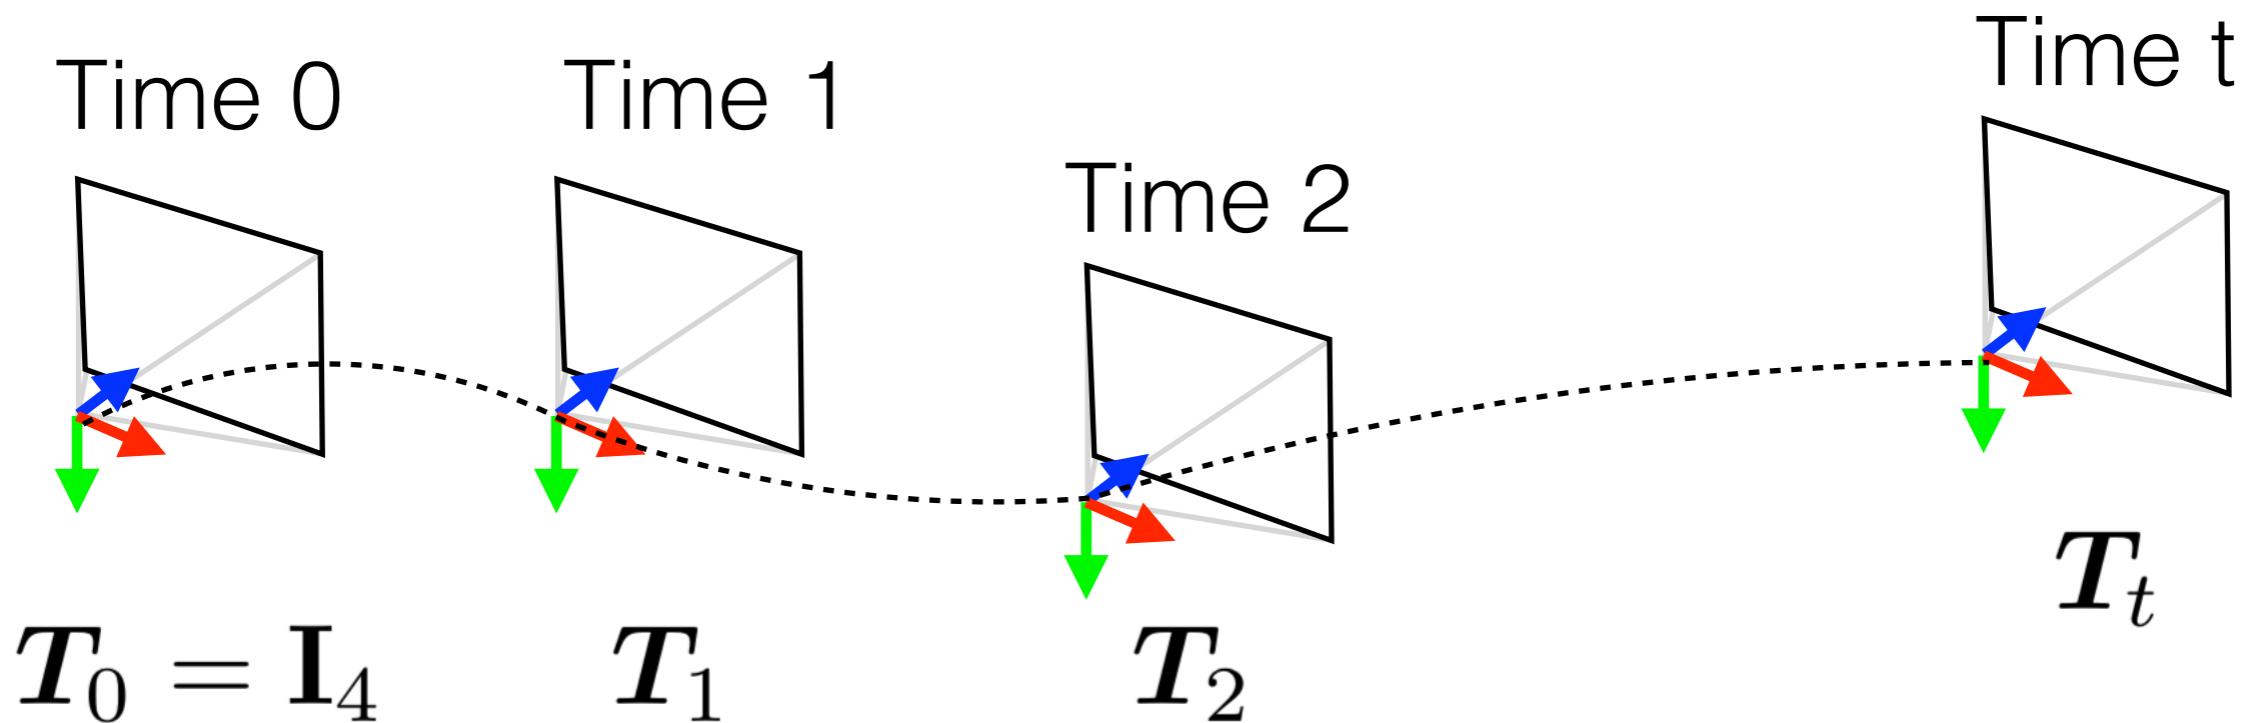
The existing literature on VIO imposes a trade-off between accuracy and computational efficiency (a detailed review is given in Section II). On the one hand, filtering approaches enable fast inference, but their accuracy is deteriorated by the accumulation of linearization errors. On the other hand, full smoothing approaches, based on nonlinear optimization, are accurate, but computationally demanding. Fixed-lag smoothing offers a compromise between accuracy for efficiency; however, it is not clear how to set the length of the estimation window so to guarantee a given level of performance.

In this work we show that it is possible to overcome this trade-off. We design a VIO system that enables fast incremental smoothing and computes the optimal *maximum a posteriori* (MAP) estimate in real time. An overview of our approach is given in Section IV.

The first step towards this goal is the development of a novel preintegration theory. The use of *preintegrated IMU measurements* was first proposed in [2] and consists of combining

Visual Odometry

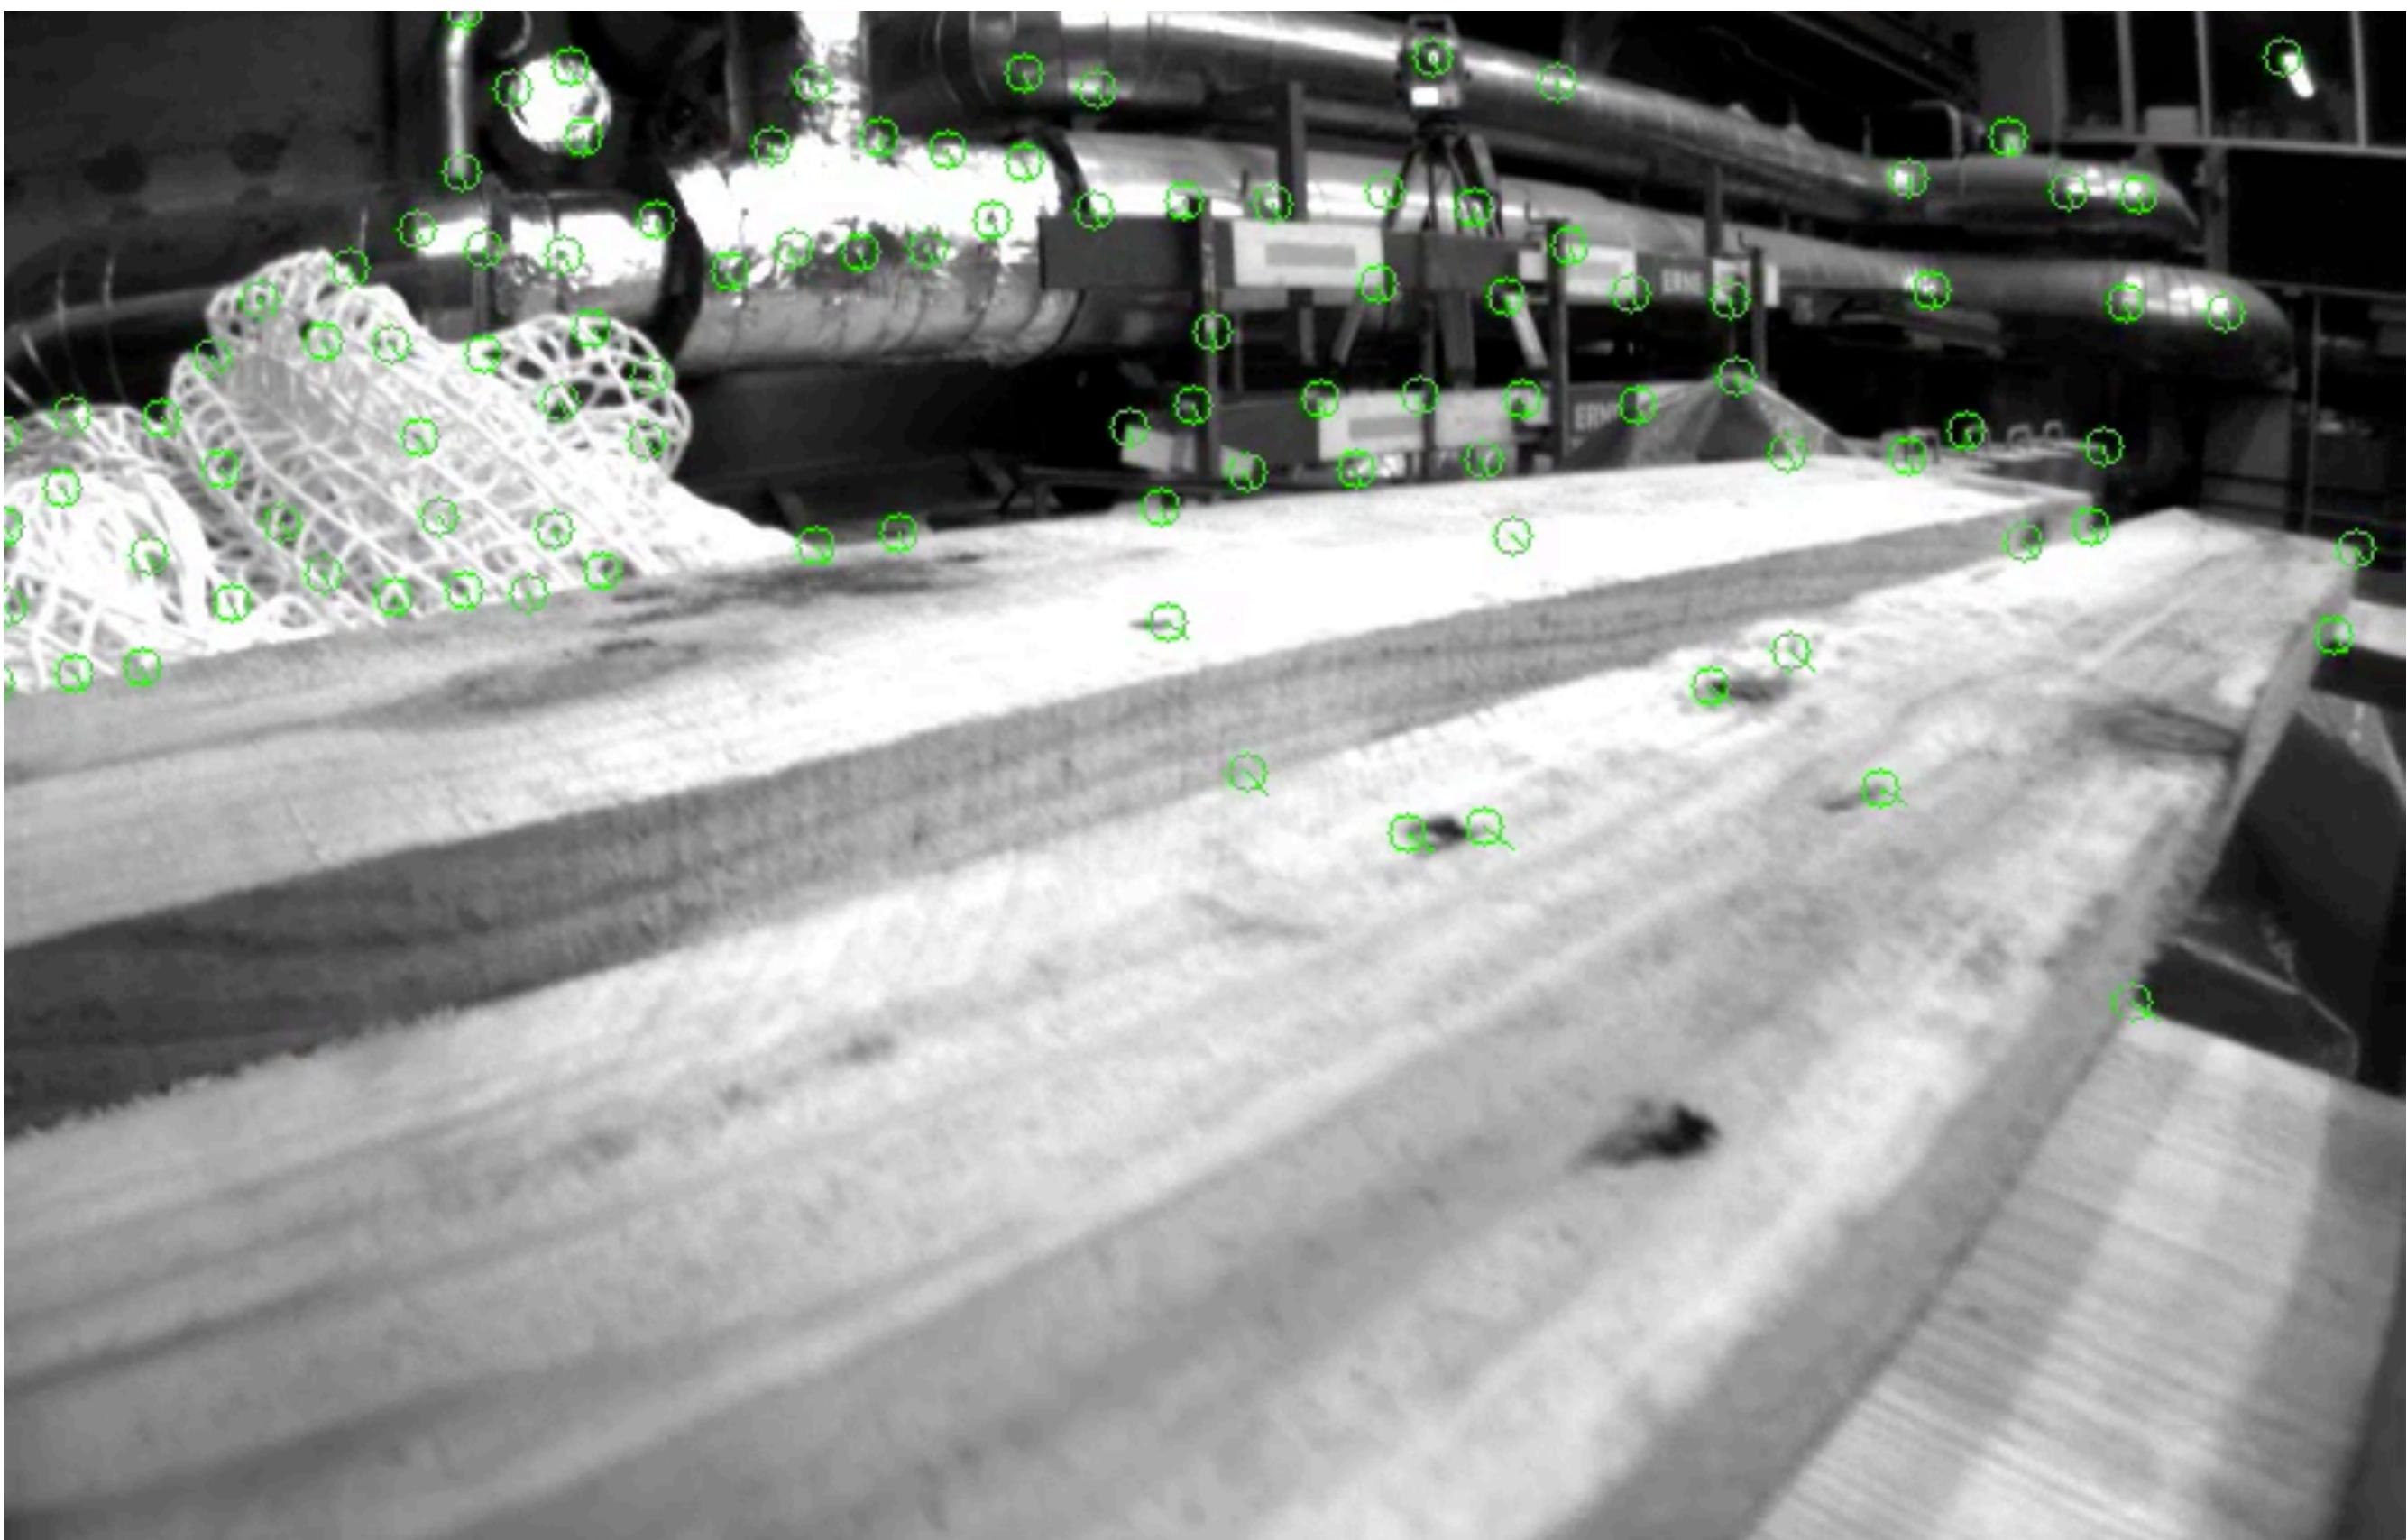
odometry: incremental motion estimation



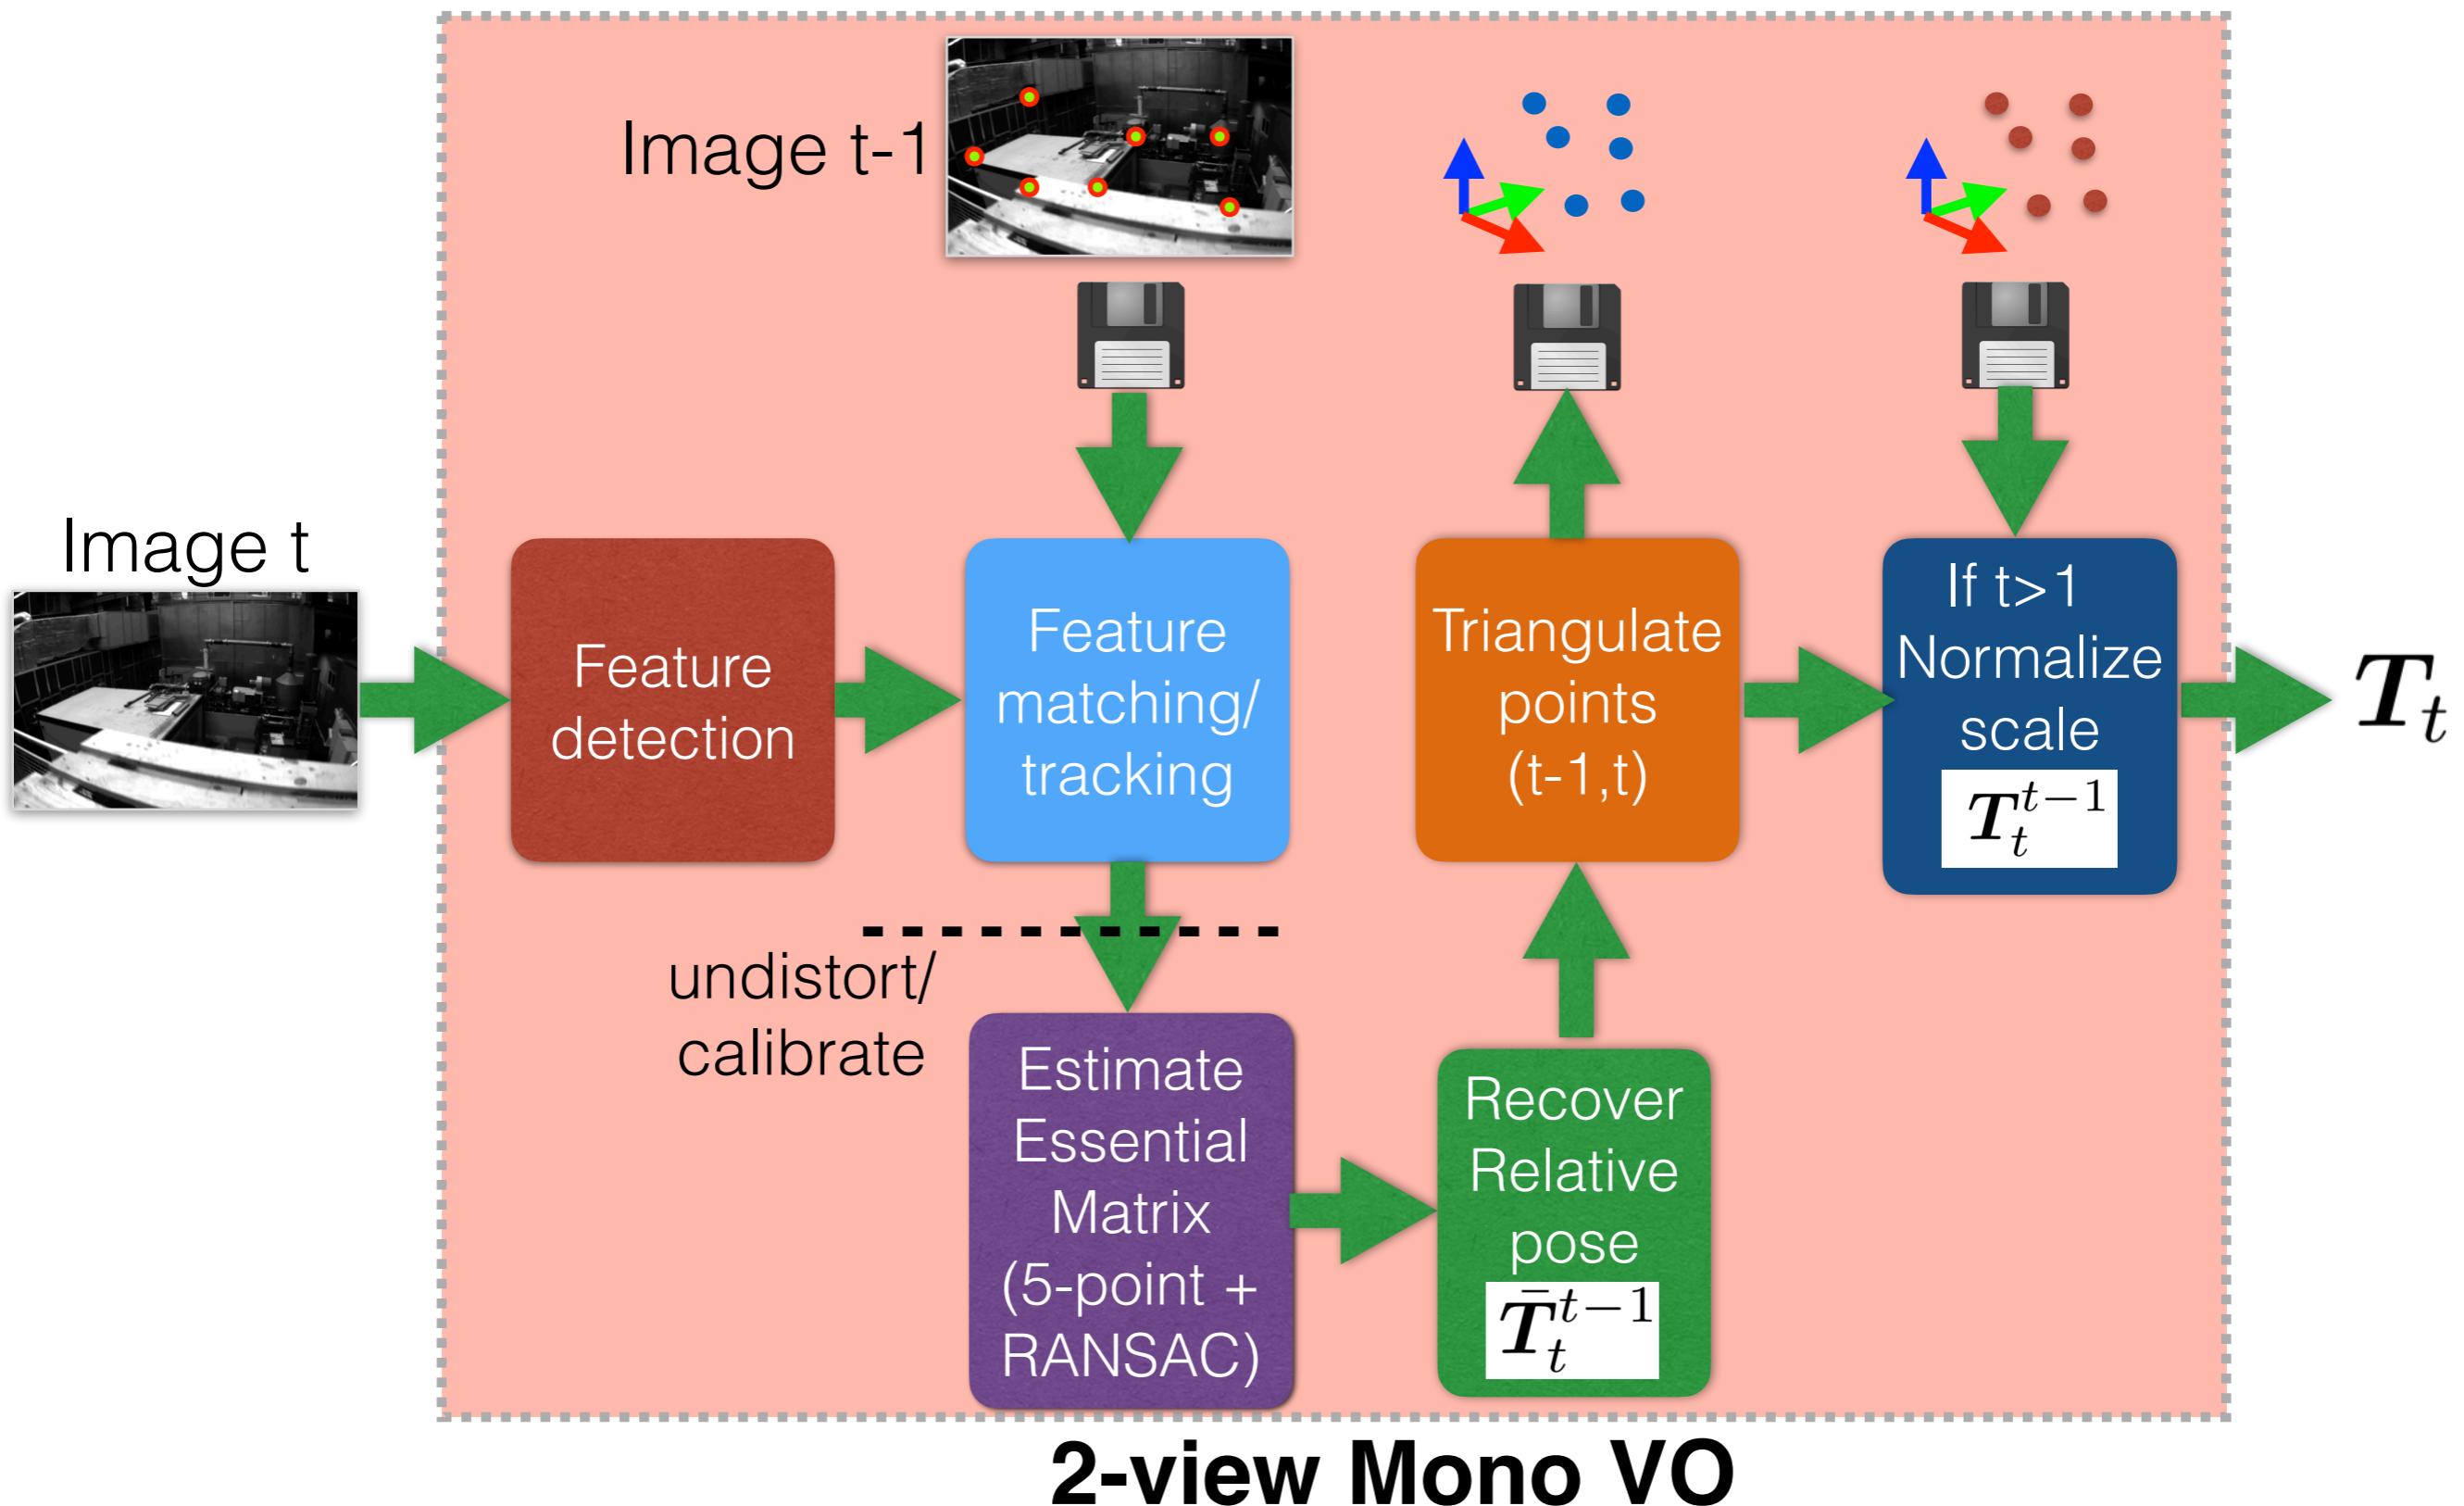
Visual odometry (VO): motion estimation based on cameras (monocular, stereo, RGB-D, ...)

others: wheel odometry, inertial, visual-inertial

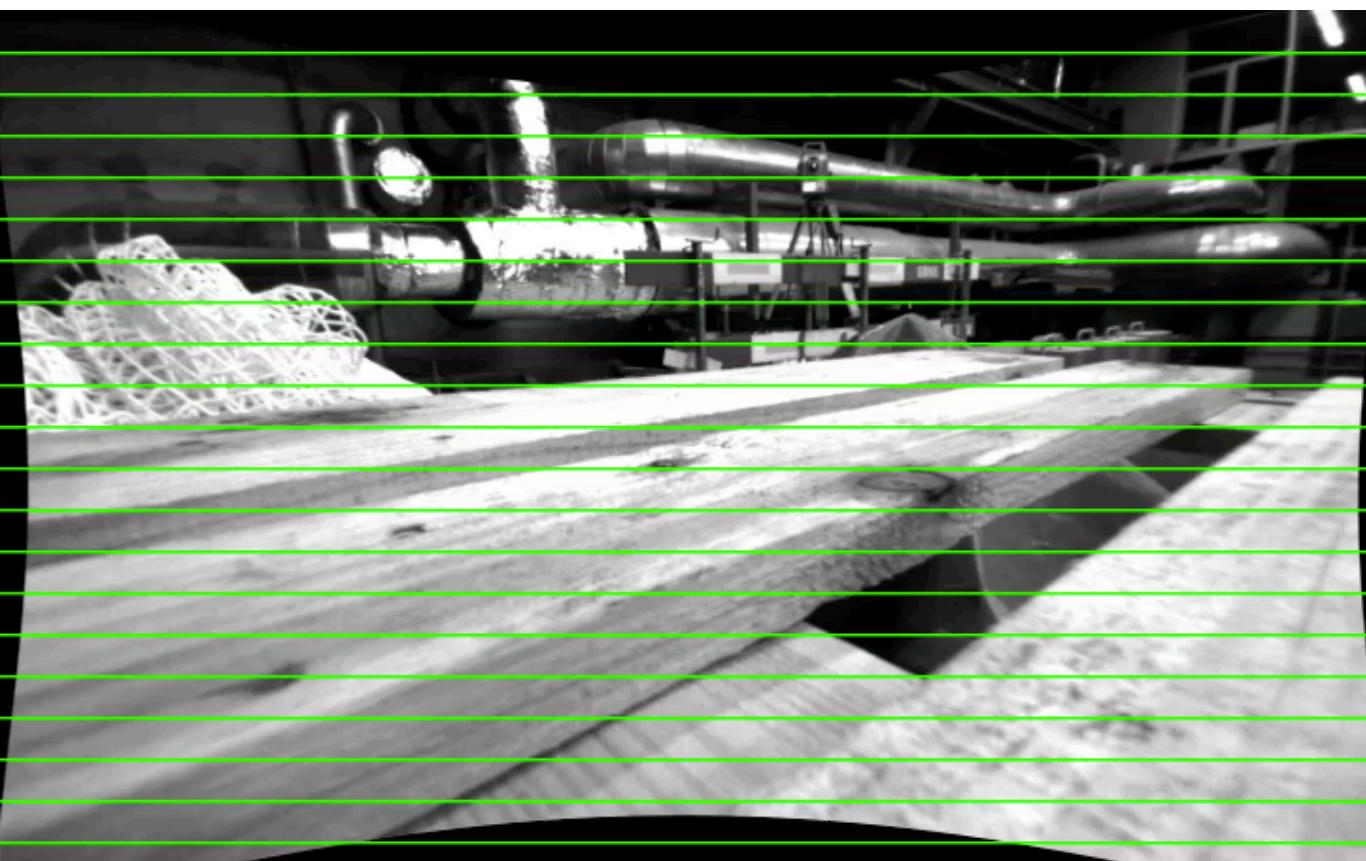
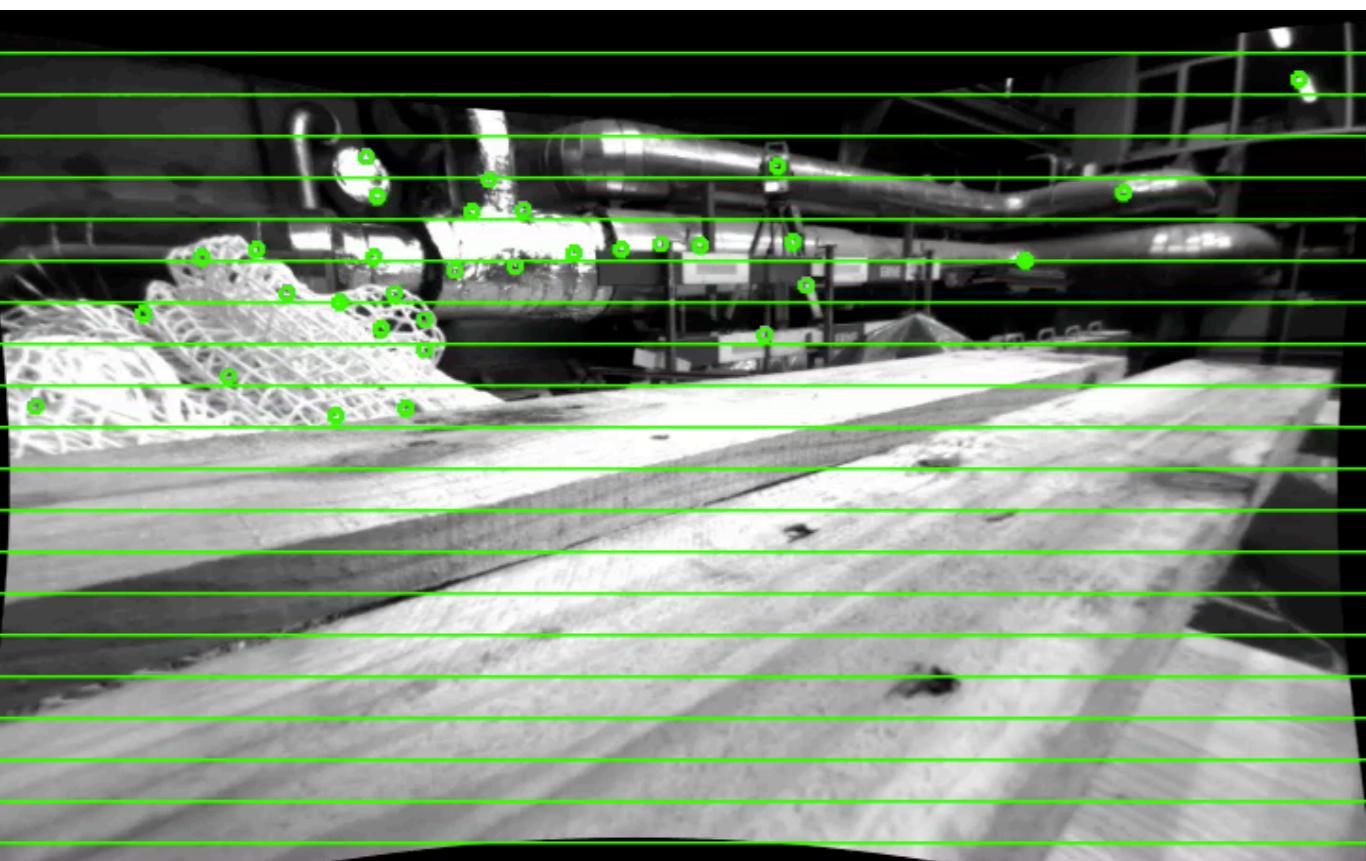
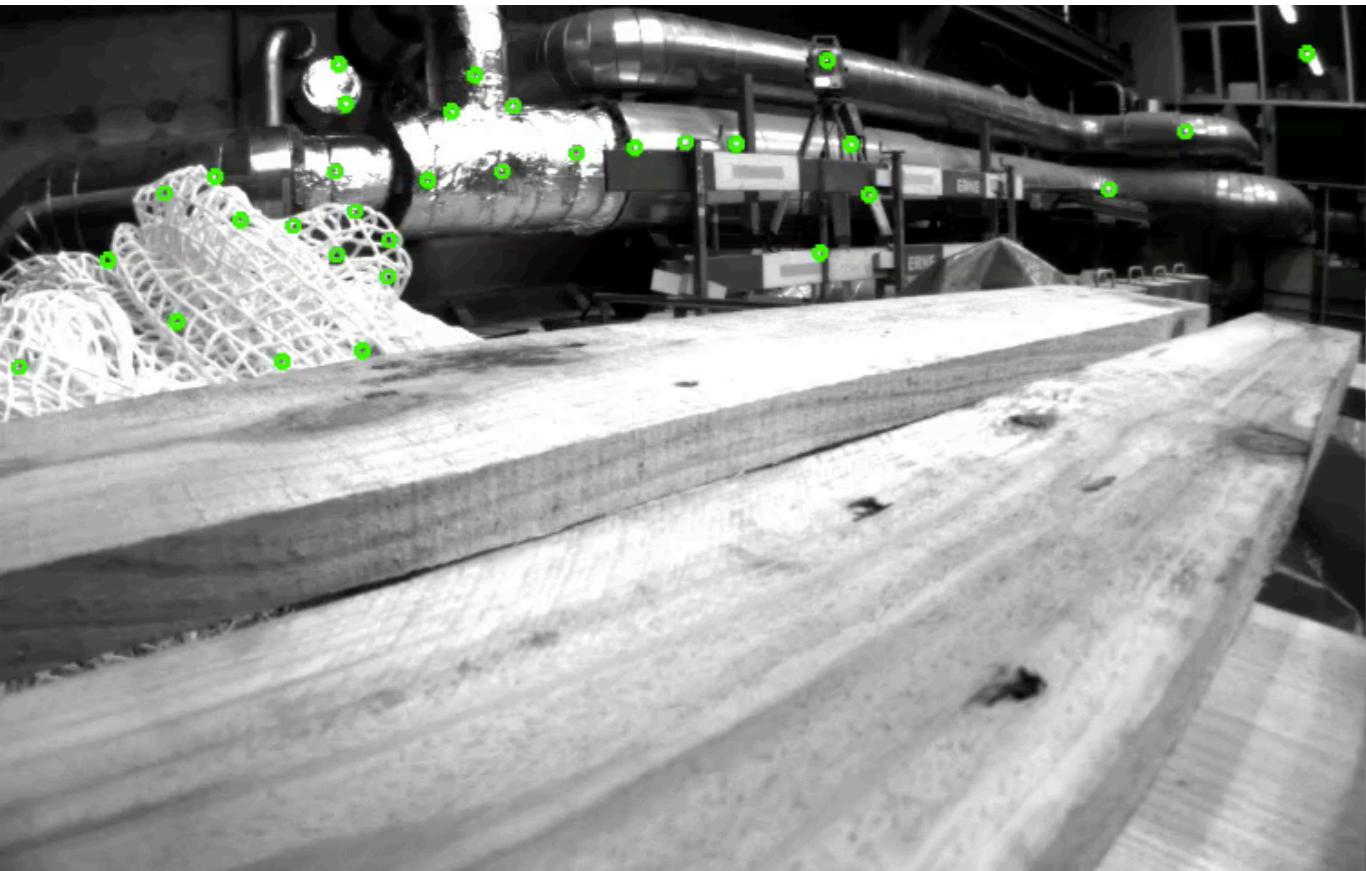
Feature Tracking



Monocular VO with 2D-2D Correspondences

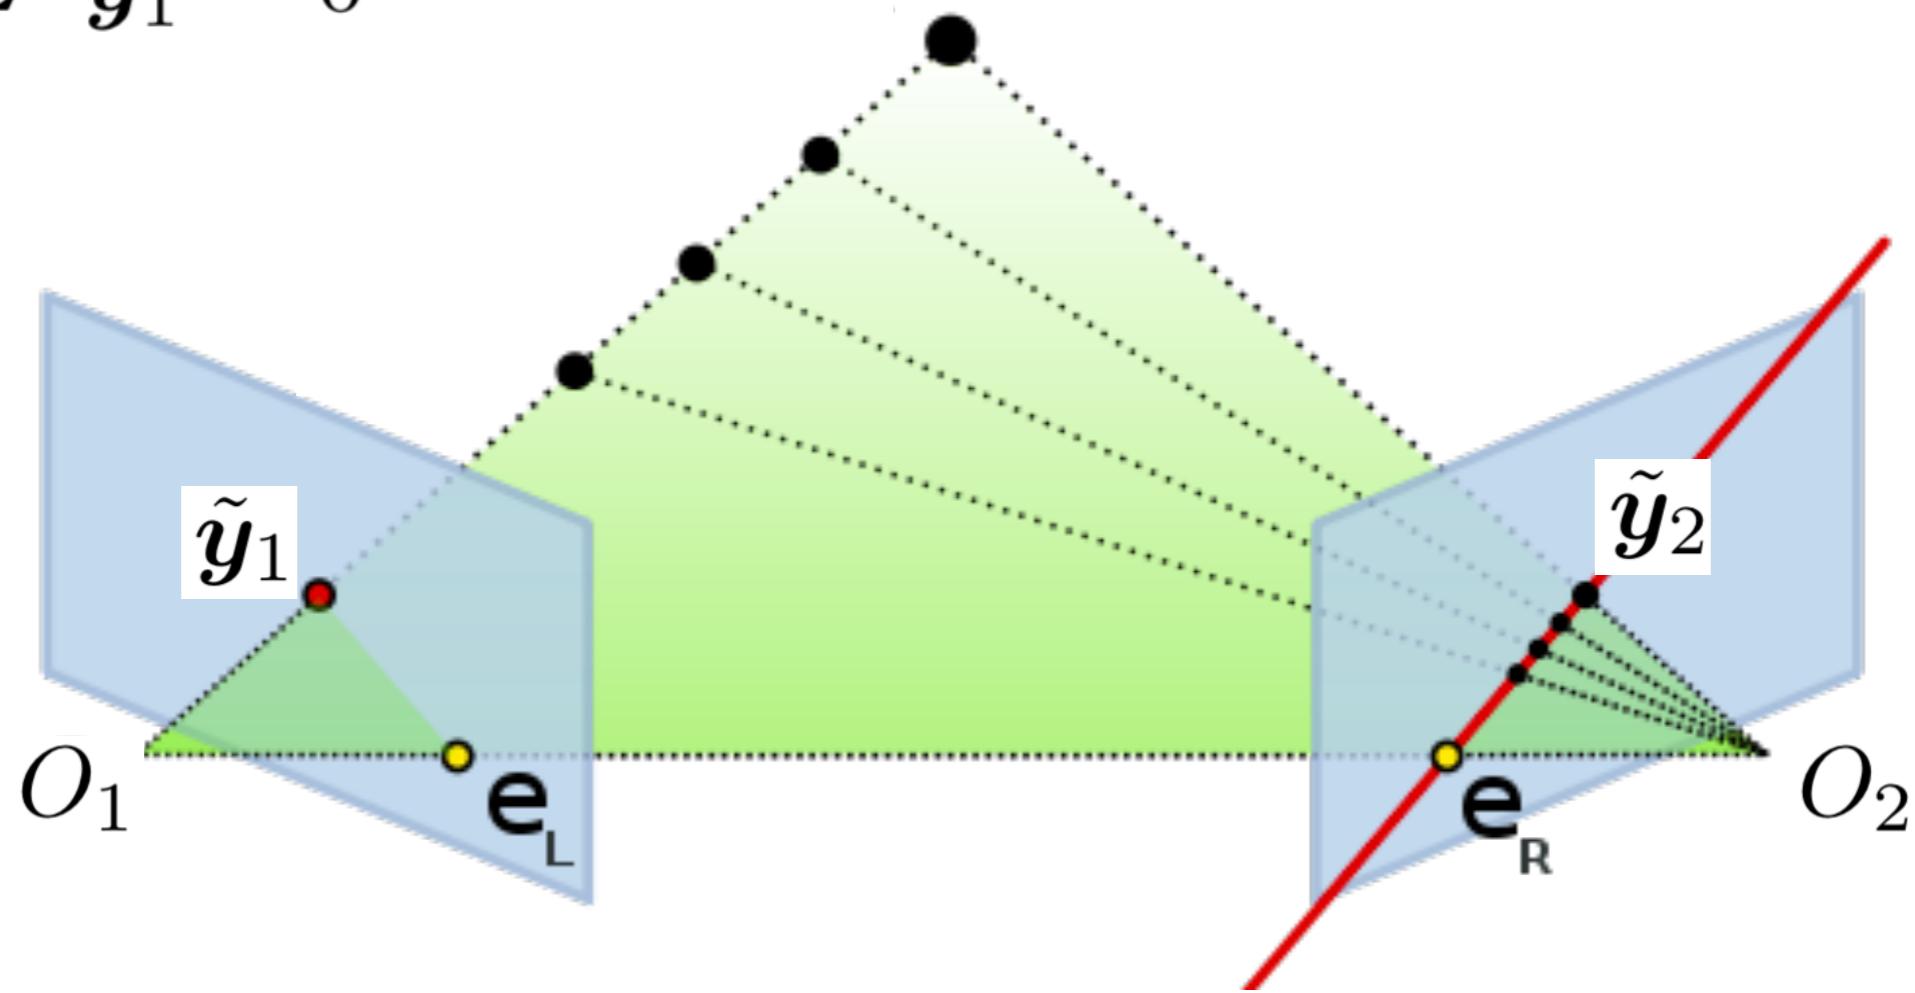


Stereo Matching



Epipolar Geometry

$$\tilde{y}_2^\top E \tilde{y}_1 = 0$$



■ epipolar plane

✓ epipolar line

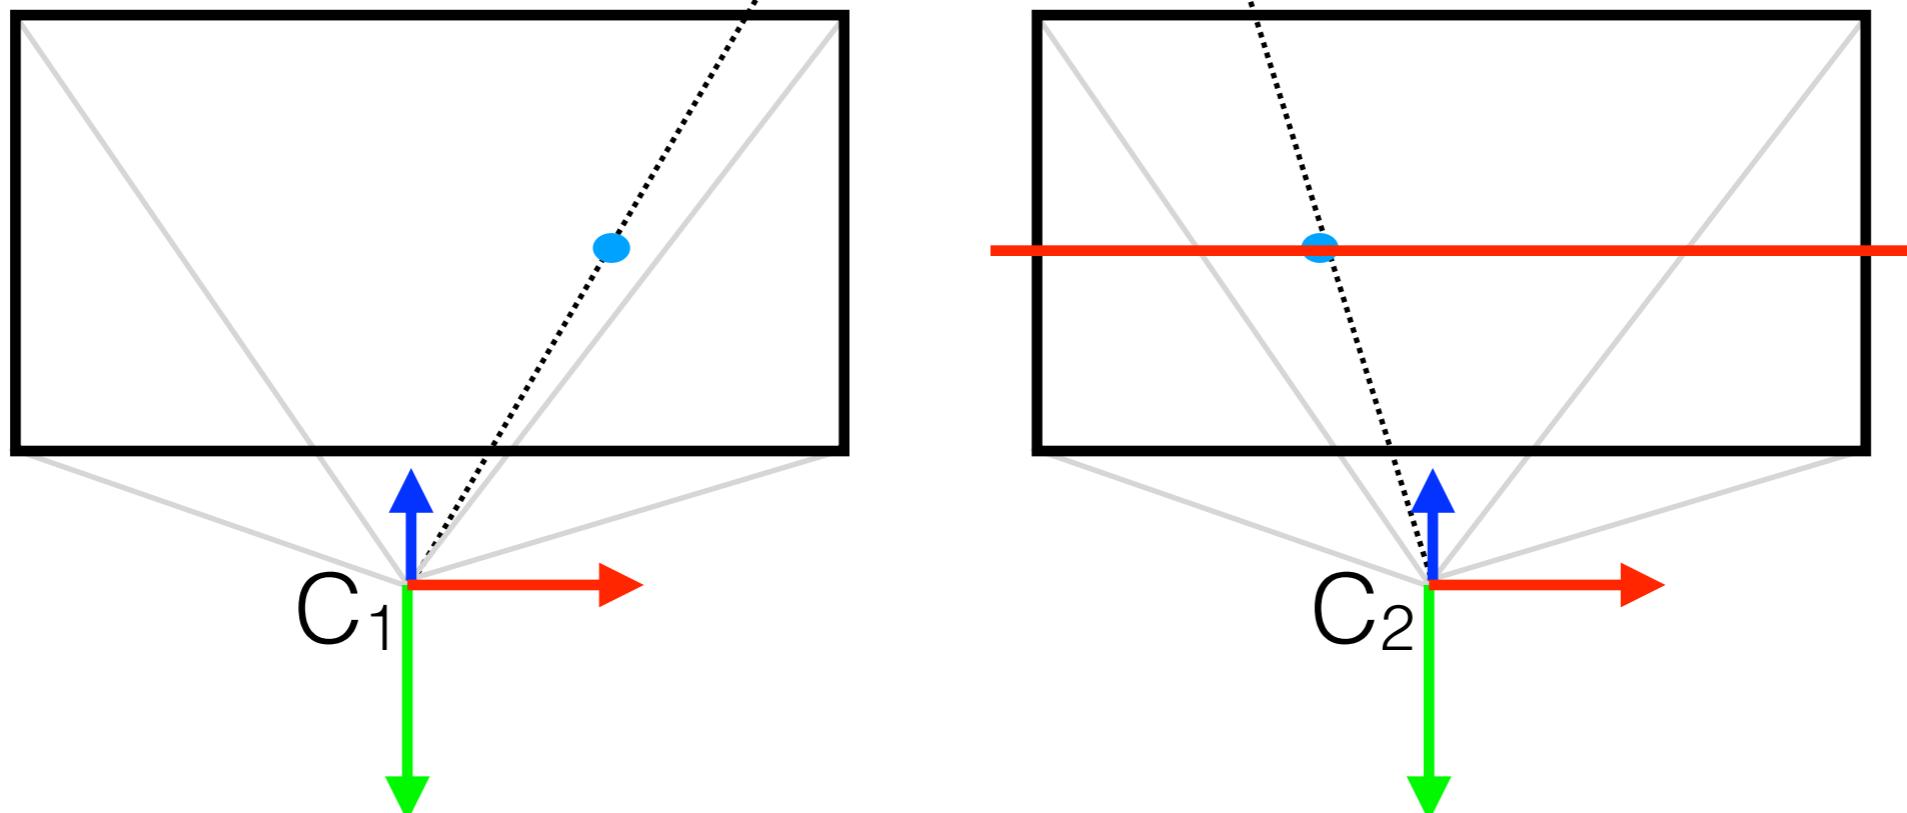
e_L, e_R : epipoles

Example: Stereo Camera



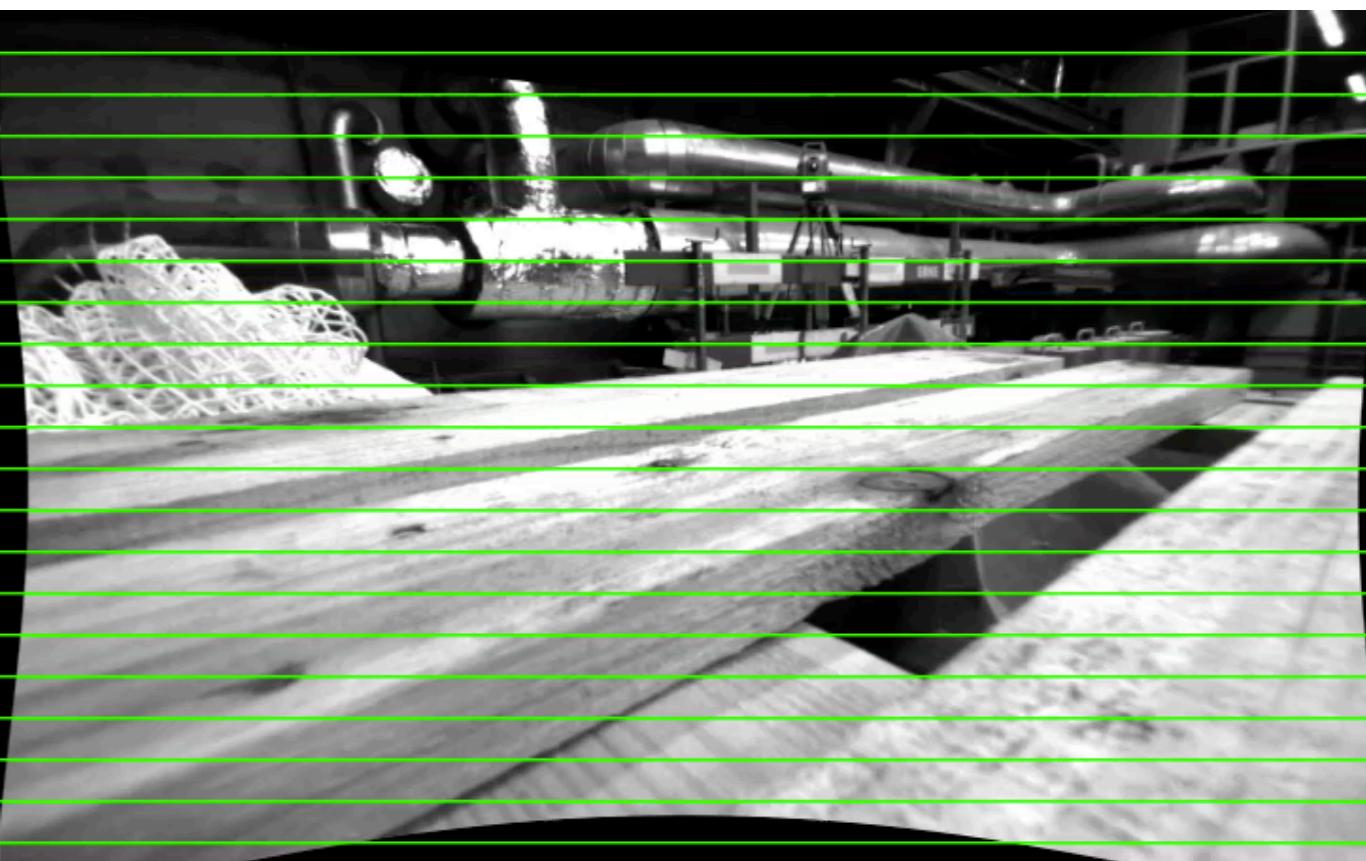
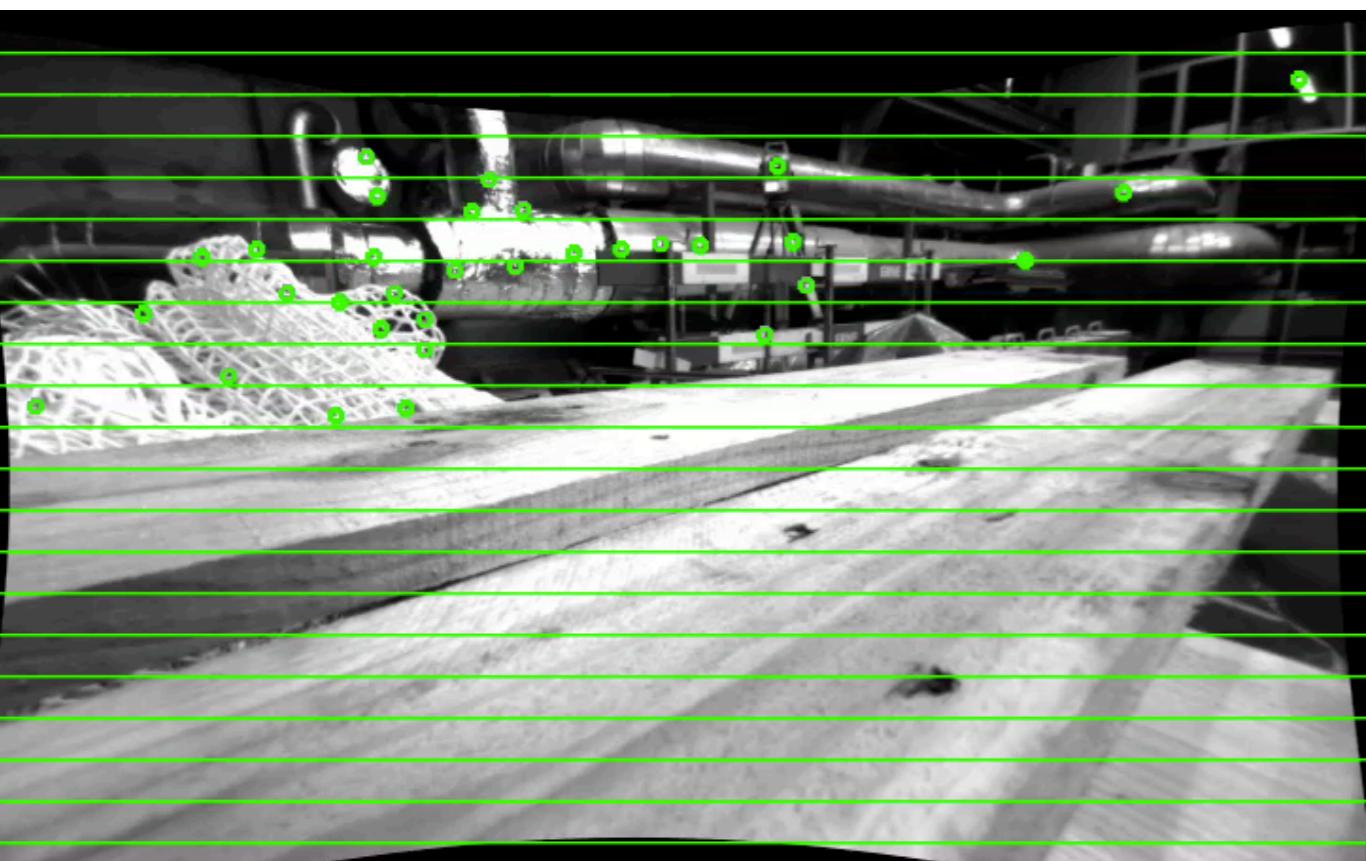
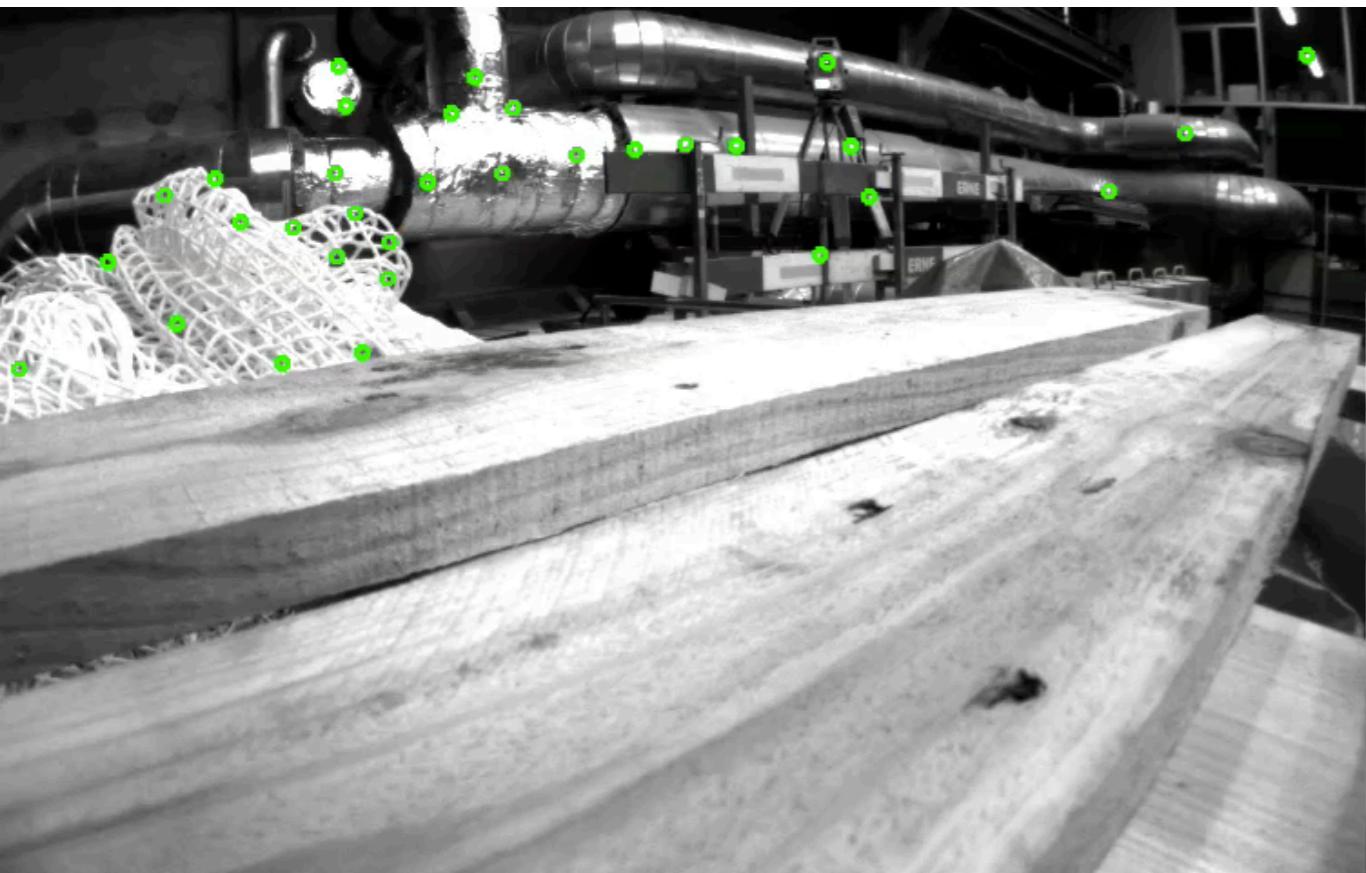
$$R_{C_1}^{C_2} = I_3$$

$$t_{C_1}^{C_2} = \begin{bmatrix} -b \\ 0 \\ 0 \end{bmatrix}$$

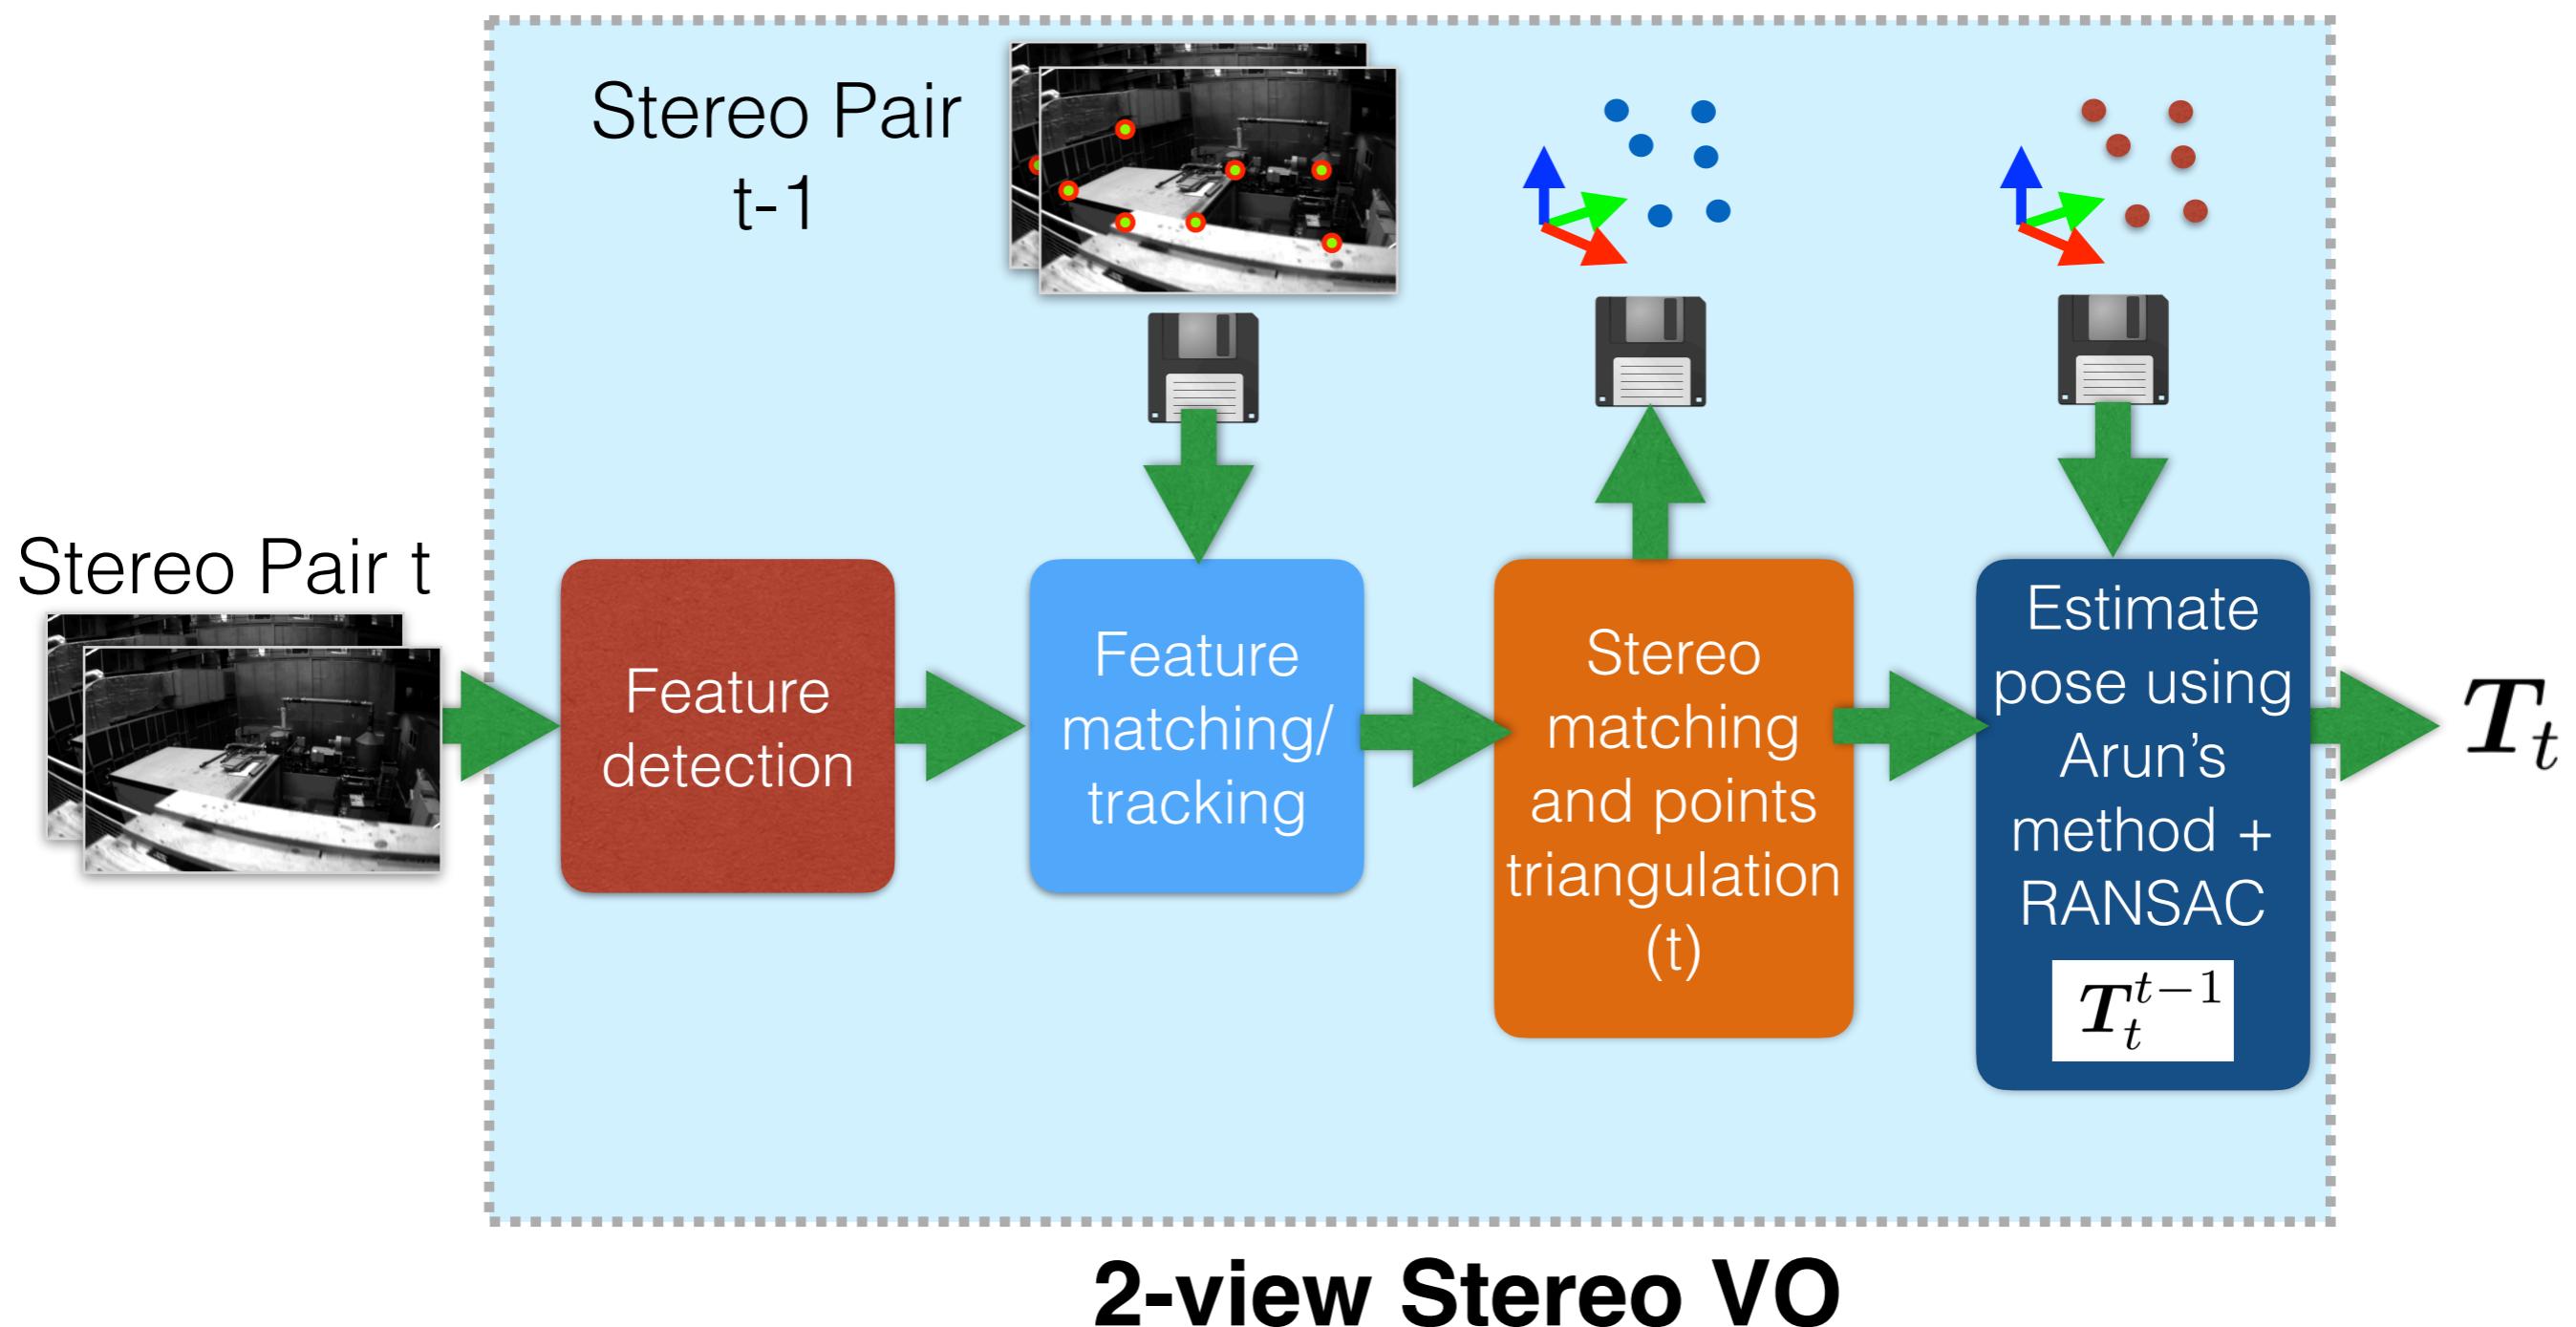


also: easy to triangulate points given geometry

Stereo Matching

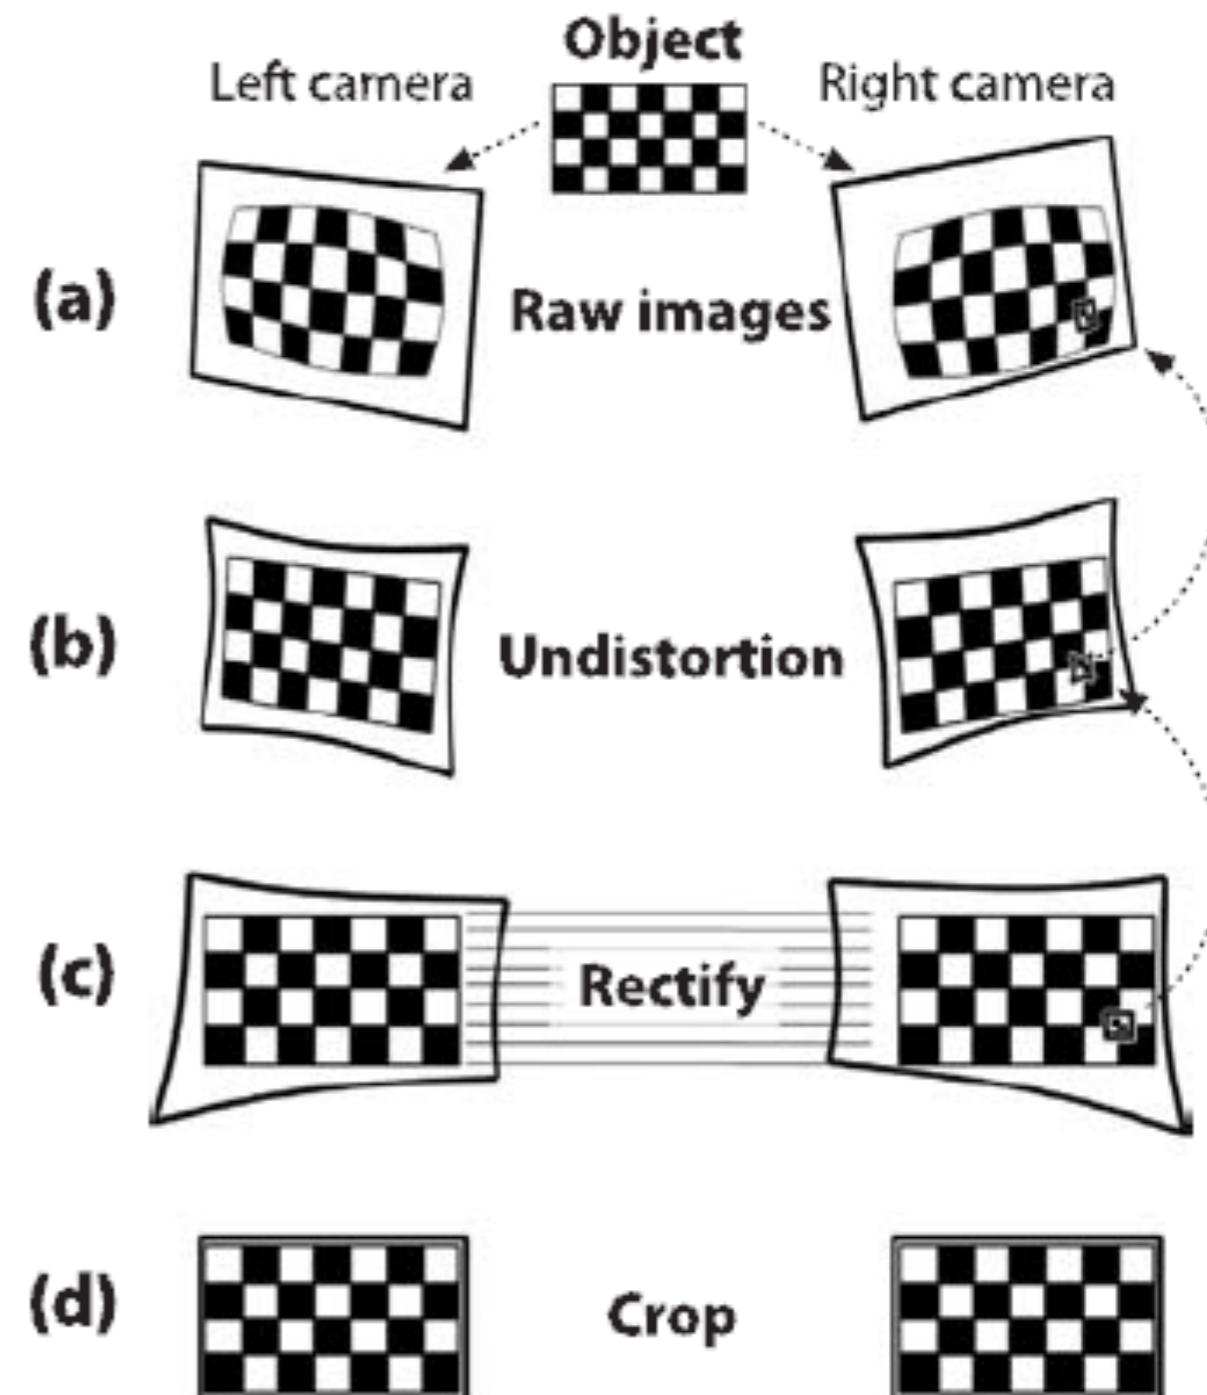
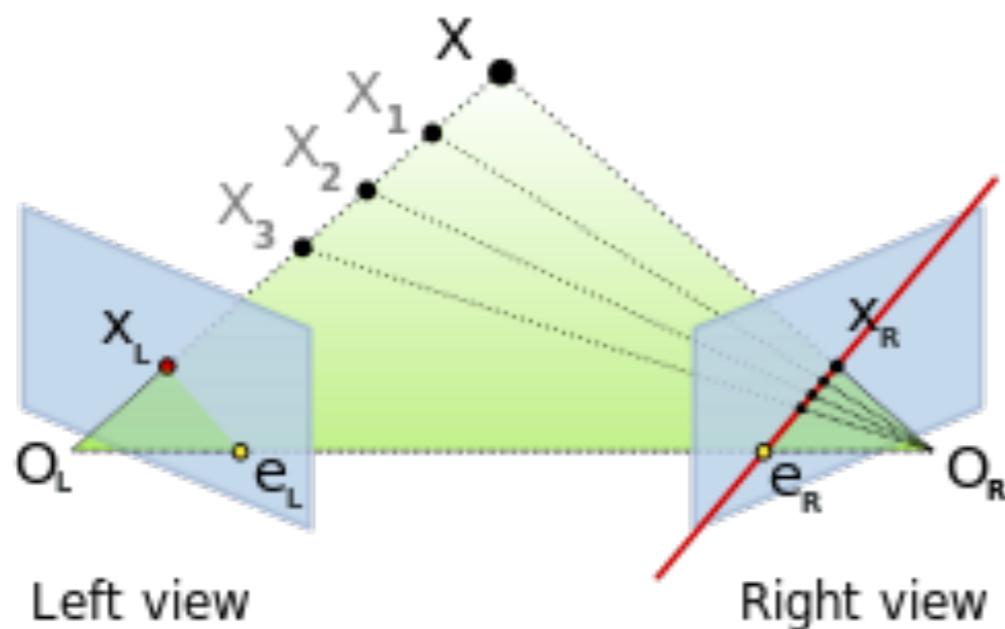
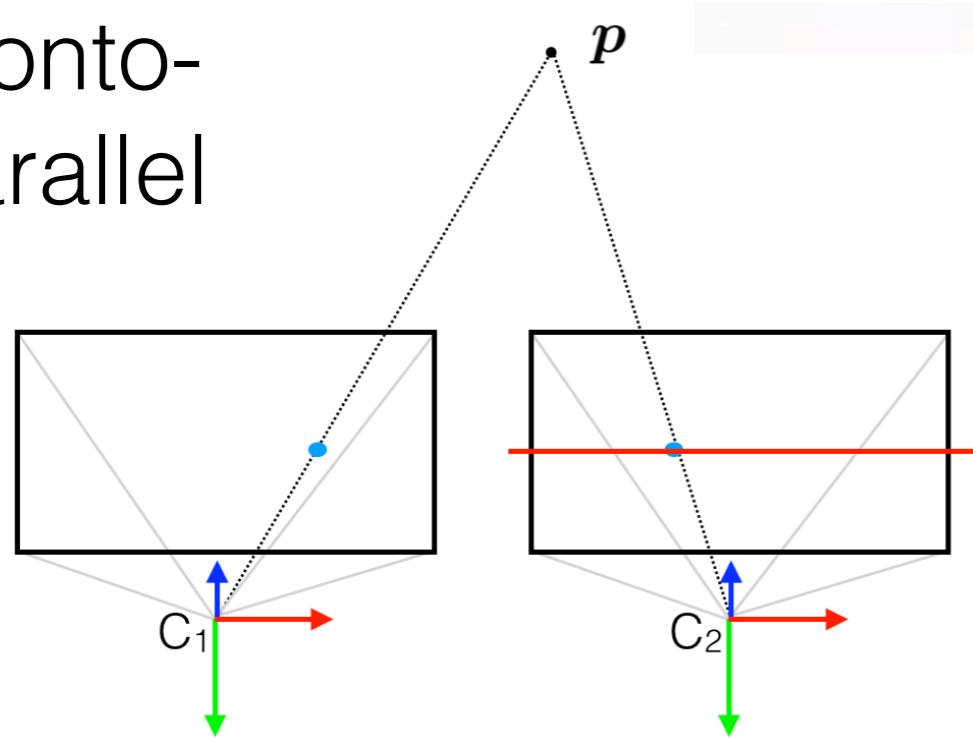


Stereo VO with **3D-3D** Correspondences



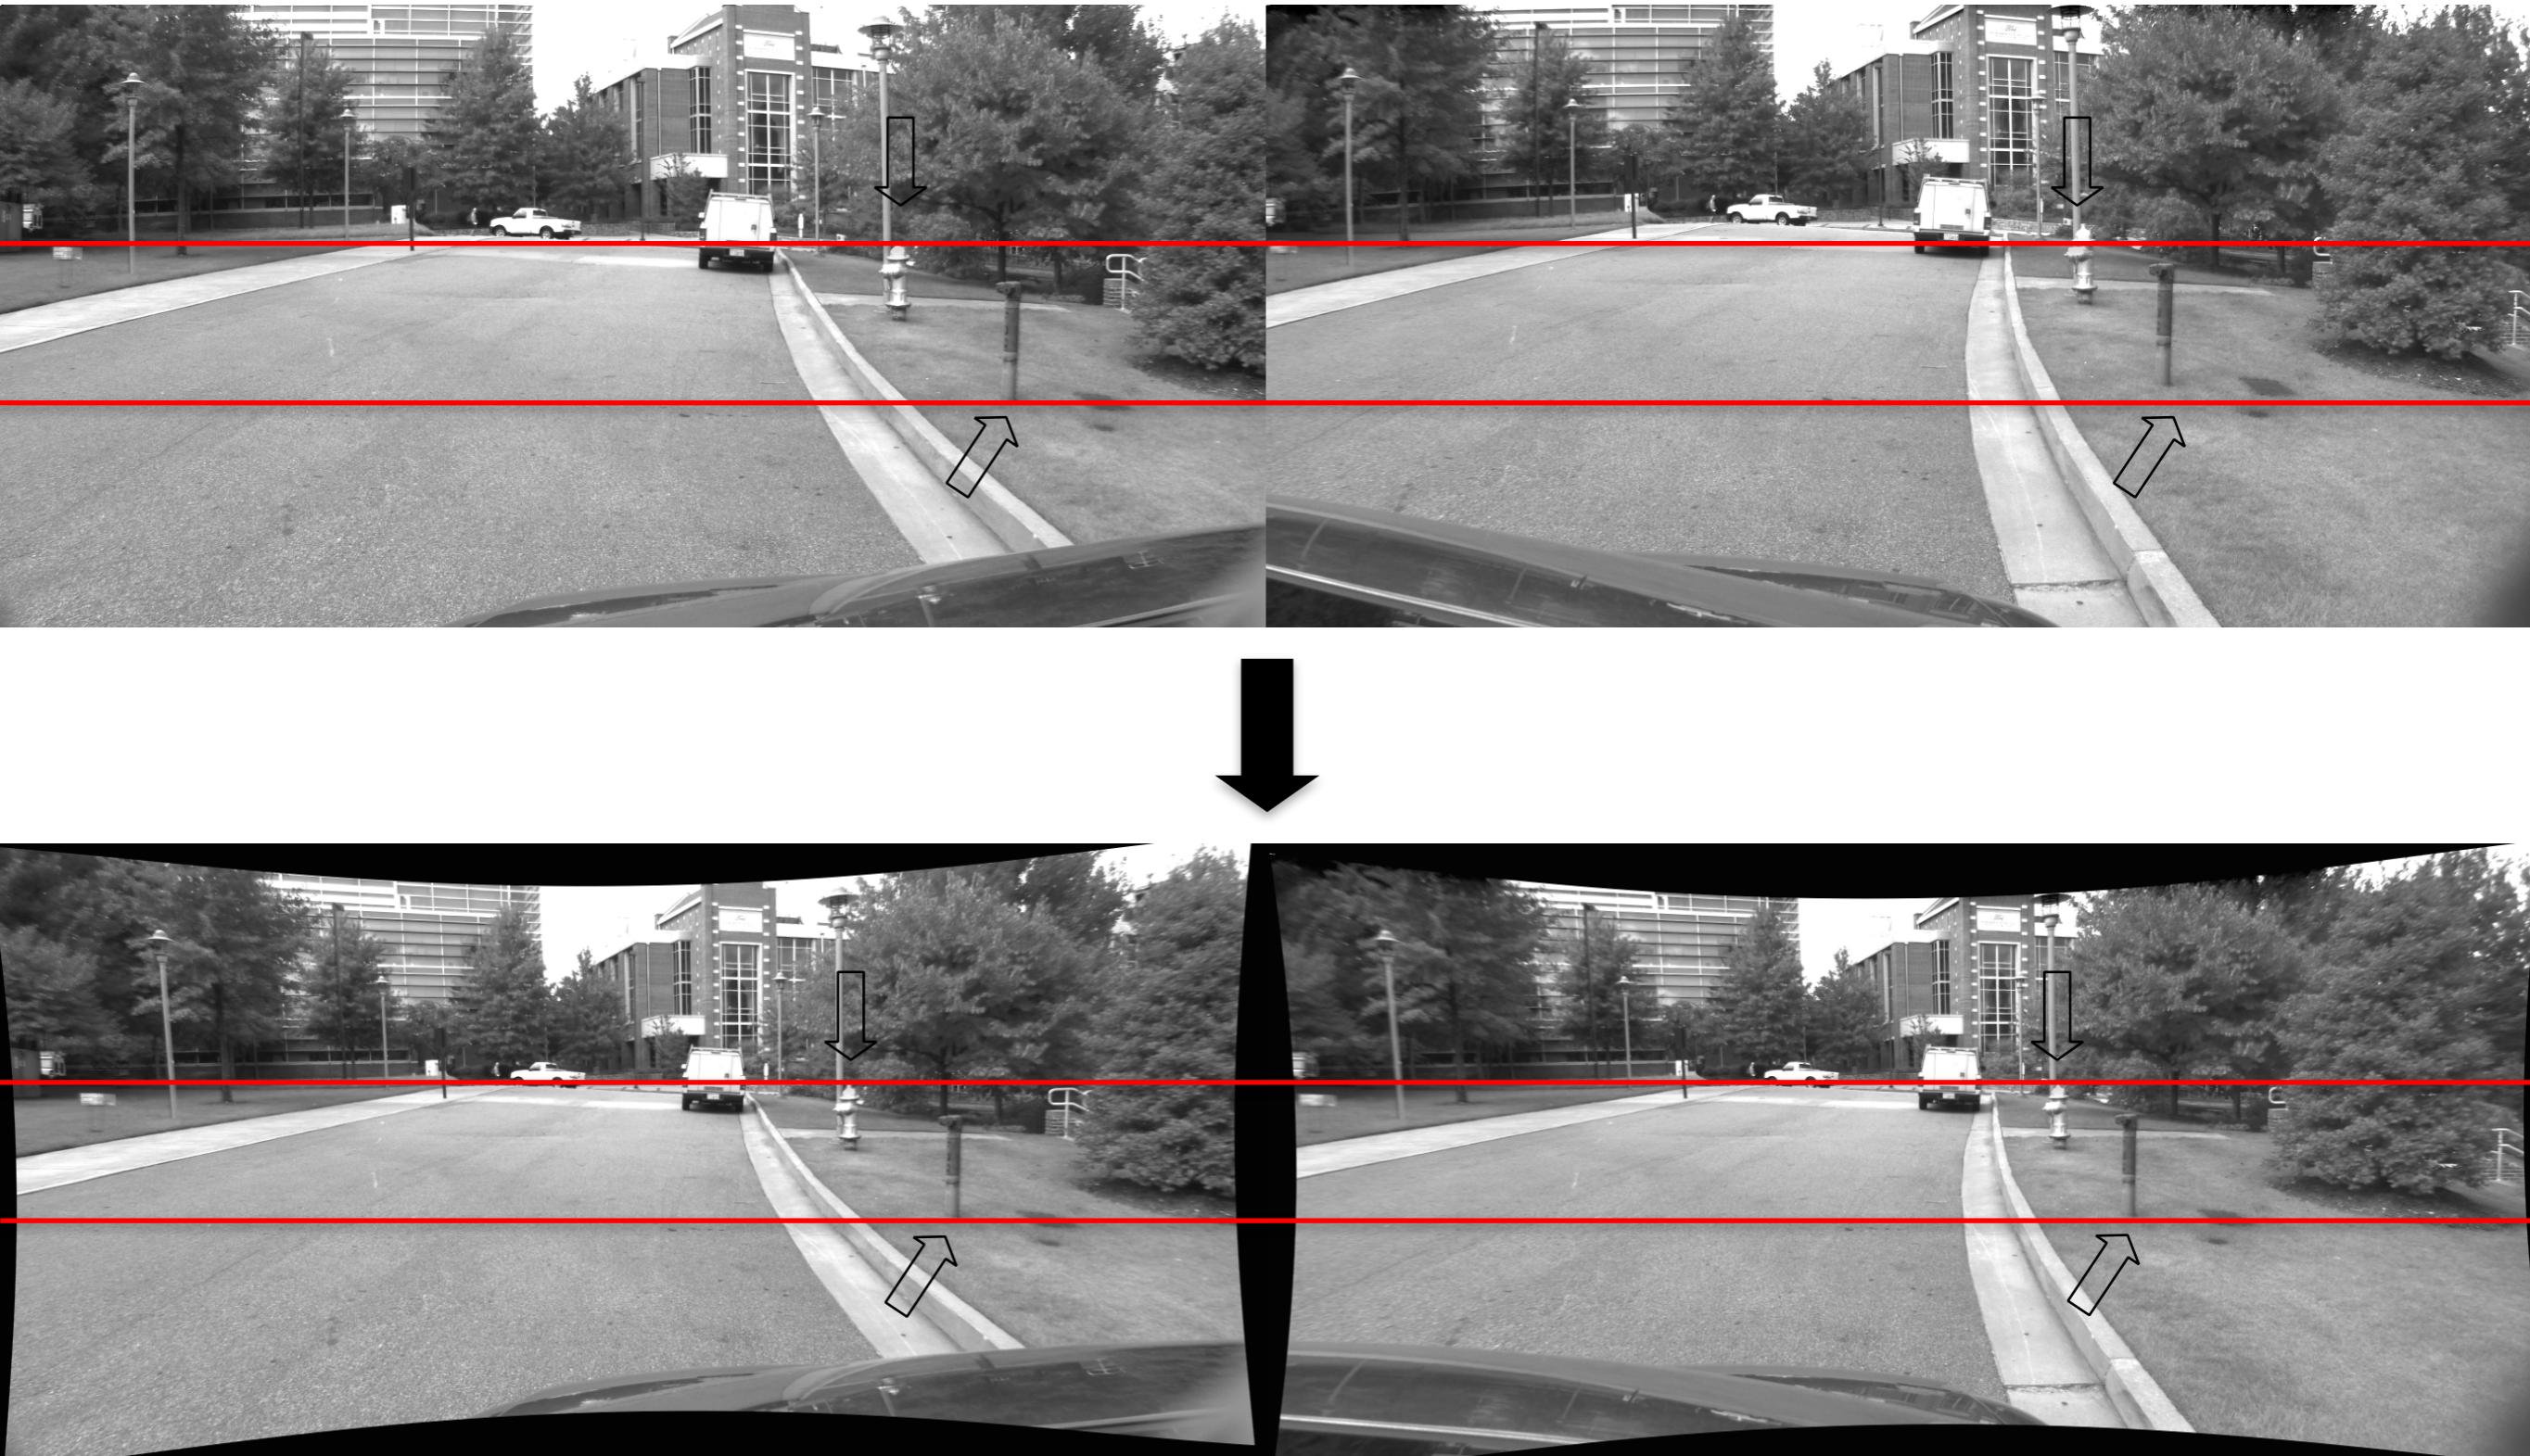
(Parenthesis on Stereo Matching)

Fronto-
parallel



OpenCV: `stereoRectify`, `initUndistortRectifyMap`

(Parenthesis on Stereo Matching)



[courtesy of Frank Dellaert]

(Parenthesis on Stereo Matching)

After **rectification**, we can restrict search for left-right matches to horizontal lines

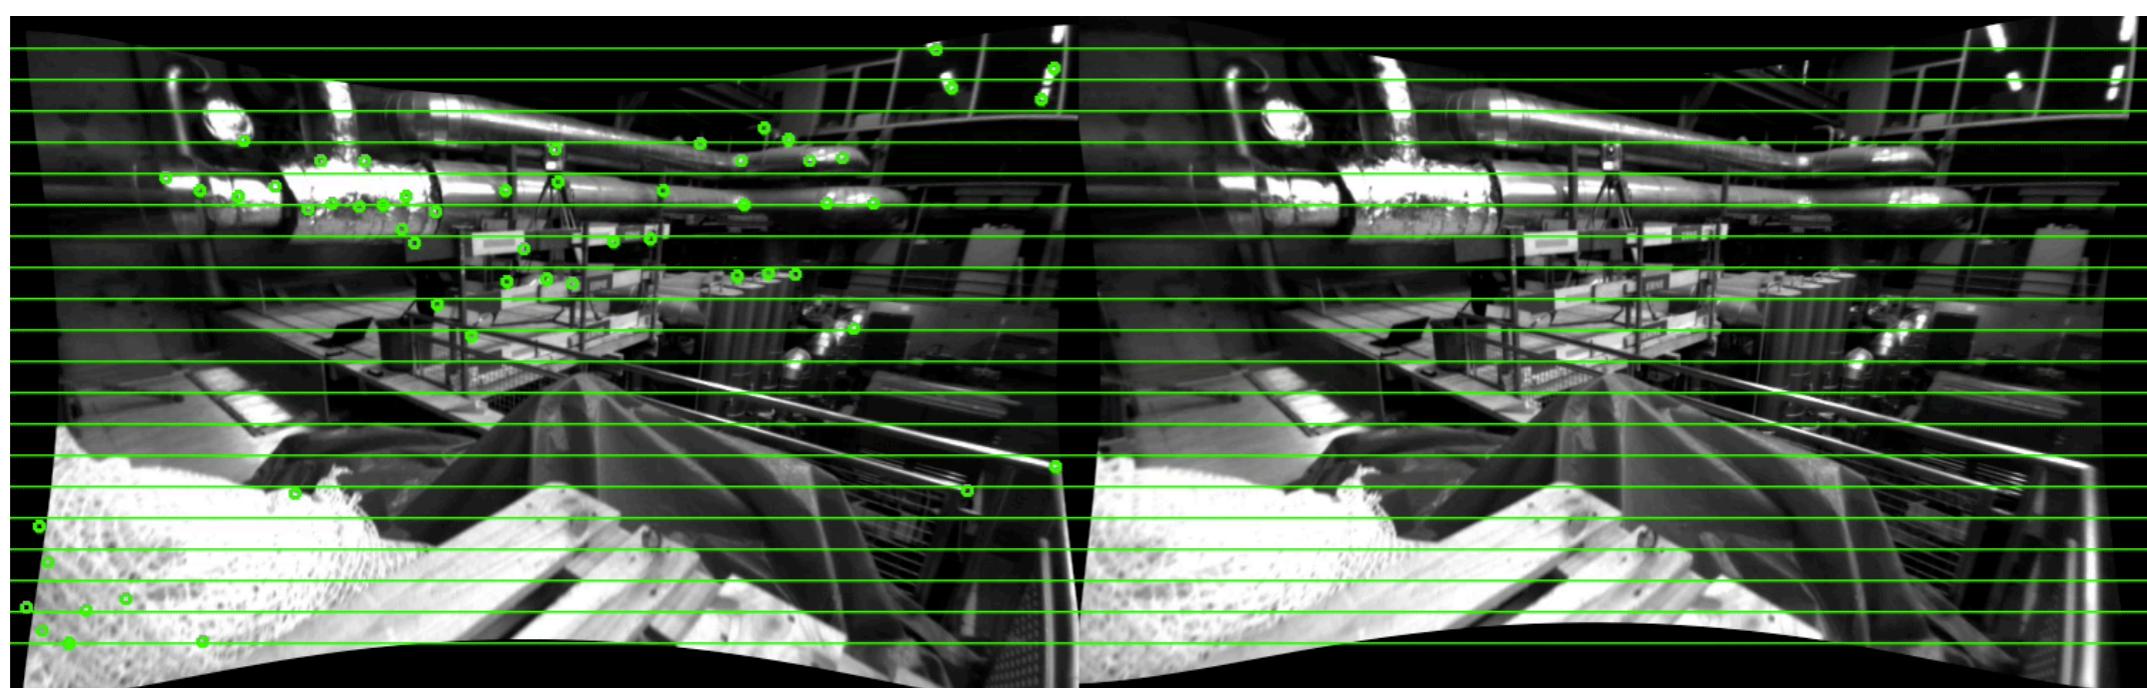
Left image



Right image



[courtesy of Frank Dellaert
and Pablo Alcantarilla]



Comparing VO approaches

Drift (error accumulation):

$$T_0 = I_4$$

$$T_1 = T_0 T_1^0$$

$$T_2 = T_1 T_2^1 = T_0 T_1^0 T_2^1$$

⋮

$$T_t = T_{t-1} T_t^{t-1} = T_0 T_1^0 T_2^1 \cdots T_t^{t-1}$$

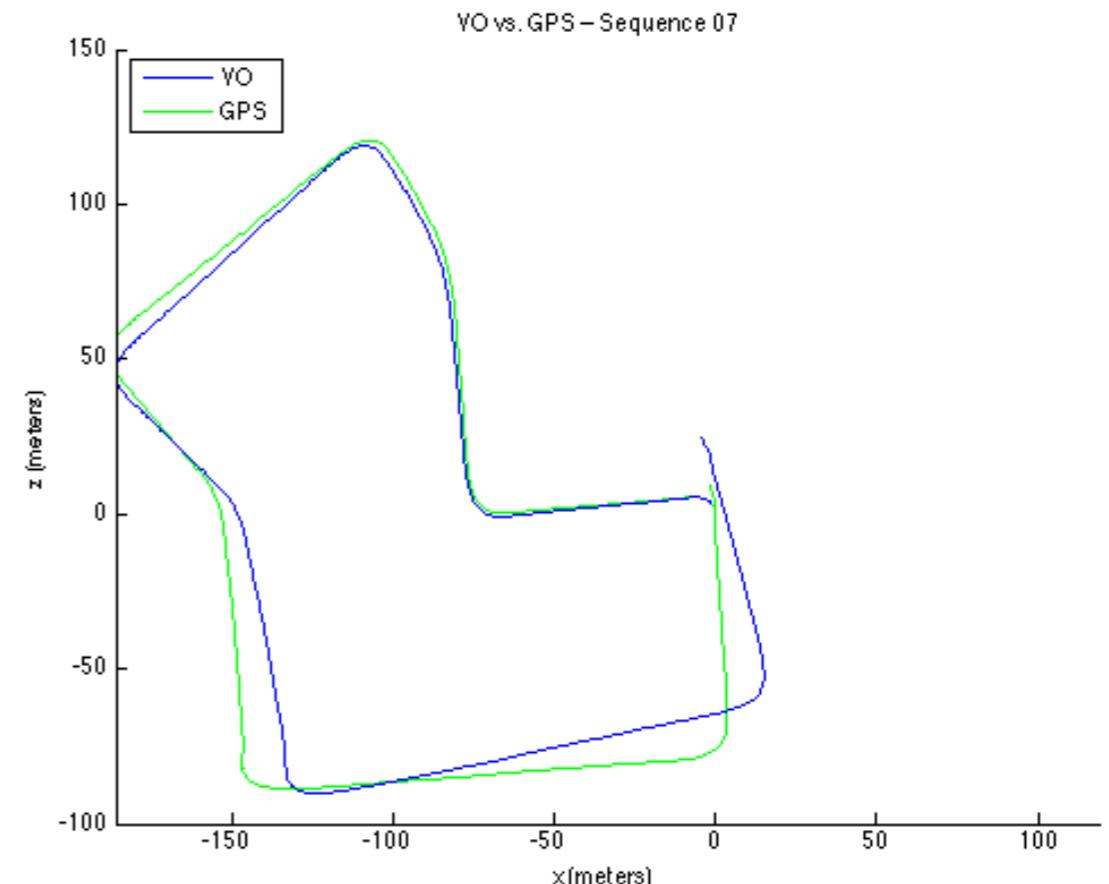
Mono VO:

- 5-point method accurate

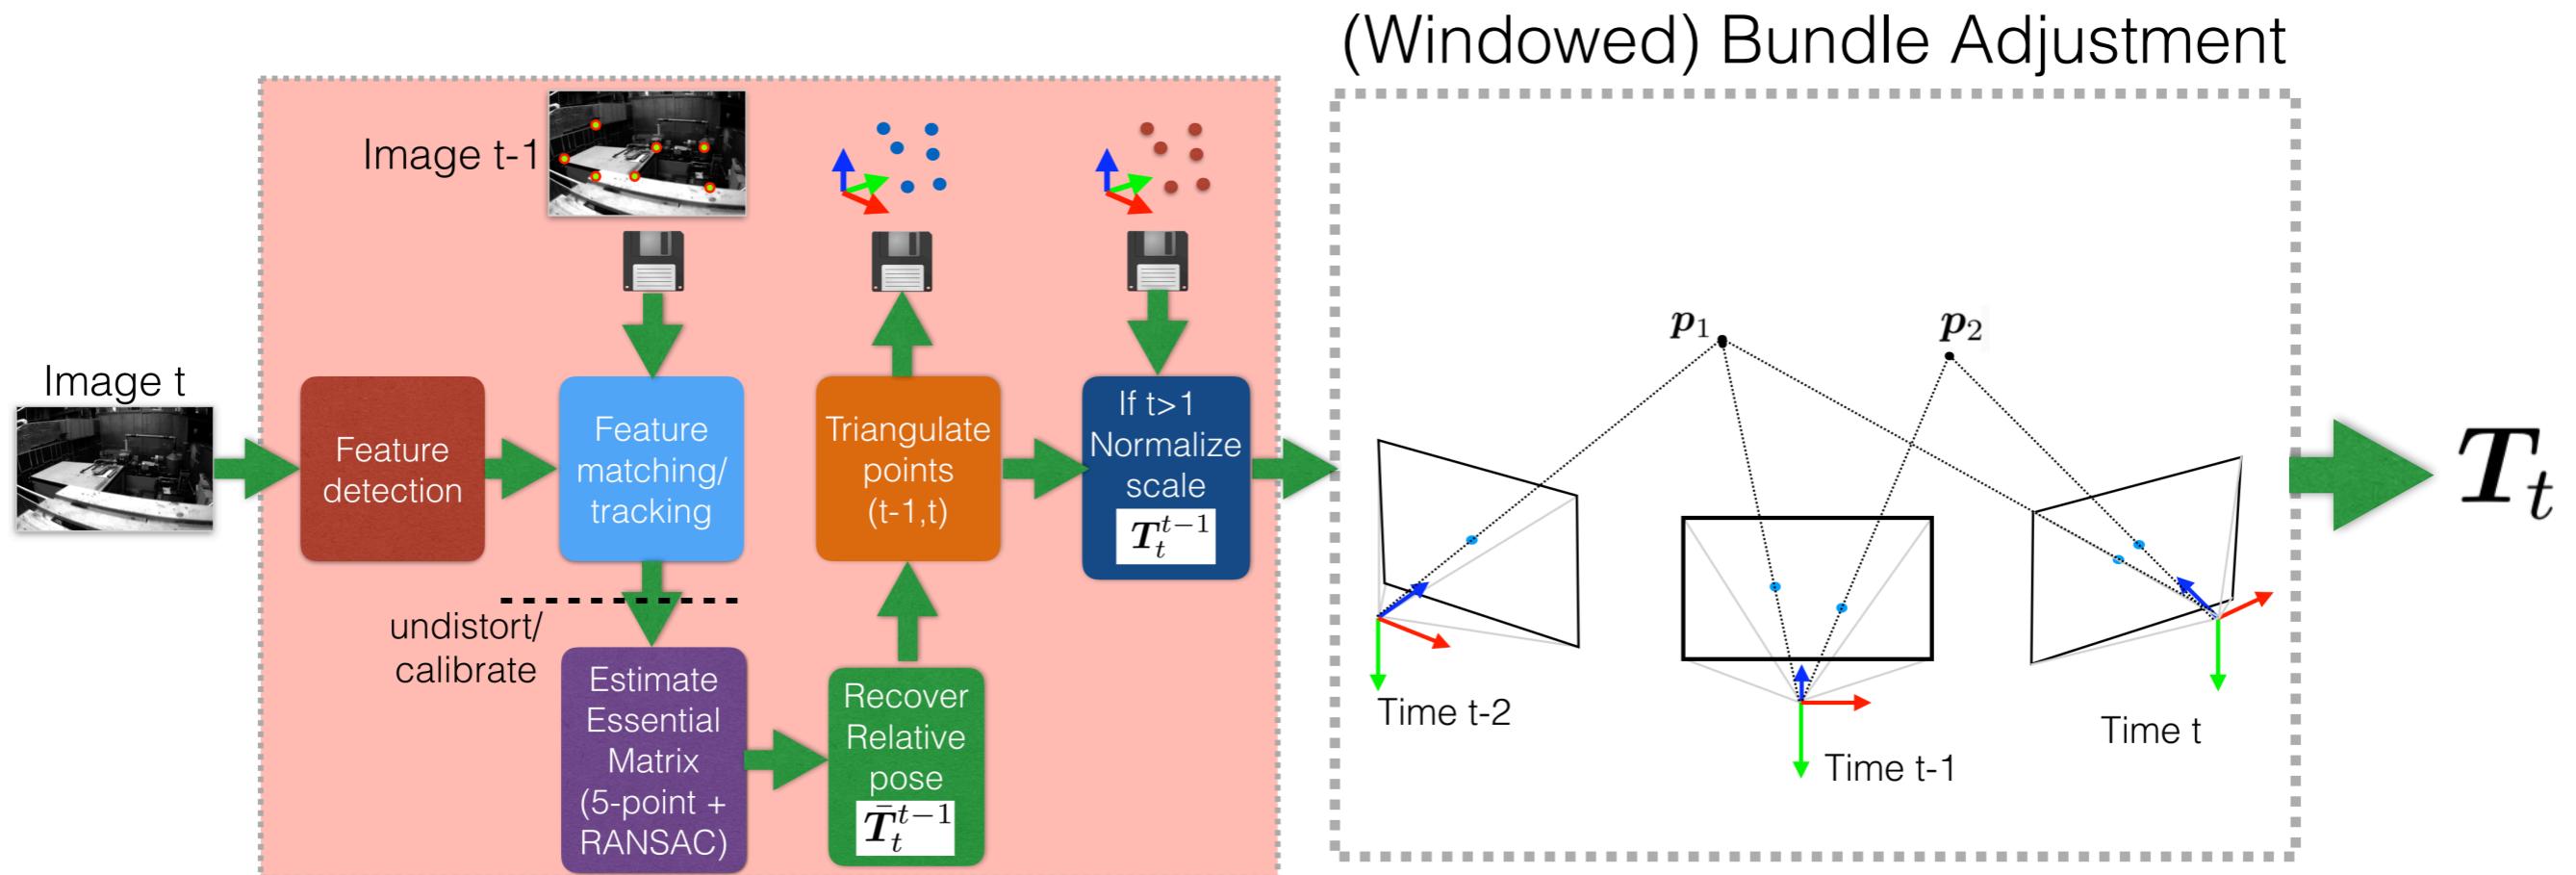
Stereo VO:

- scale

Can we do better?



Refinement: Bundle Adjustment

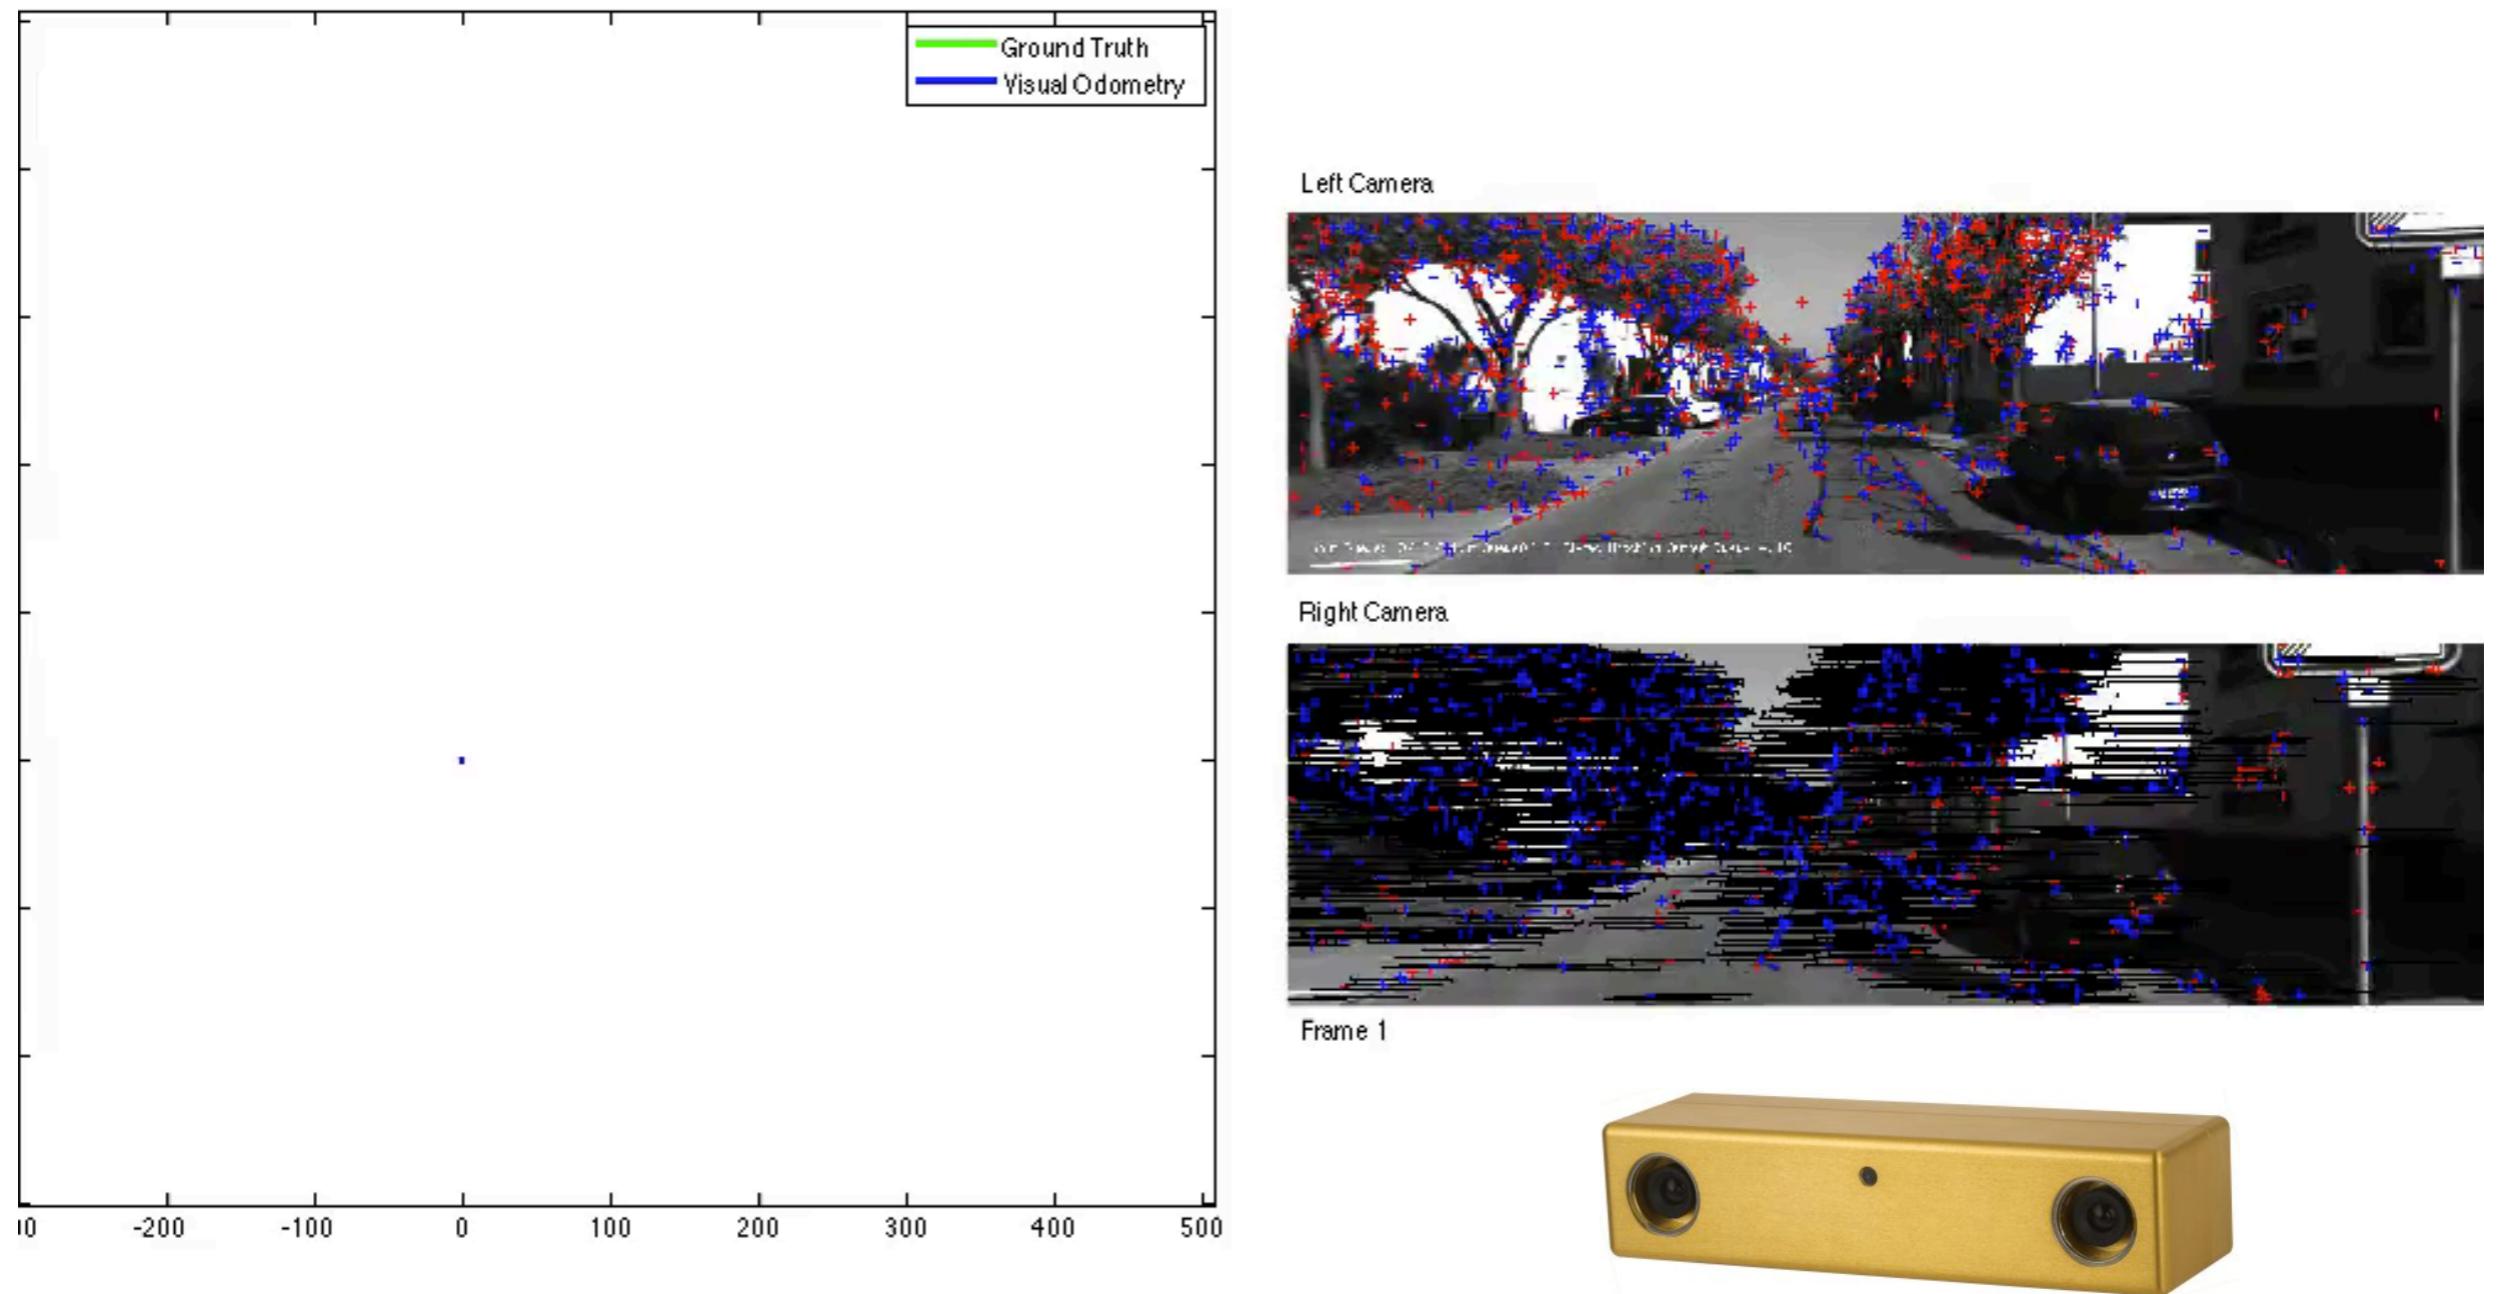


Windowed Bundle Adjustment: optimization of the most recent camera poses and points via non-linear least squares

$$\min_{\substack{\mathbf{T}_i, i=1, \dots, N_C \\ \mathbf{p}_k, k=1, \dots, N}} \sum_{k=1}^N \sum_{i \in \mathcal{C}_k} \|\mathbf{x}_{k,i} - \pi(\mathbf{T}_i, \mathbf{p}_k)\|^2$$

Can be applied to all the pipelines discussed today

Stereo VO example (2)



Typical drifts: 0.1% to 2% of trajectory travelled

[courtesy of Frank Dellaert]

Challenges for VO (1/3): Illumination and Features



Feature detection,
tracking,
matching ...



Challenges for VO (2/3): Dynamic Scenes

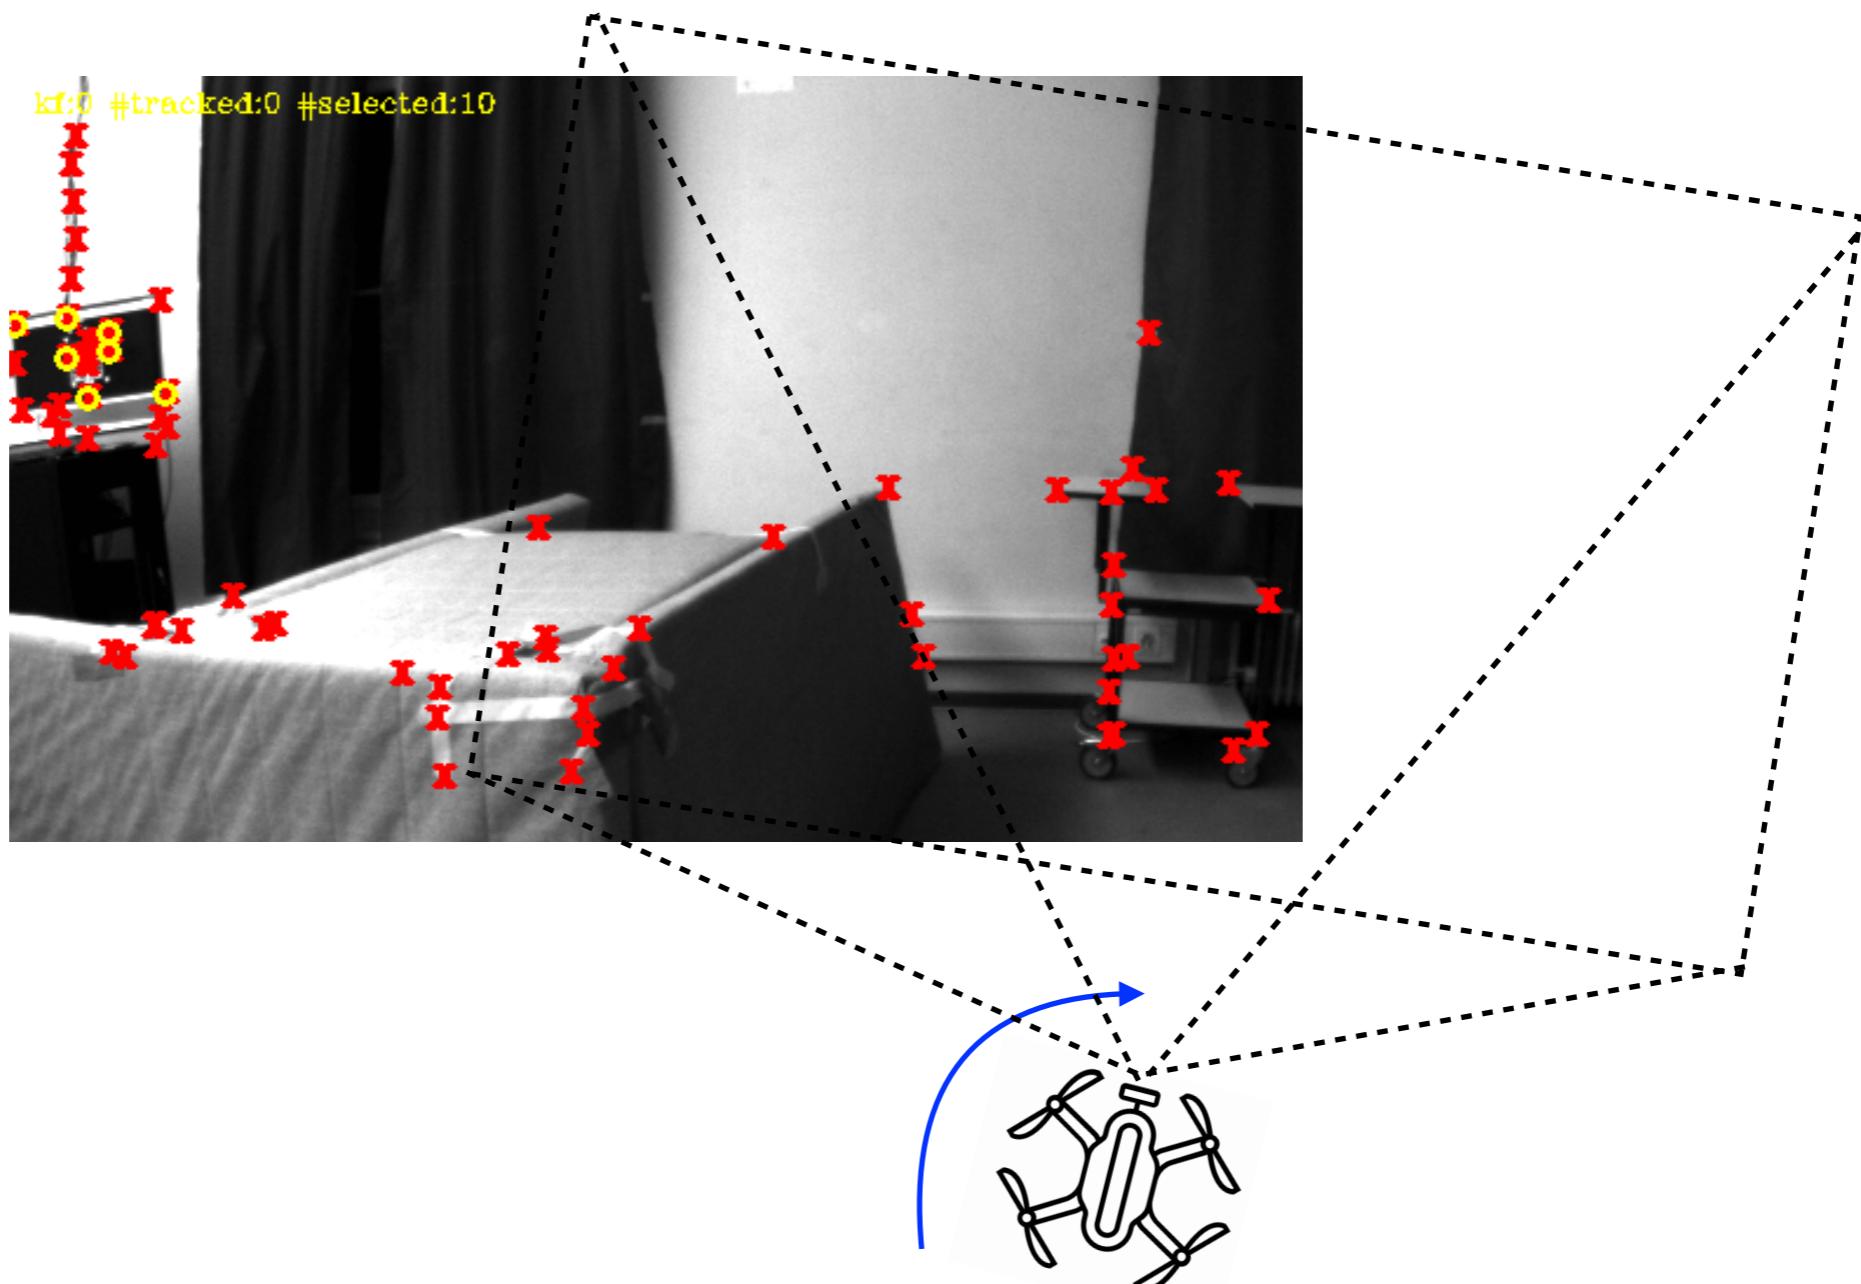
- Dynamic, crowded scenes present a real challenge
- Can't rely on RANSAC to always recover the correct inliers
- Example: Large van “steals” inlier set in passing



[courtesy of Frank Dellaert]

Challenges for VO (3/3): Fast Motion

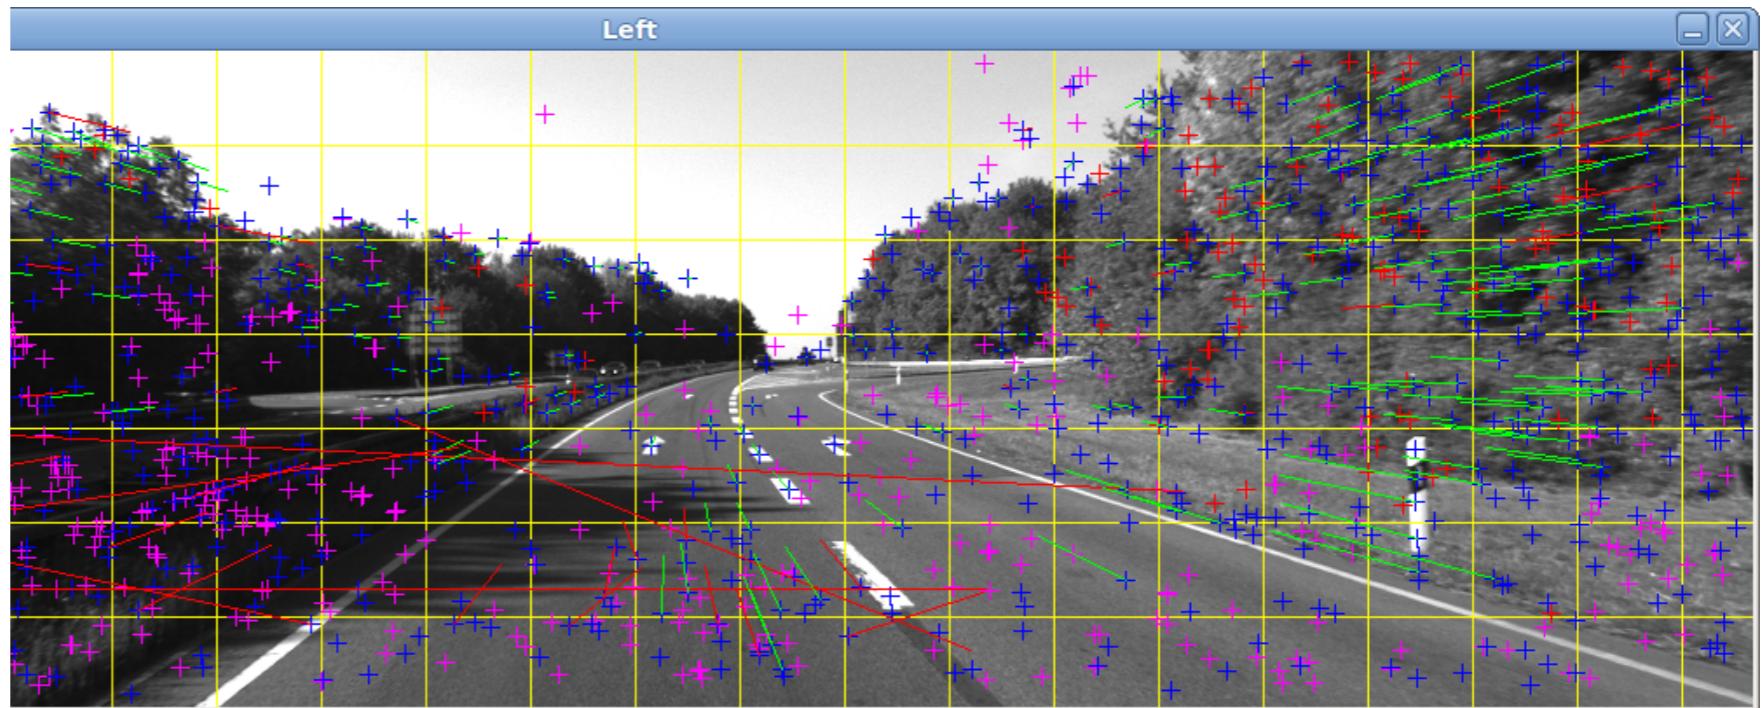
Need good overlap between consecutive images



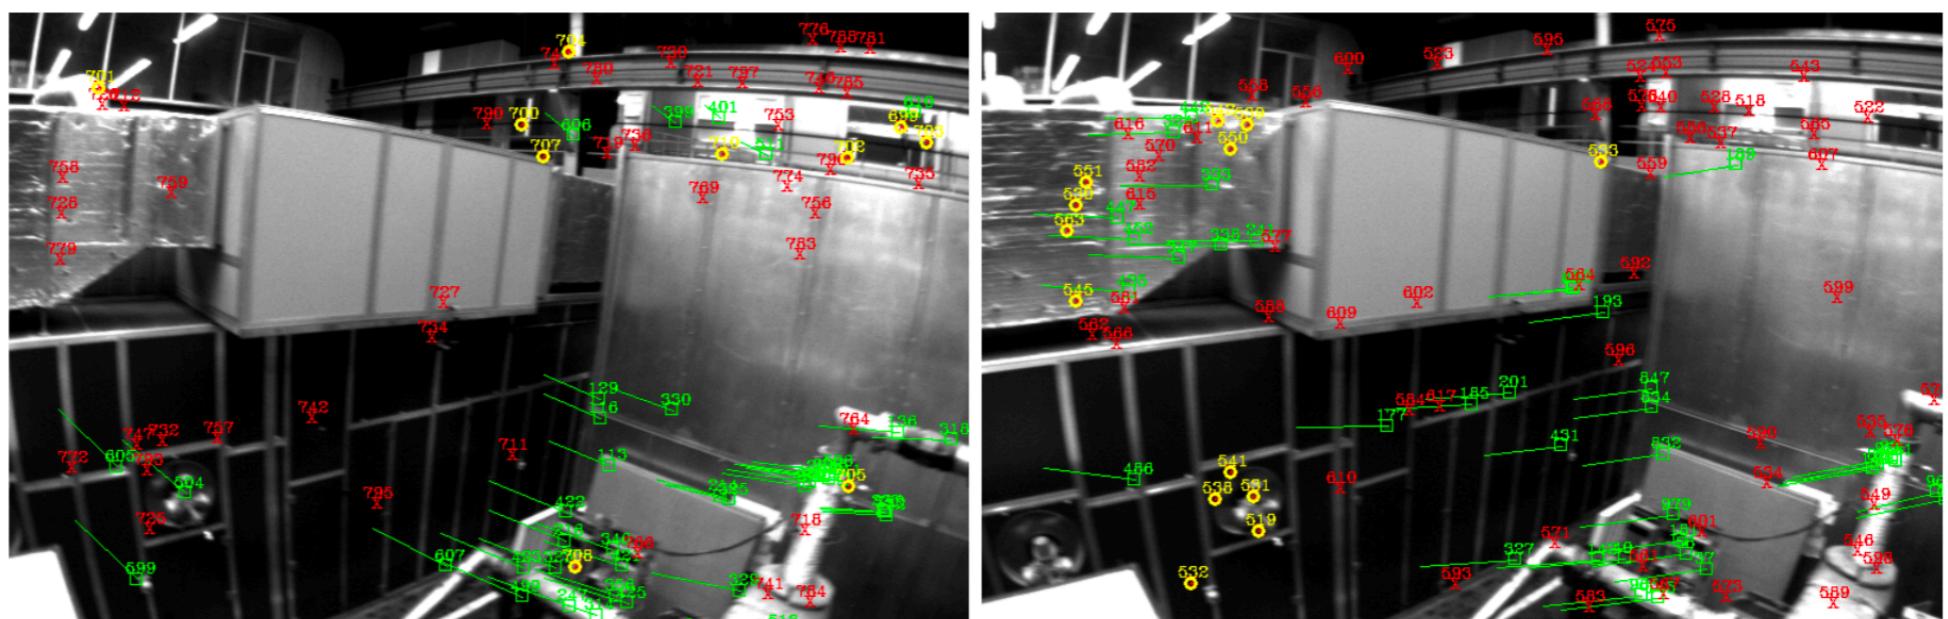
Robot speed, camera framerate, ...

VO Tricks (1/2): Feature Distribution

Feature Binning:



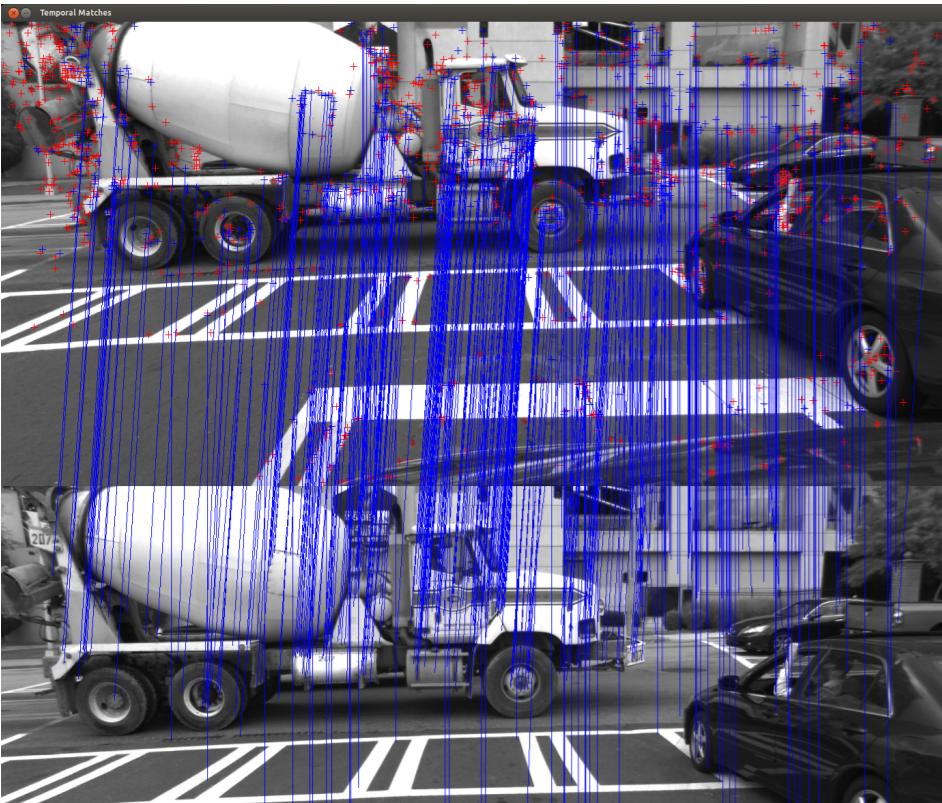
Attention & Anticipation: (Carlone '17)



select features depending on motion of the robot

VO Tricks (2/2): Domain Knowledge and Keyframes

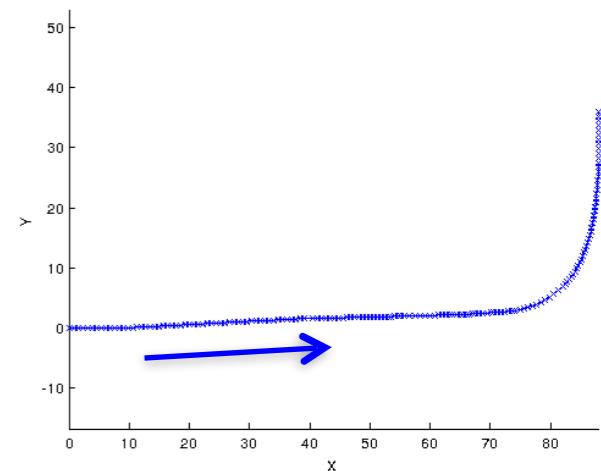
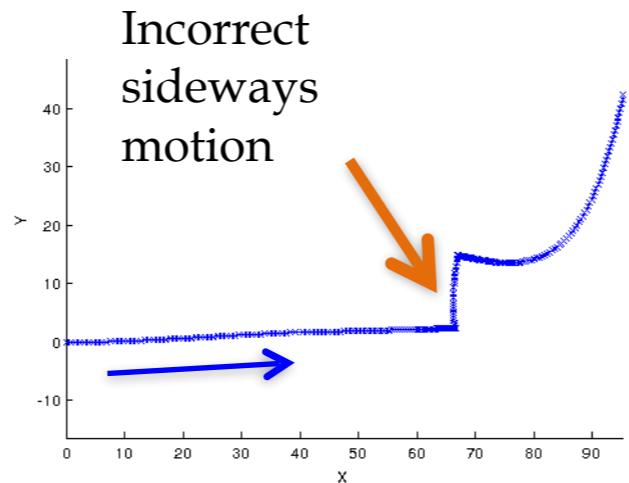
- Stereo VO Example: Cross-traffic while waiting to turn left at light



Without keyframing

Only accept incremental pose if:

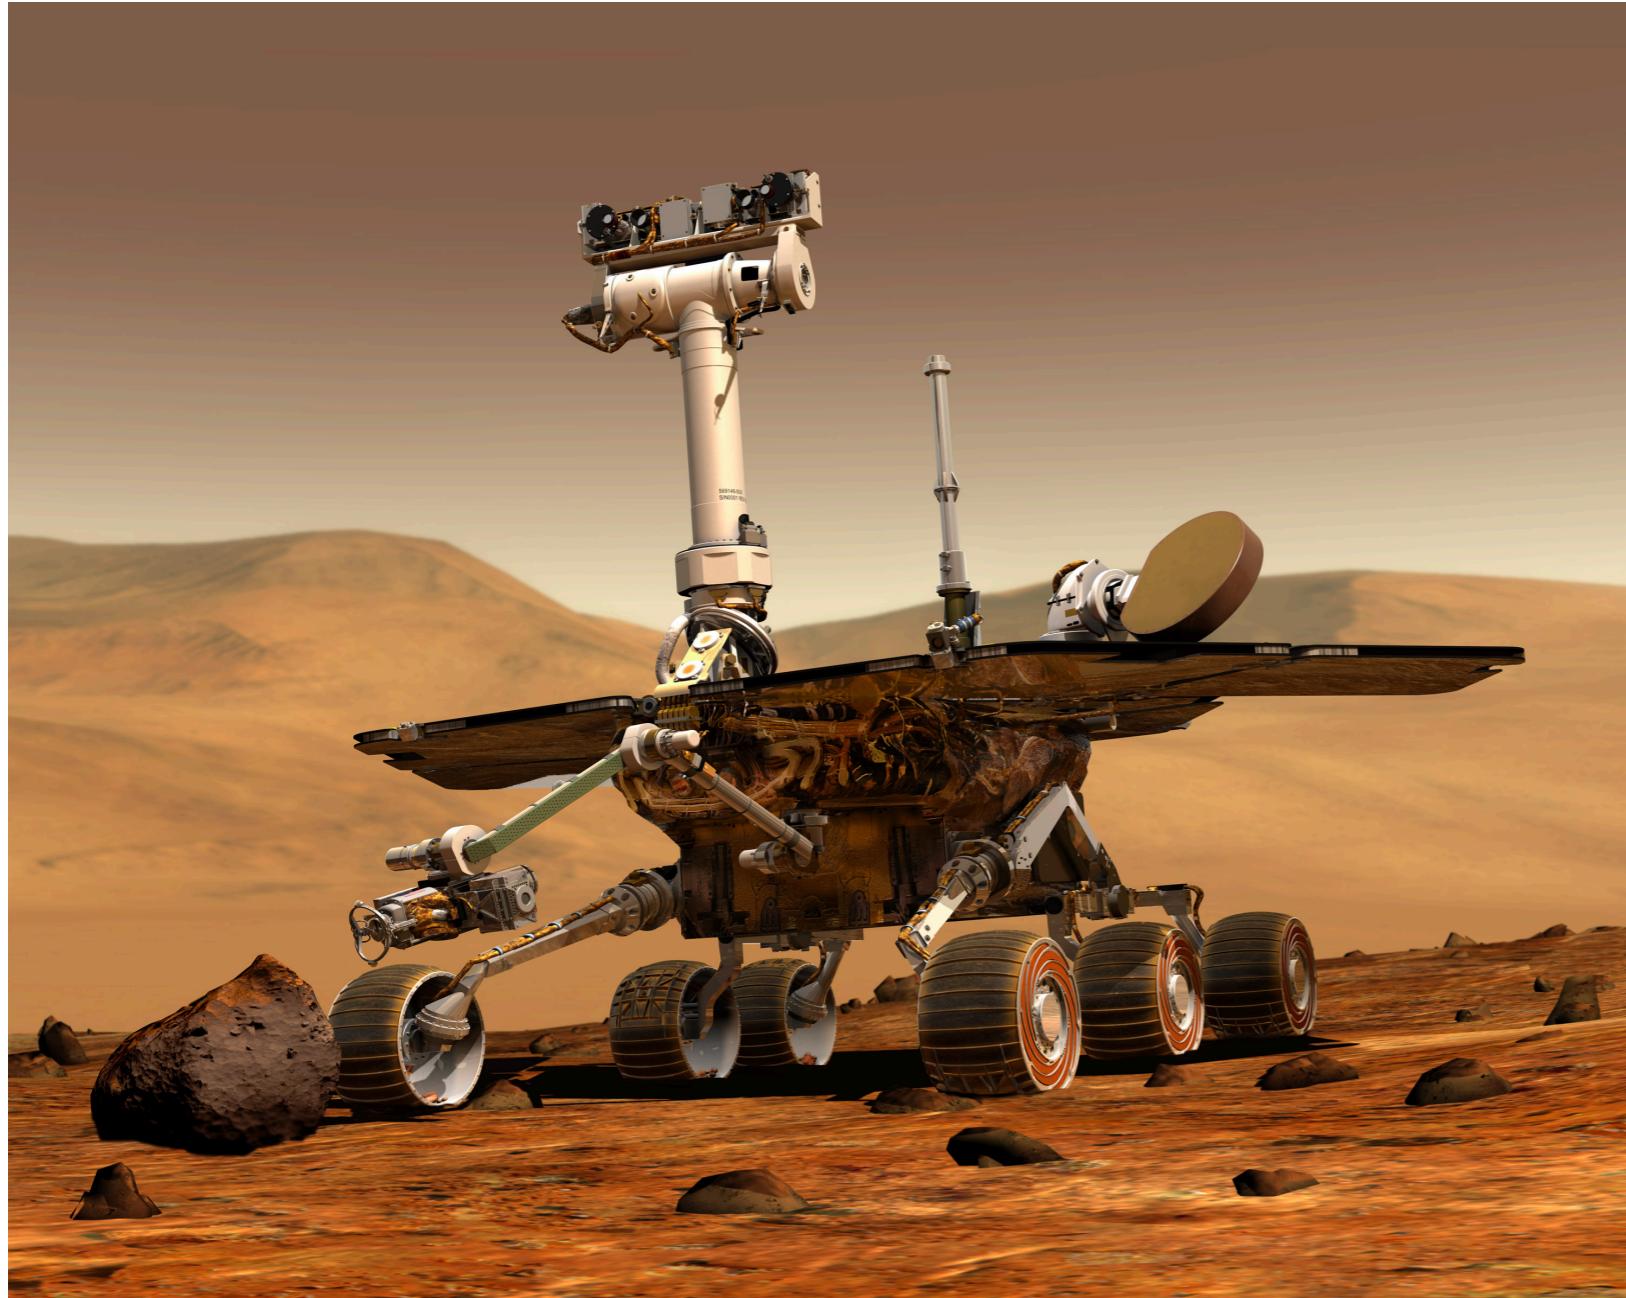
- Translation $> 0.5m$
- Dominant direction is forward



With keyframing

[courtesy of Frank Dellaert]

Stereo VO example (1)



Spirit and Opportunity Mars rovers:

- stereo VO
- 20-MHz CPU
- up to three minutes for 2-view VO
- Drift ~0.5% of trajectory travelled

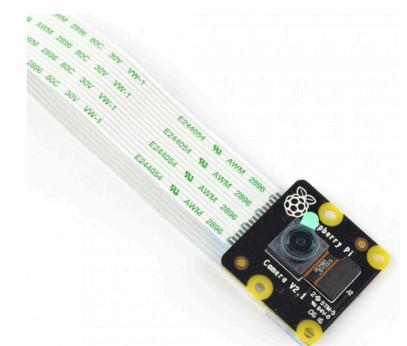
Earlier implementation: Moravec's PhD Thesis (1980)

Beyond VO

How to get scale and improve robustness?

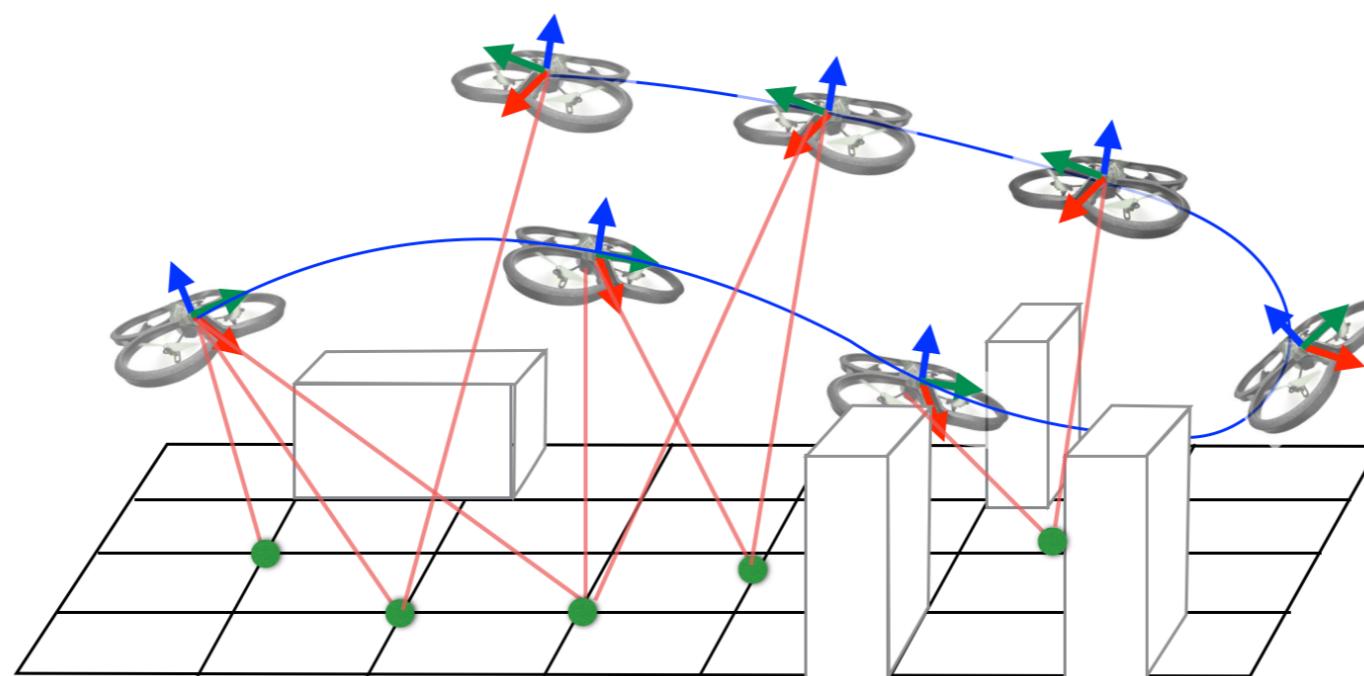
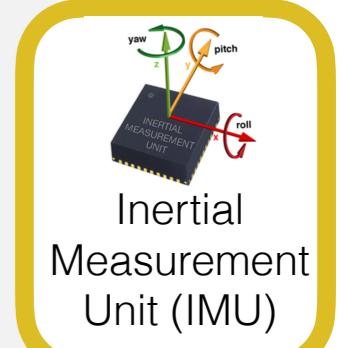
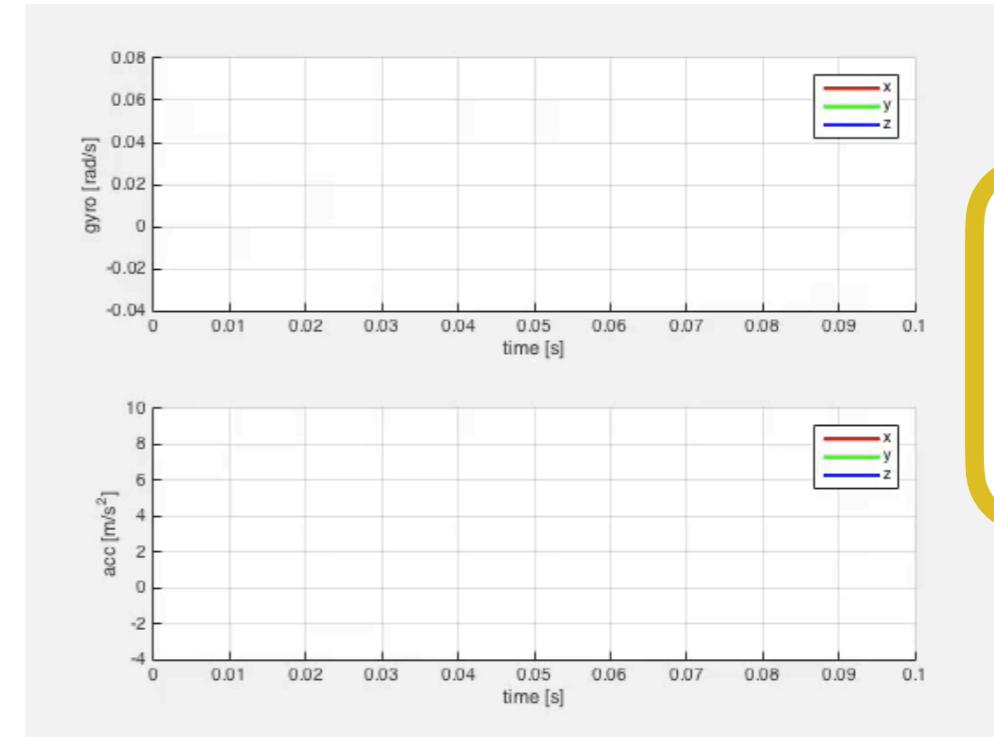
add more sensors!

- ▶ wheel odometry
- ▶ GPS
- ▶ Lidar
- ▶ Inertial Measurement Unit (IMU)

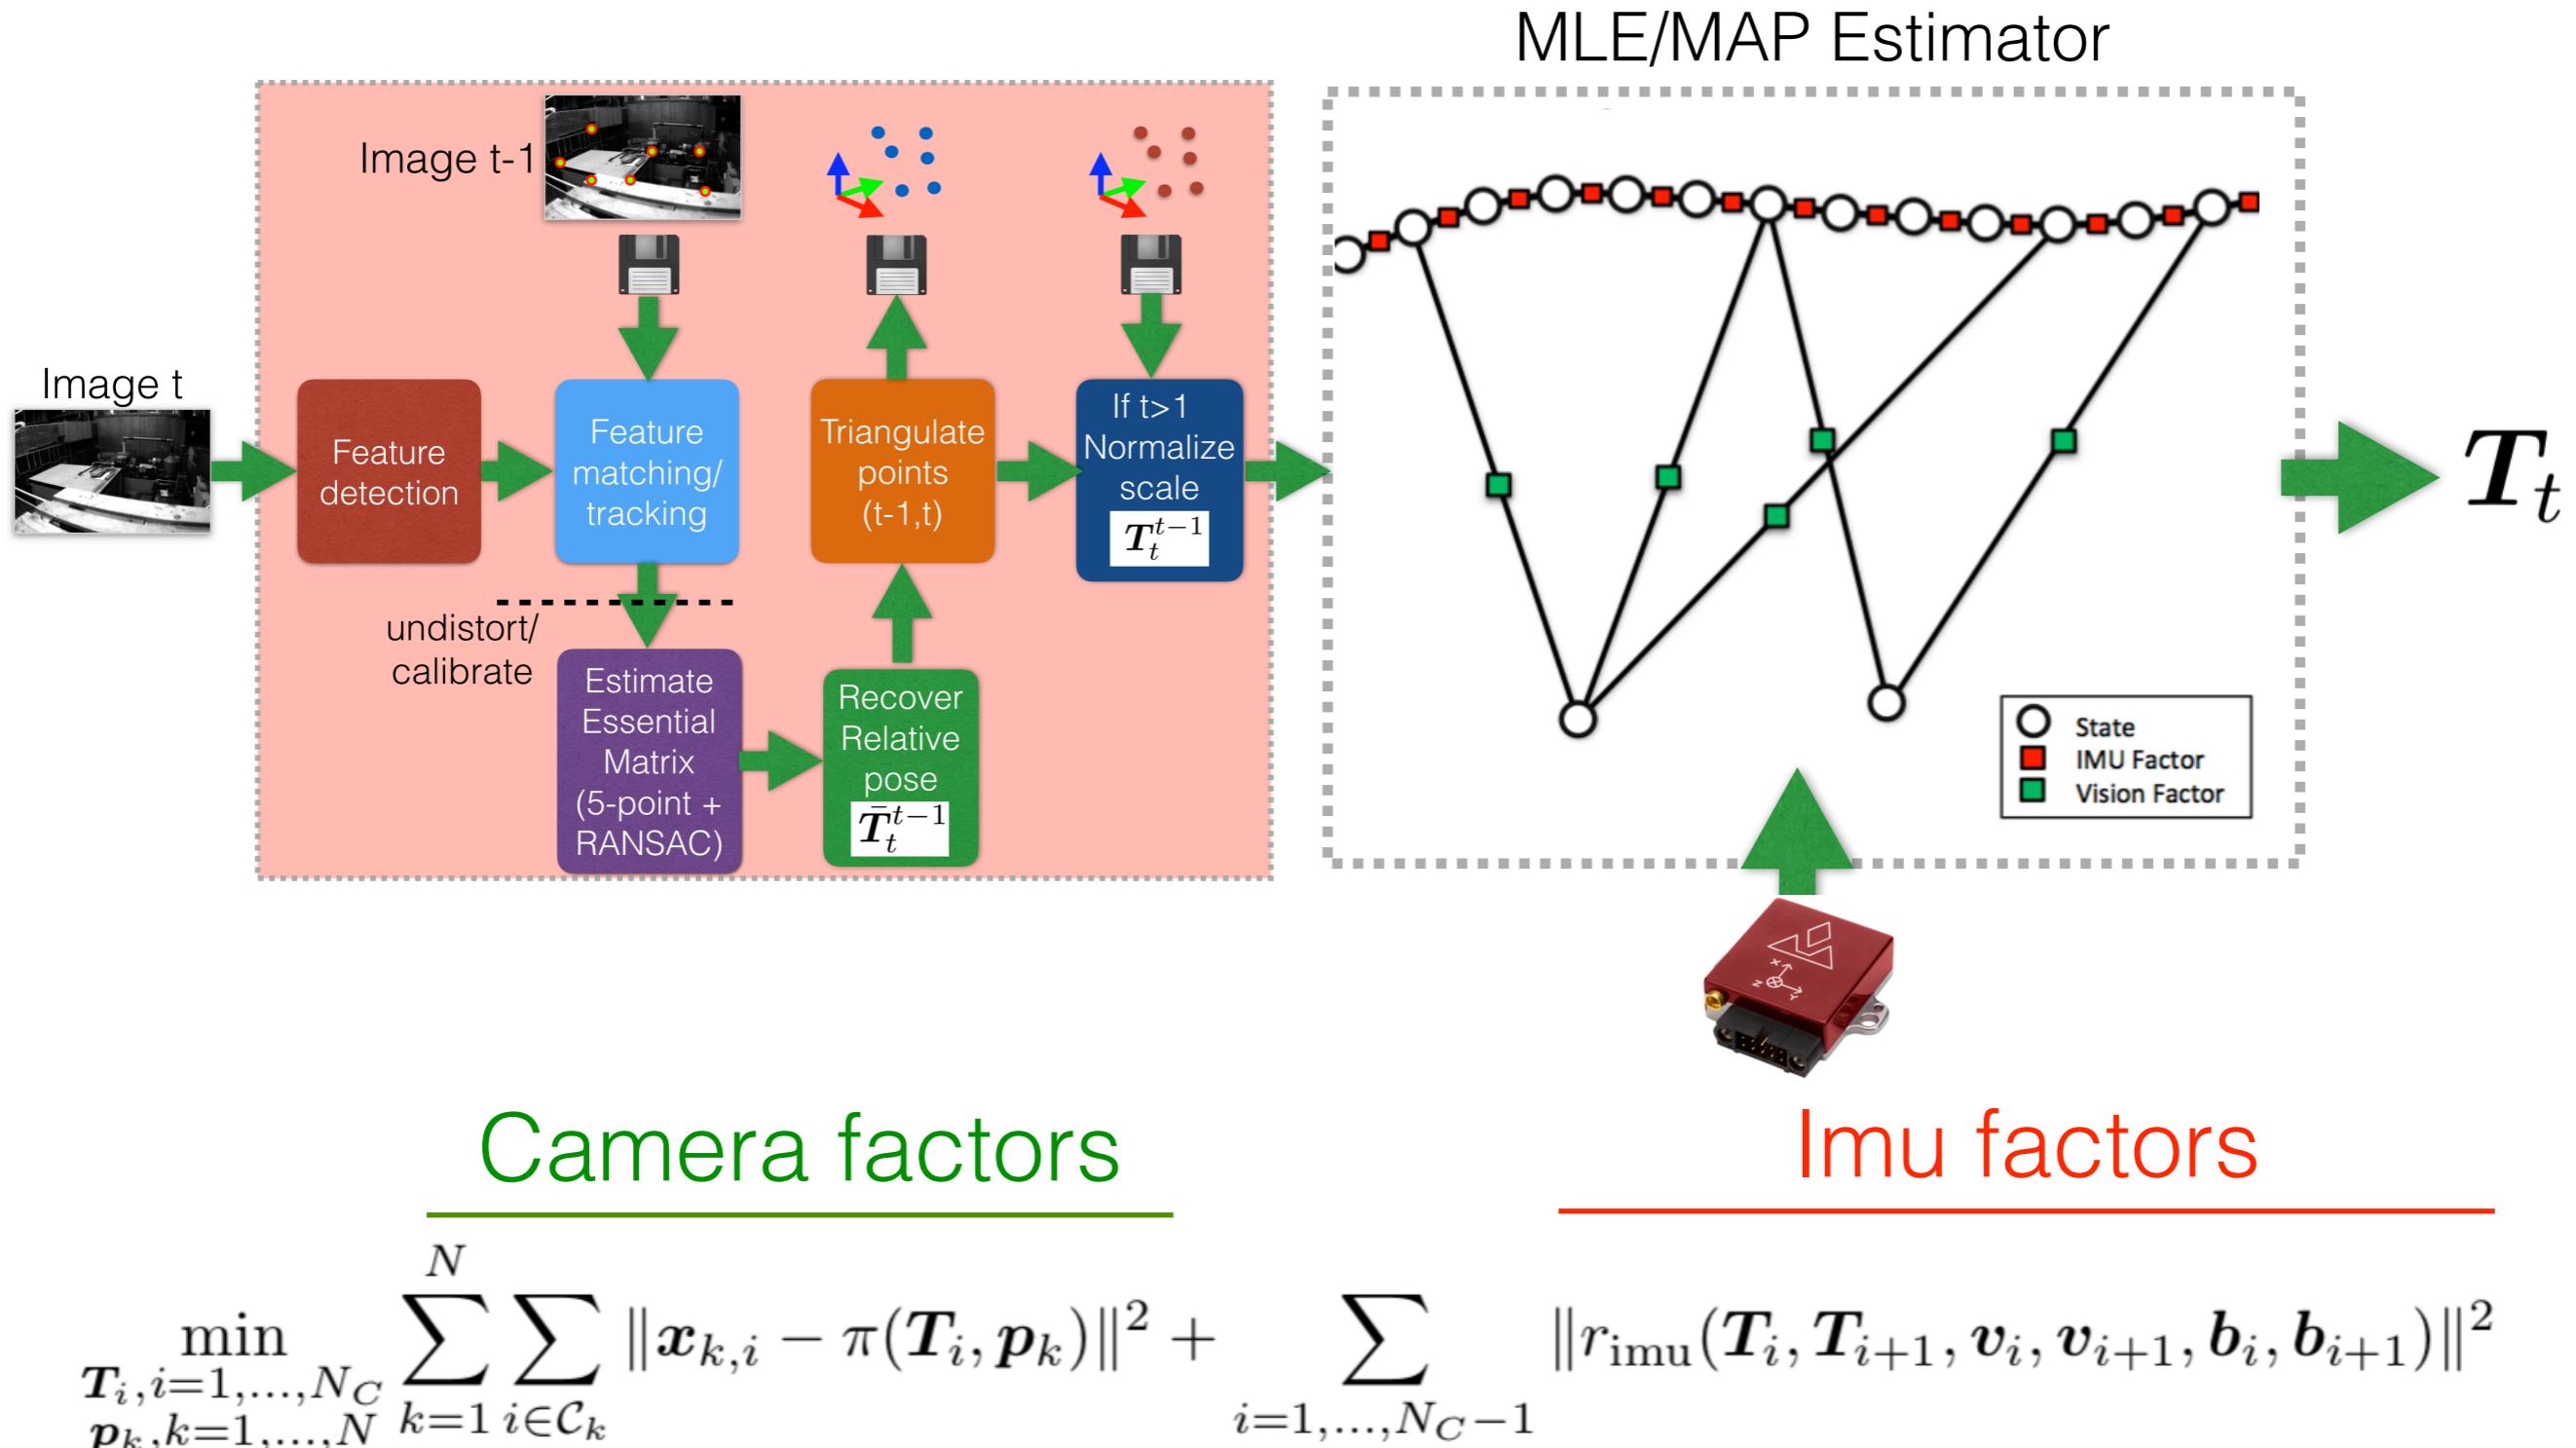


830g	160g	4g	3g
8 W	2.5 W	0.3W	~1 W

Visual-Inertial Navigation (VIN)

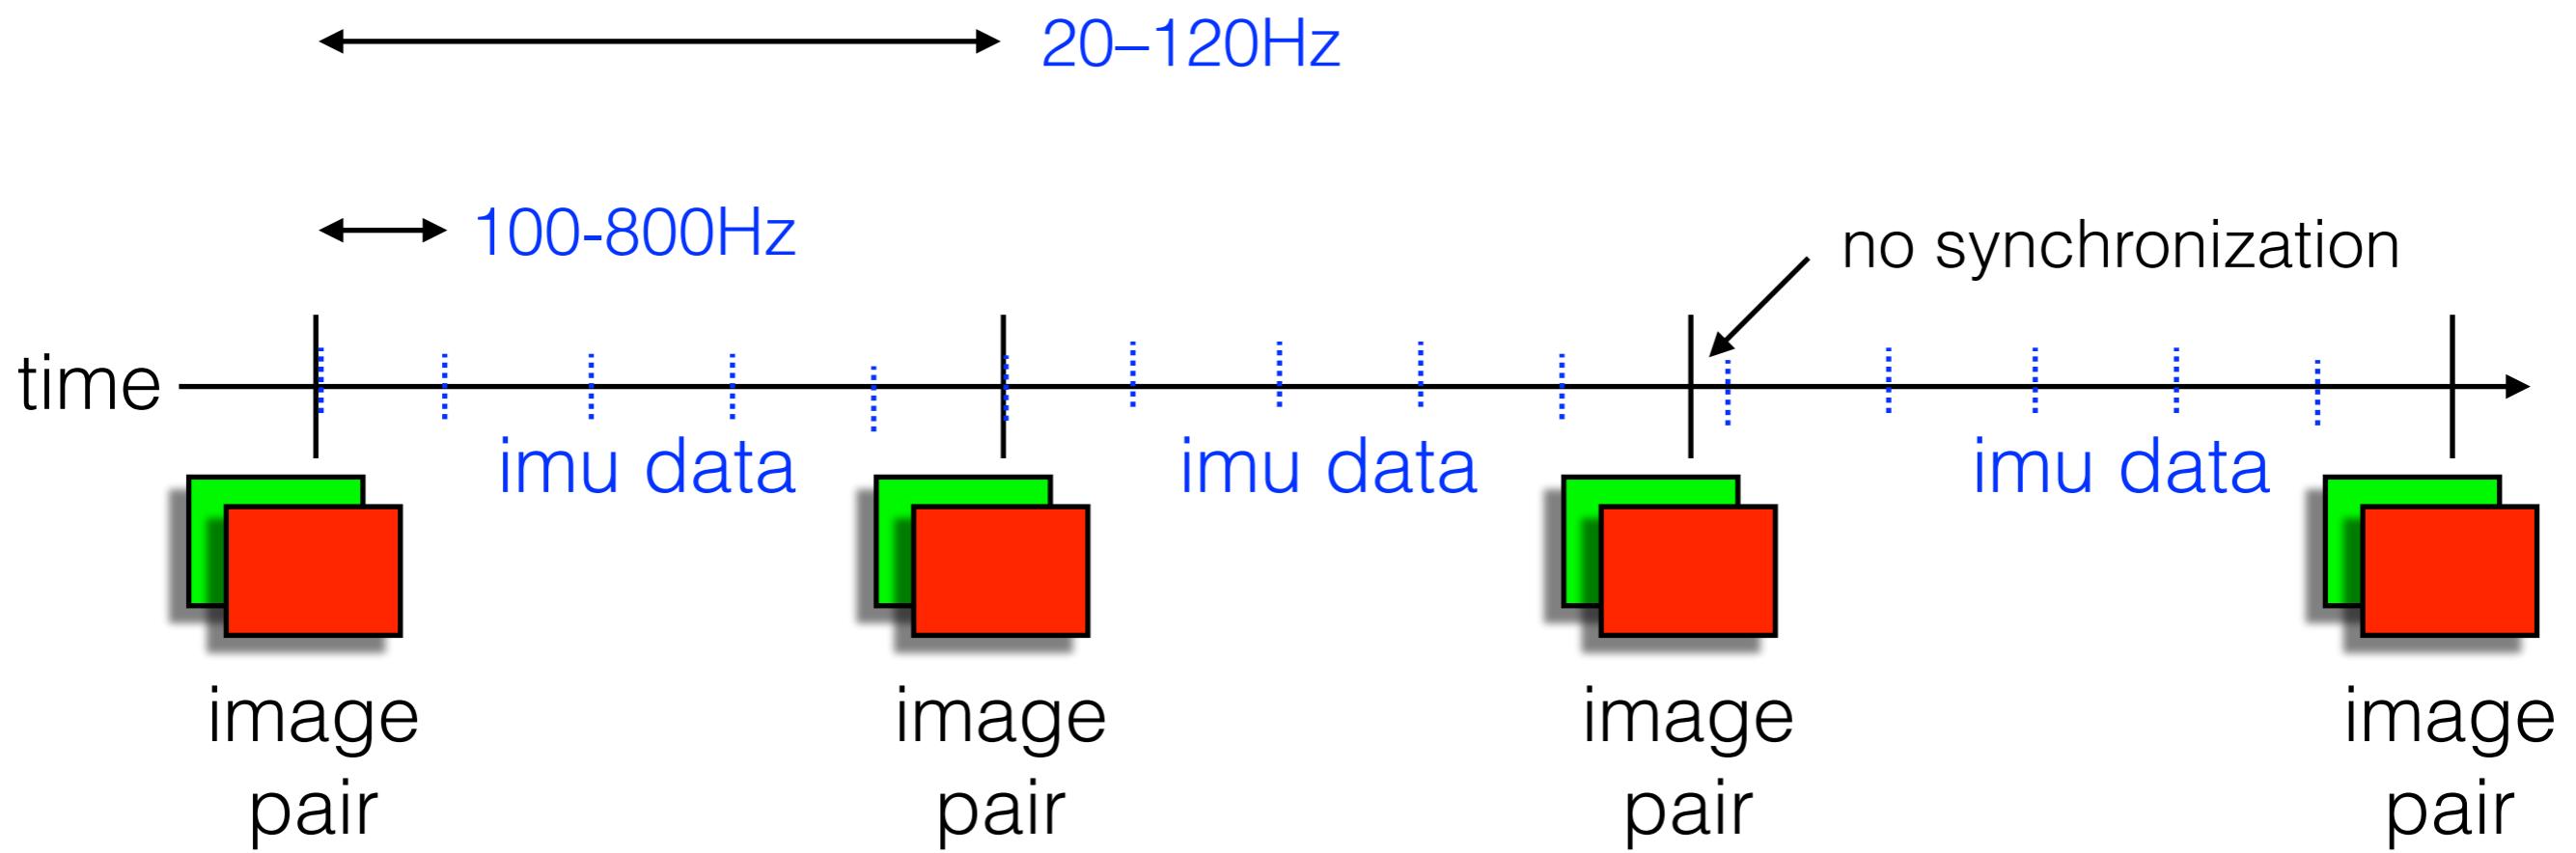


Visual-Inertial Odometry



Need to include velocities and IMU biases in the state ...

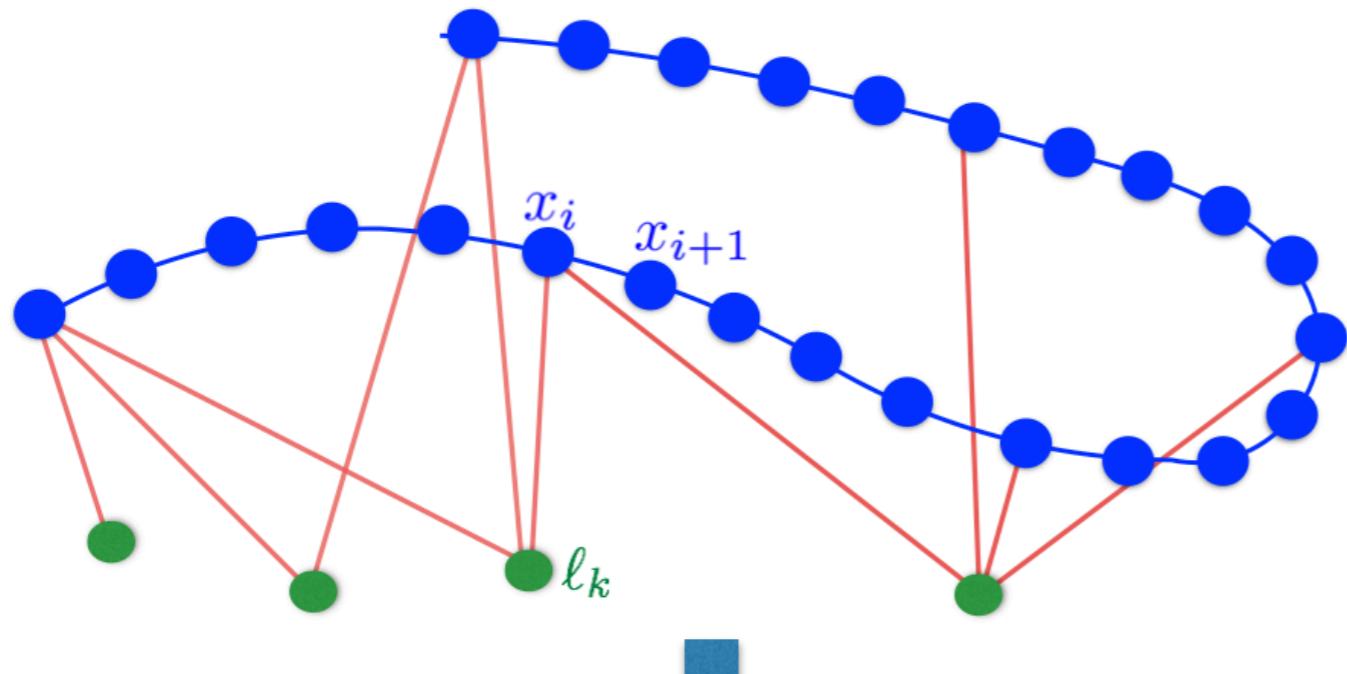
Visual-Inertial Odometry



Challenges:

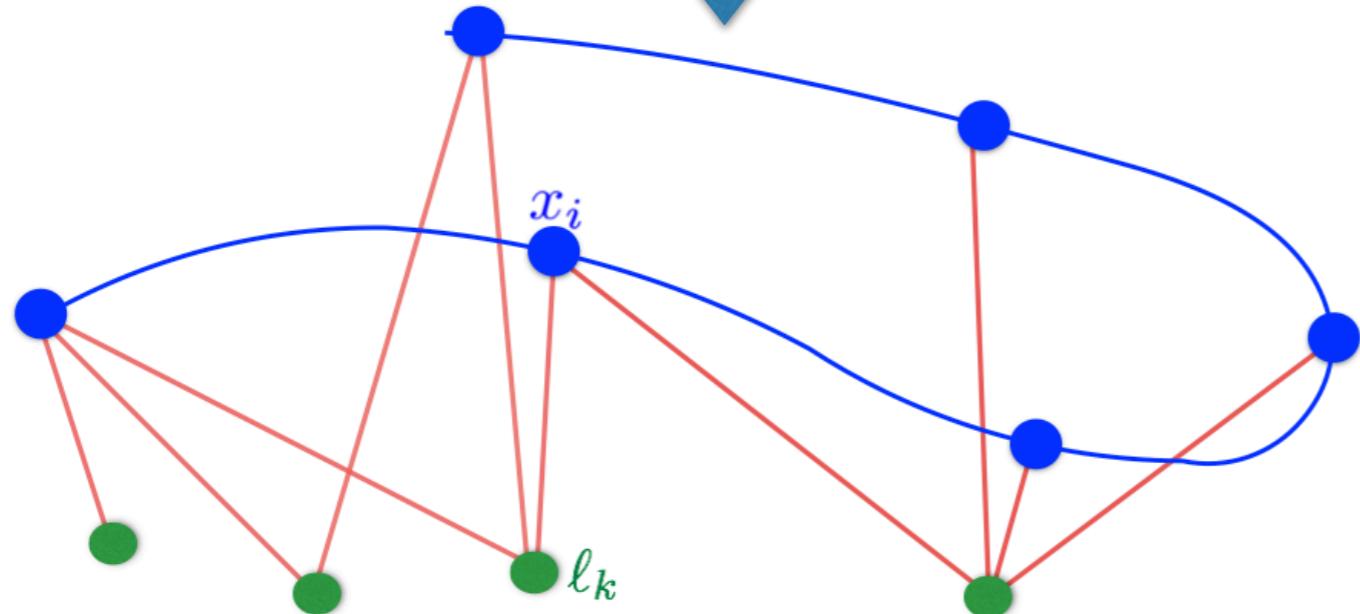
- IMU measurements arrive at high-rate (~200Hz)  **IMU preintegration**
- camera observes hundreds of landmarks per frame  **structureless vision factors**
- need to solve optimization problem quickly

Pre-integration



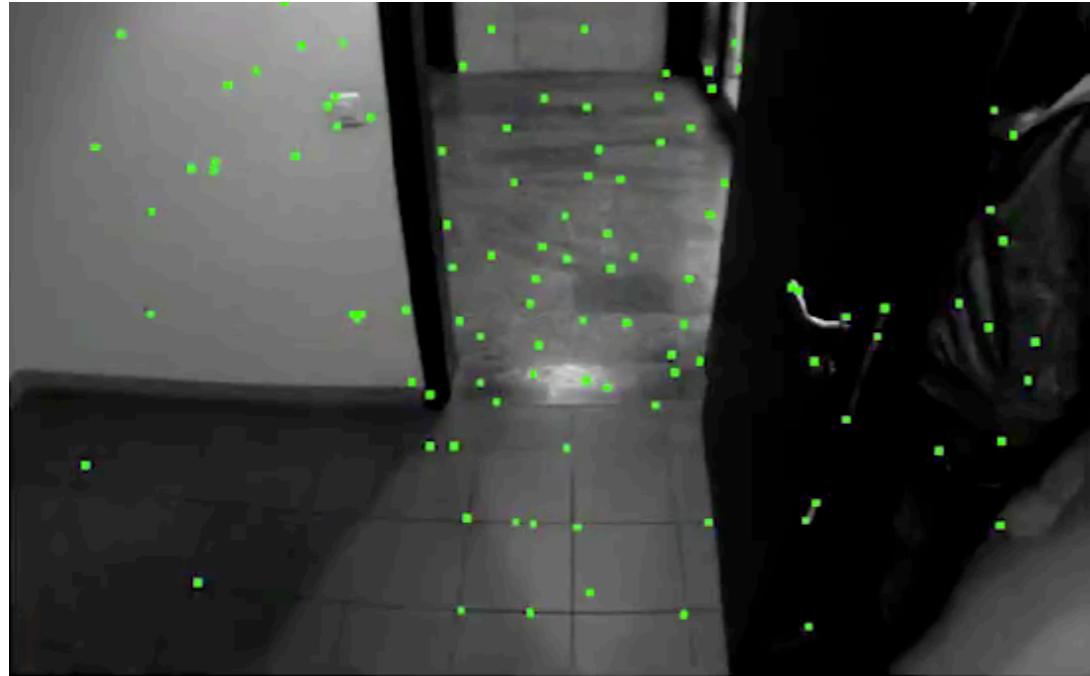
After 10 seconds, original problem has $\sim 10^4$ states

Preintegration

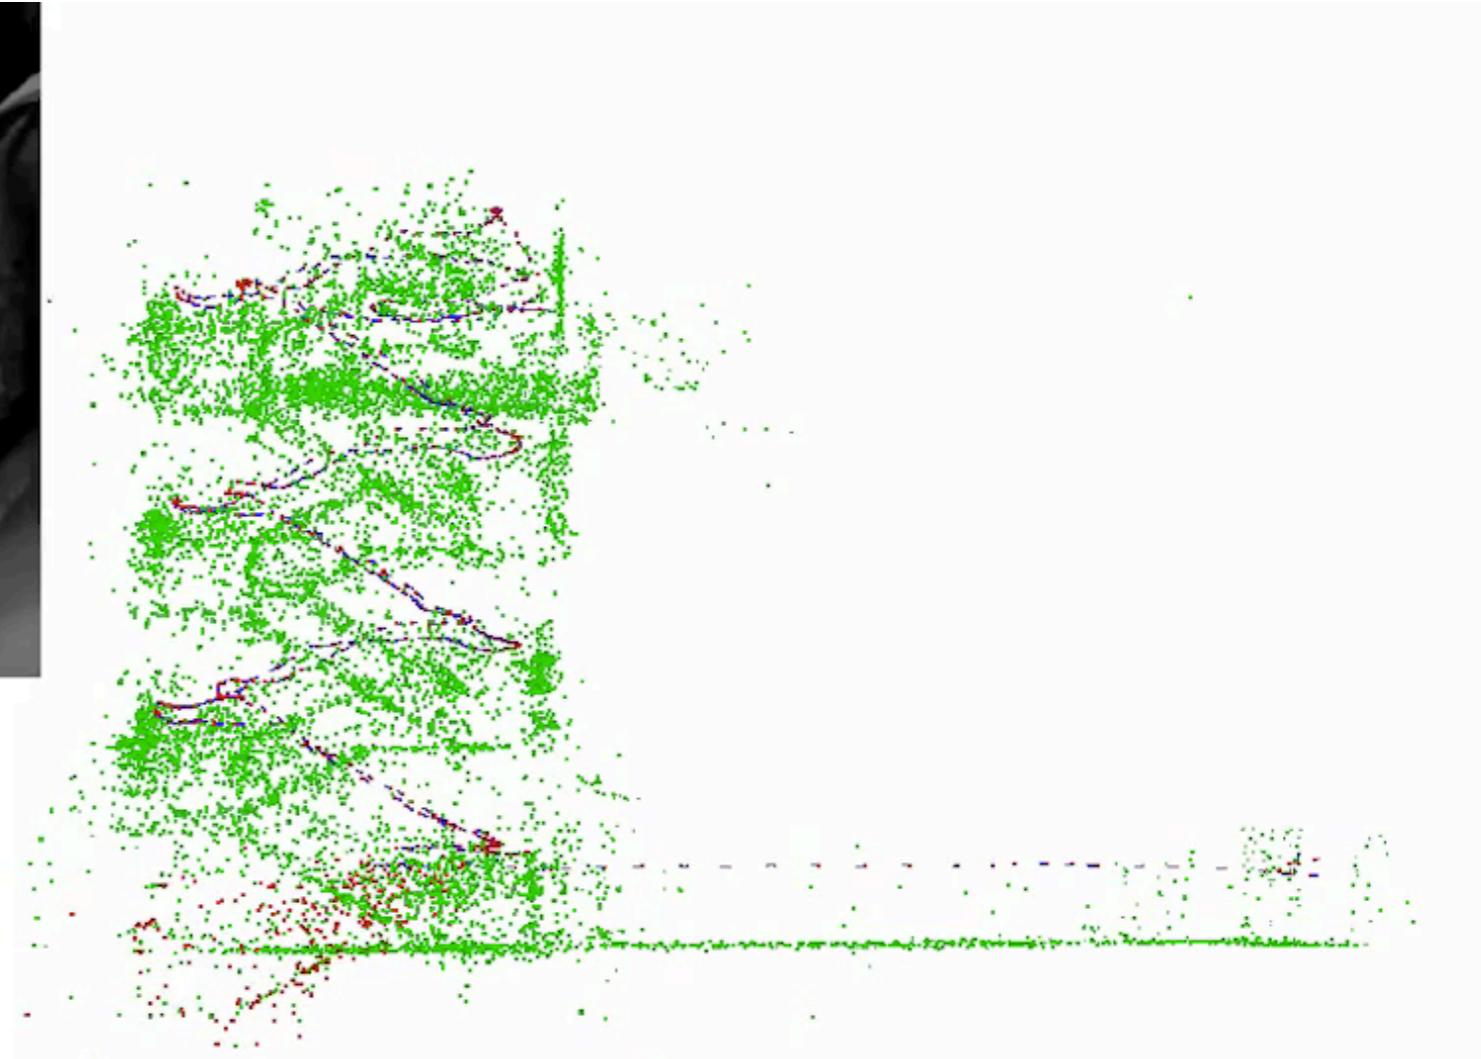


After 10 seconds,
preintegrated problem
has $\sim 10^2$ states

Visual-Inertial Odometry

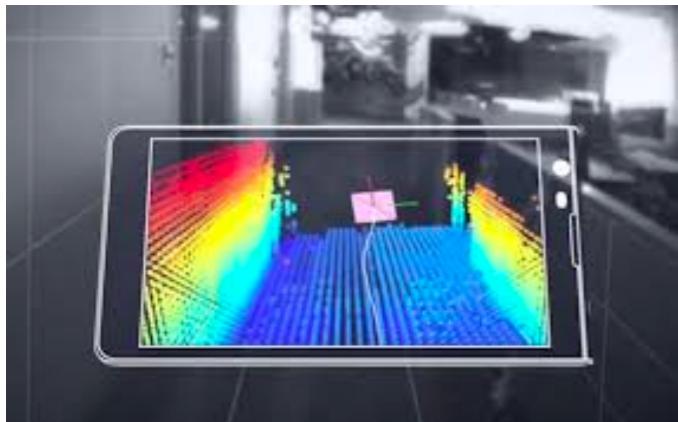


Hand-held
sensor



**Implemented
in GTSAM
(ImuFactor)**

Recent Implementations / Products



2014

Reinvented as
ARCore in 2017



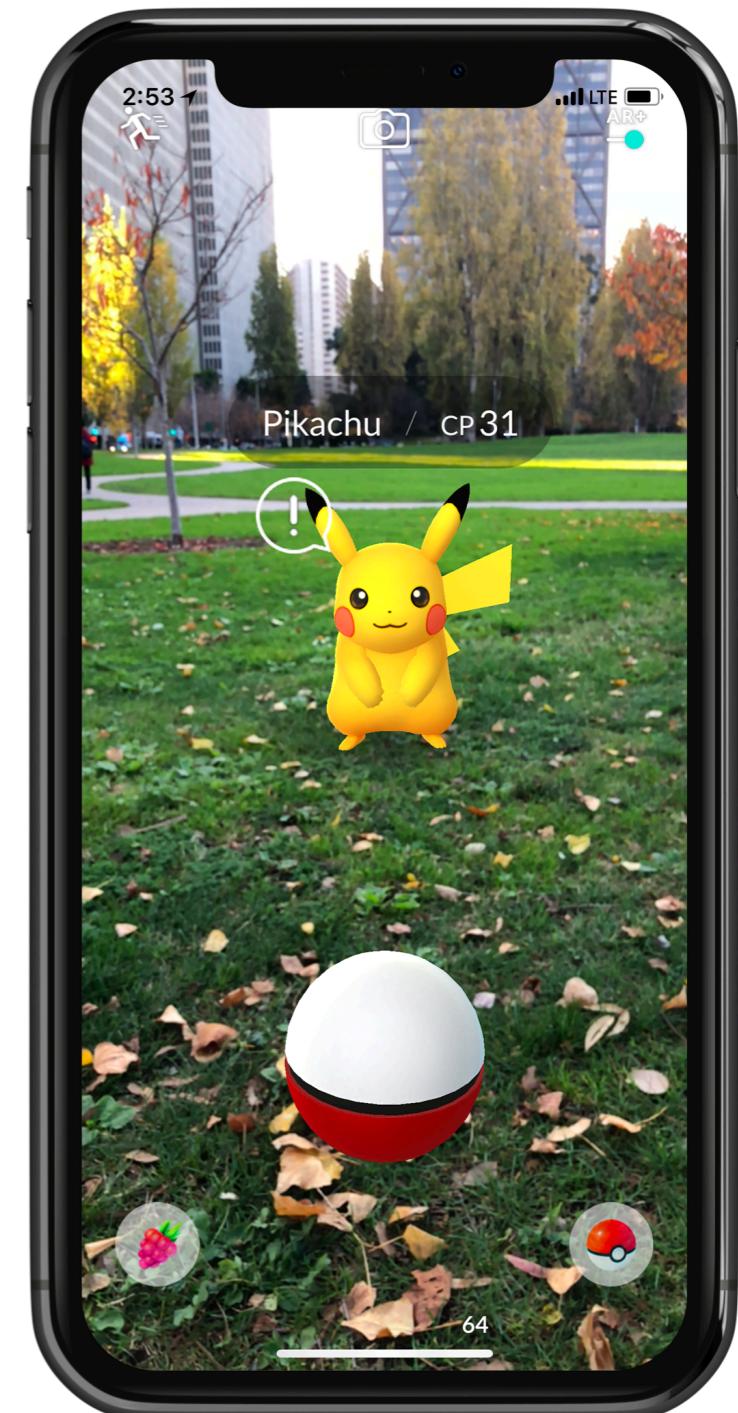
Oculus Rift

Announced in 2012.
Acquired by
Facebook in 2014

Navion Chip
2017

(<http://navion.mit.edu/>)

Pokemon Go



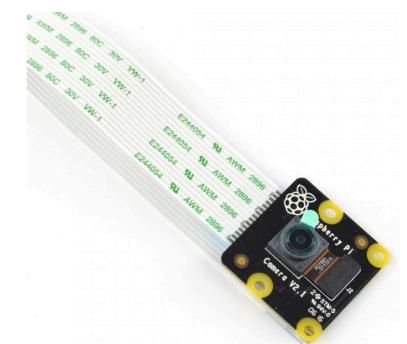
Beyond VO

How to get scale and improve robustness?

add more sensors!

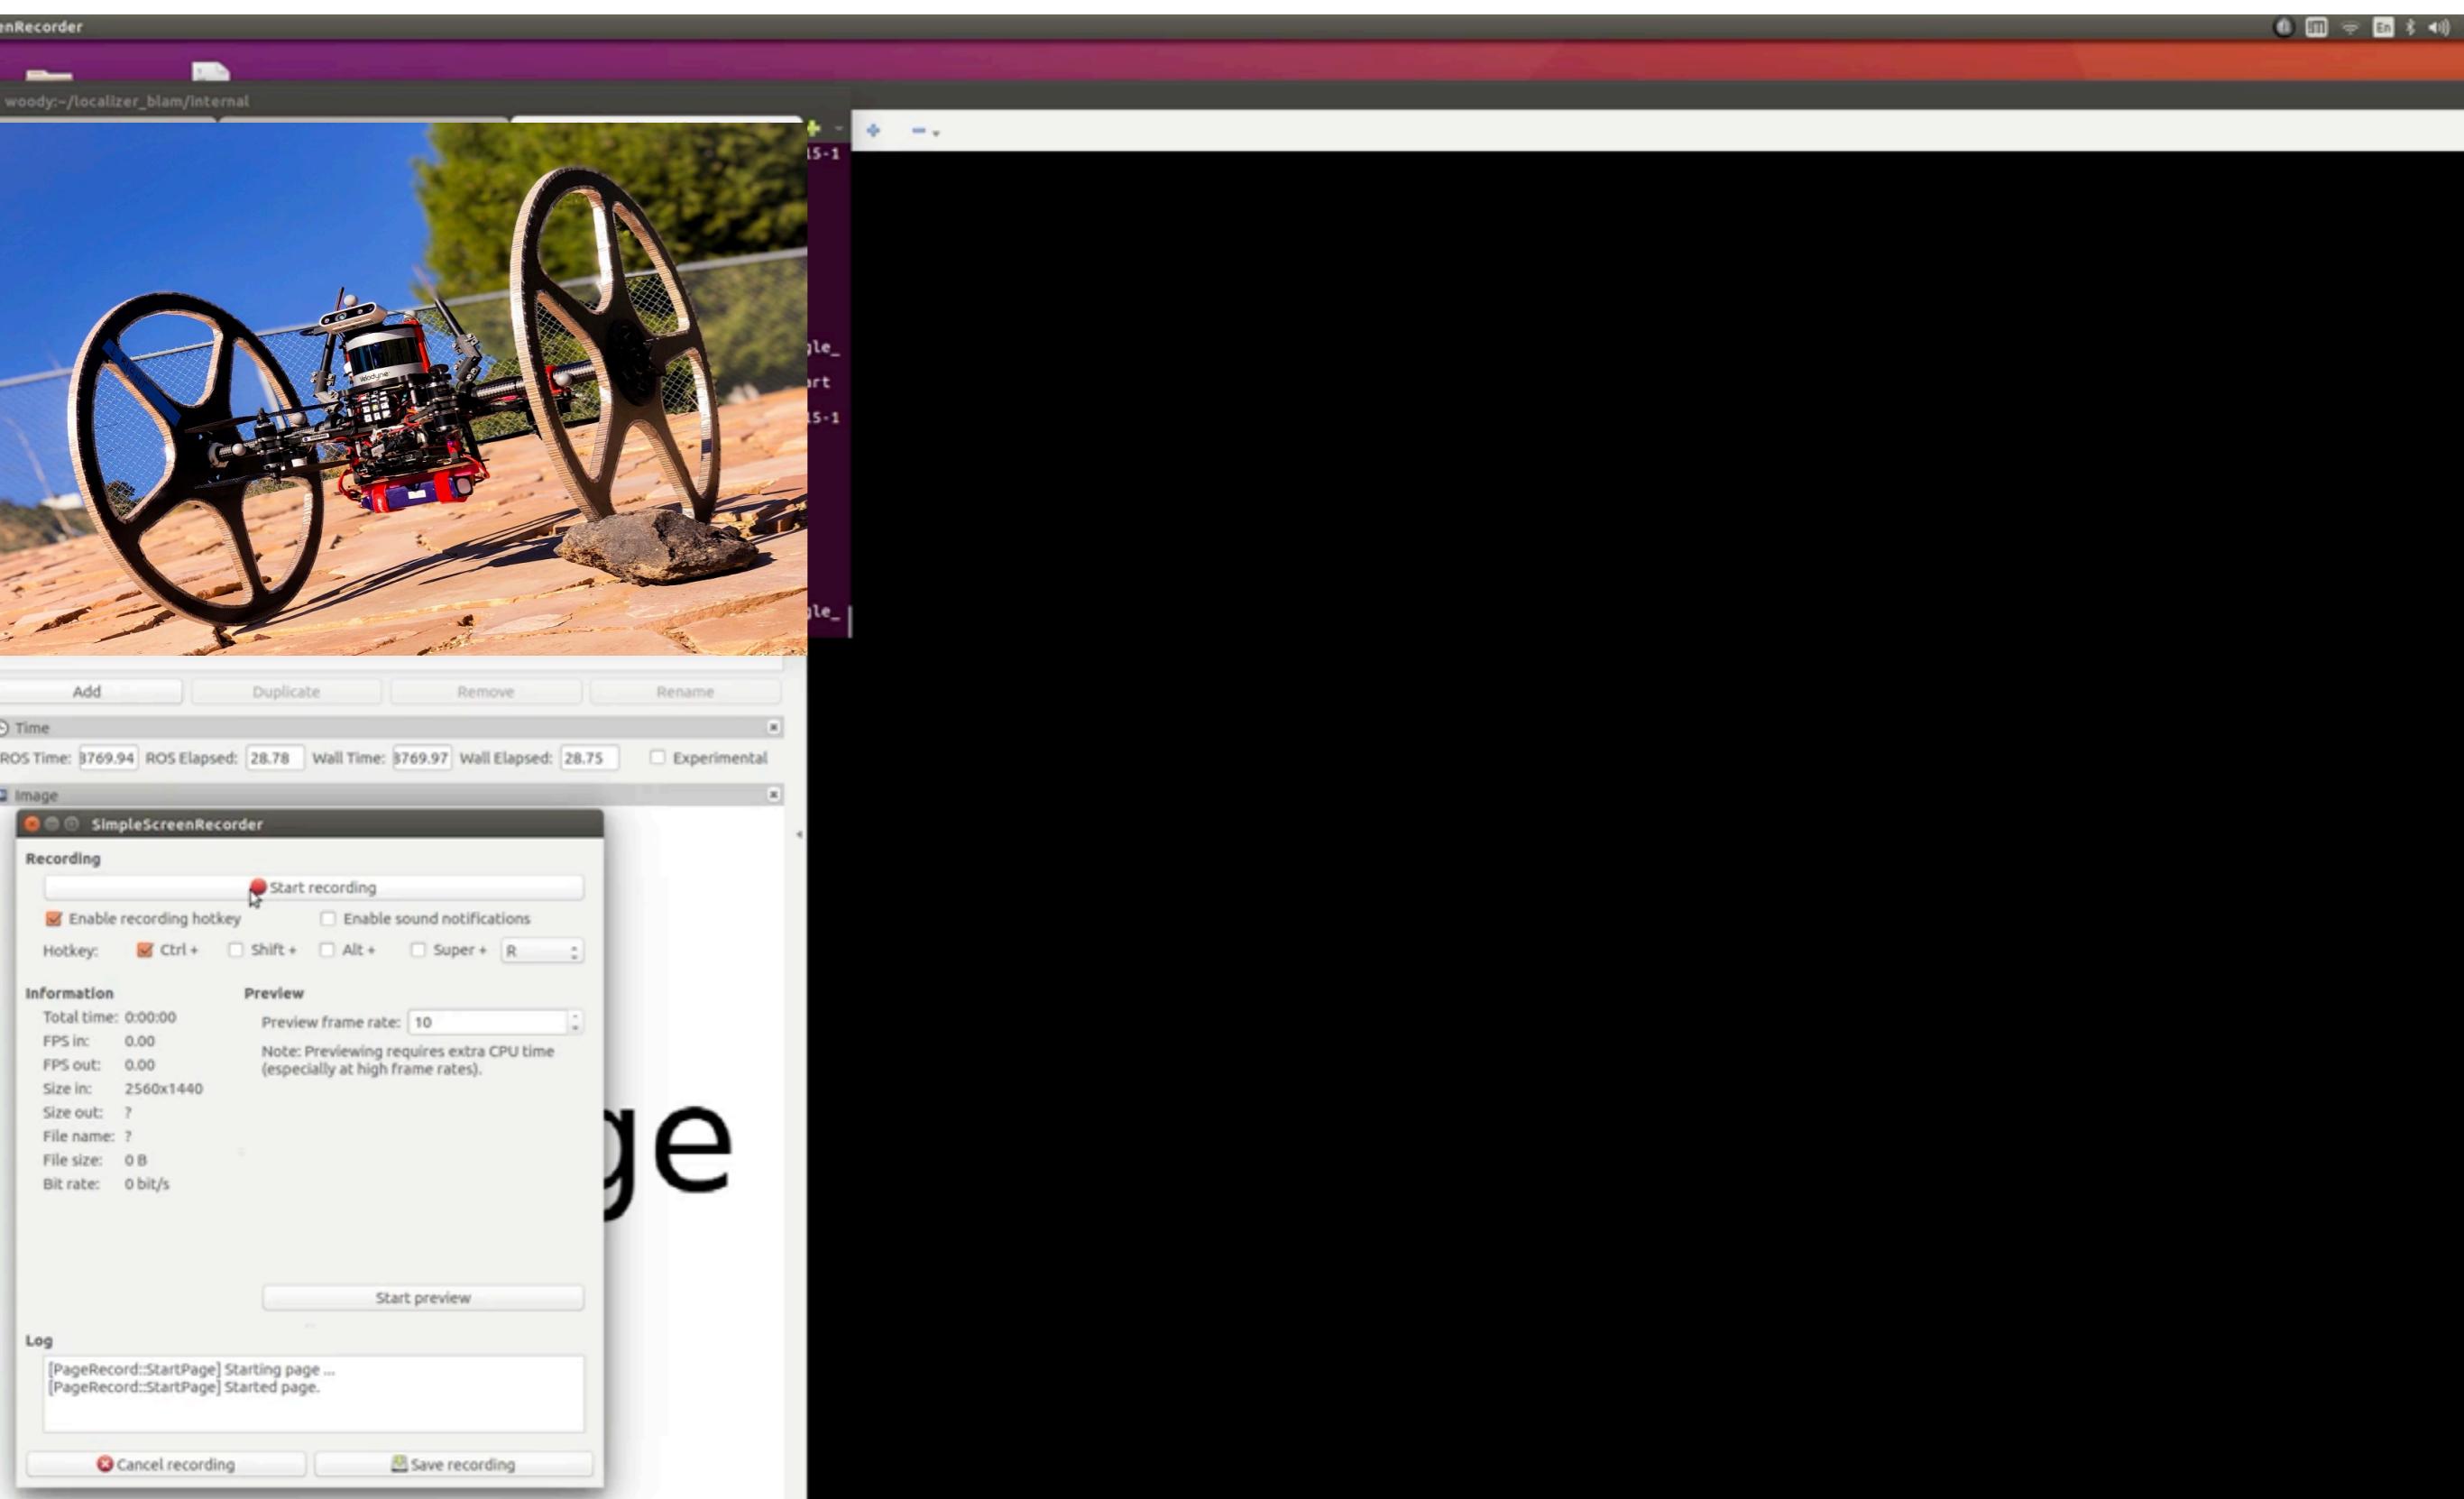
- ▶ wheel odometry
- ▶ GPS
- ▶ **Lidar**
- ▶ Inertial

Measurement
Unit (IMU)



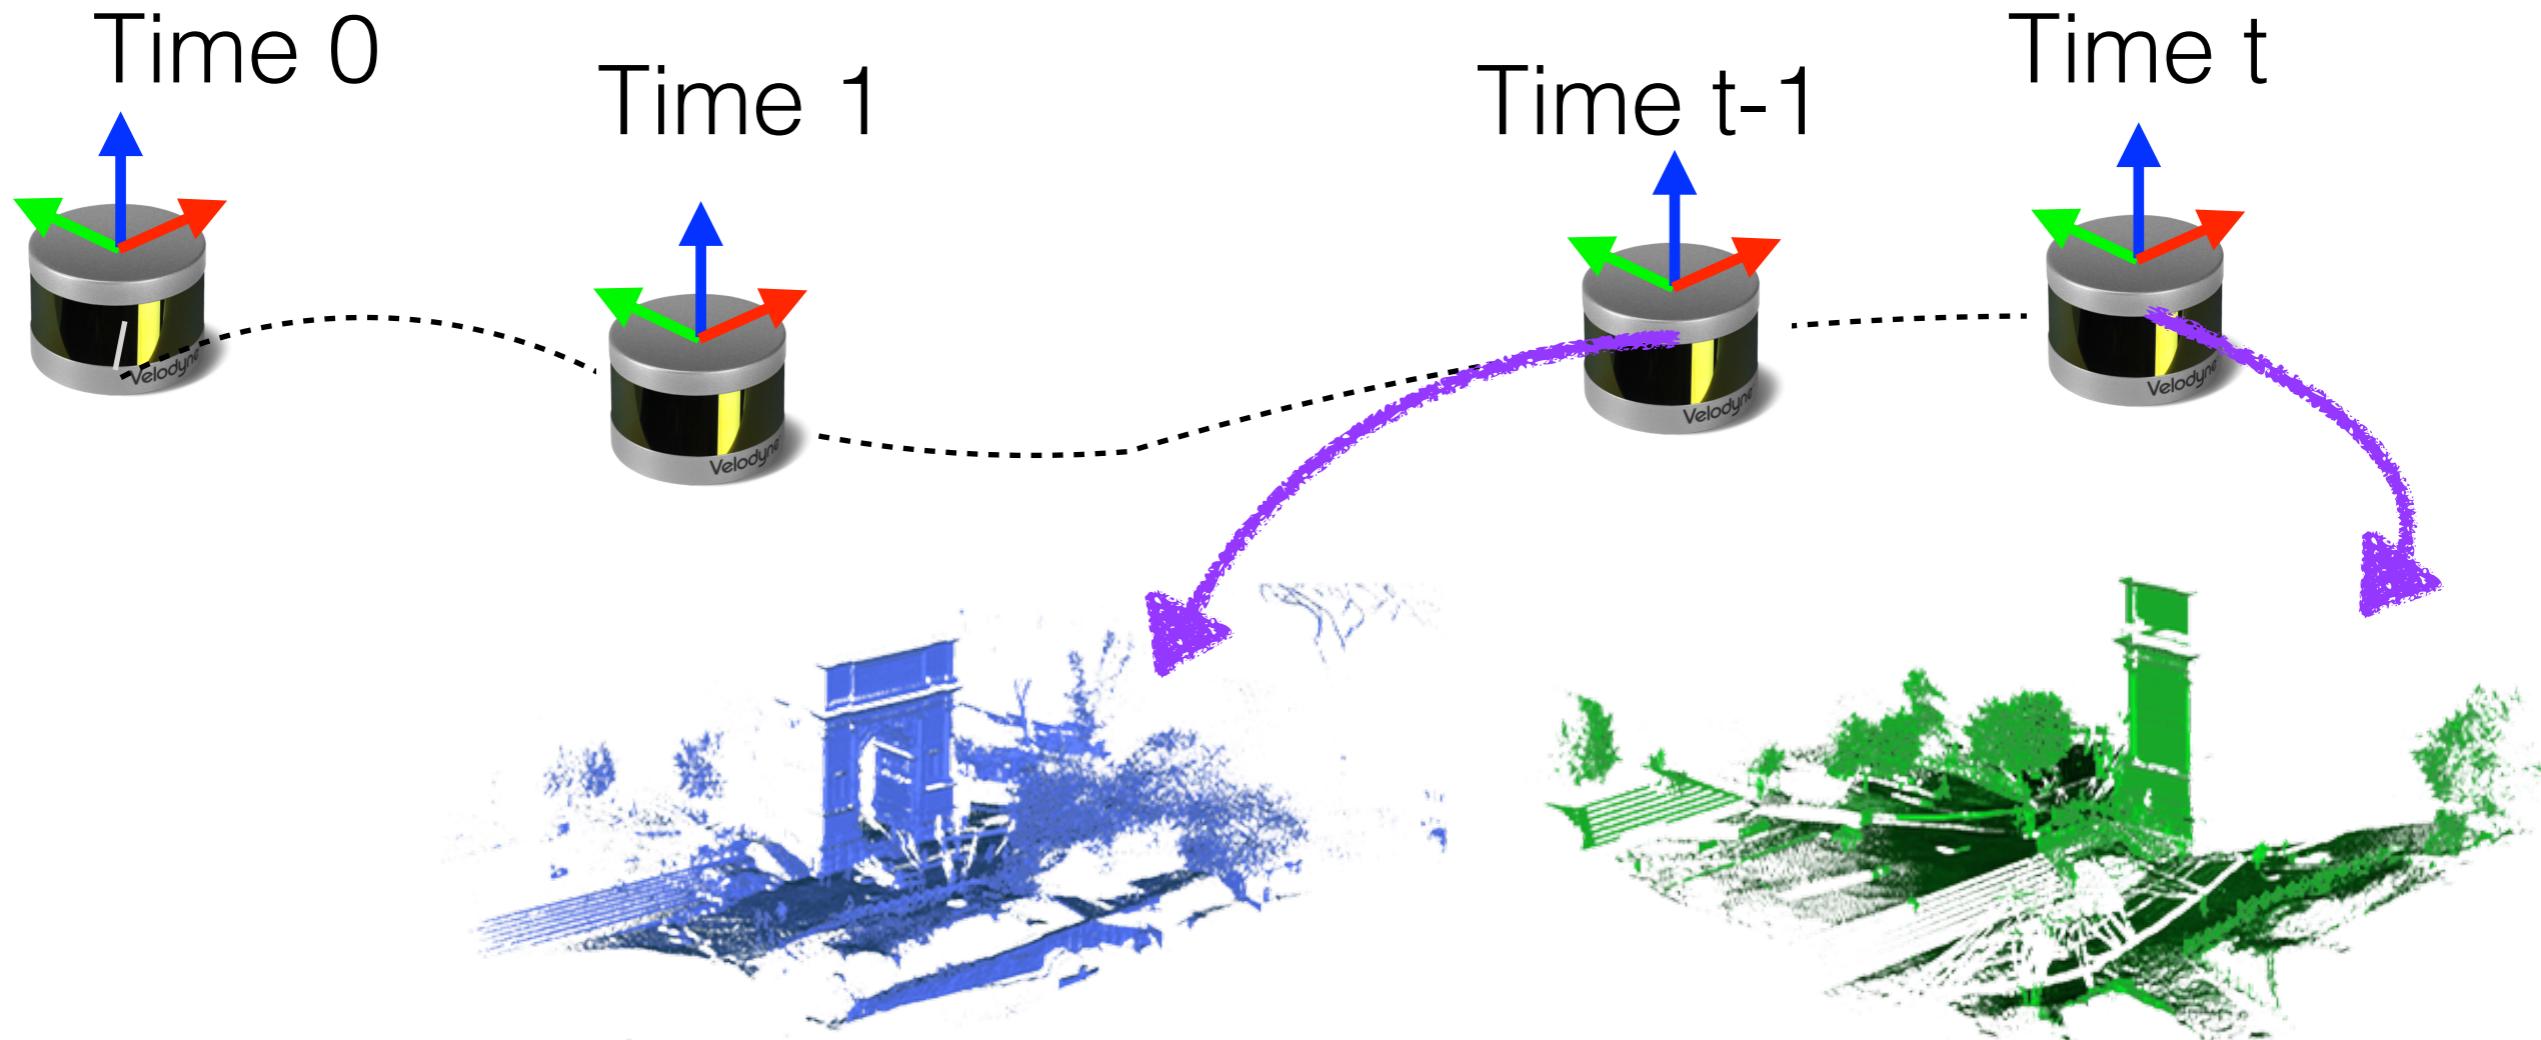
830g	160g	4g	3g
8 W	2.5 W	0.3W	~1 W

Lidar Odometry & Lidar SLAM

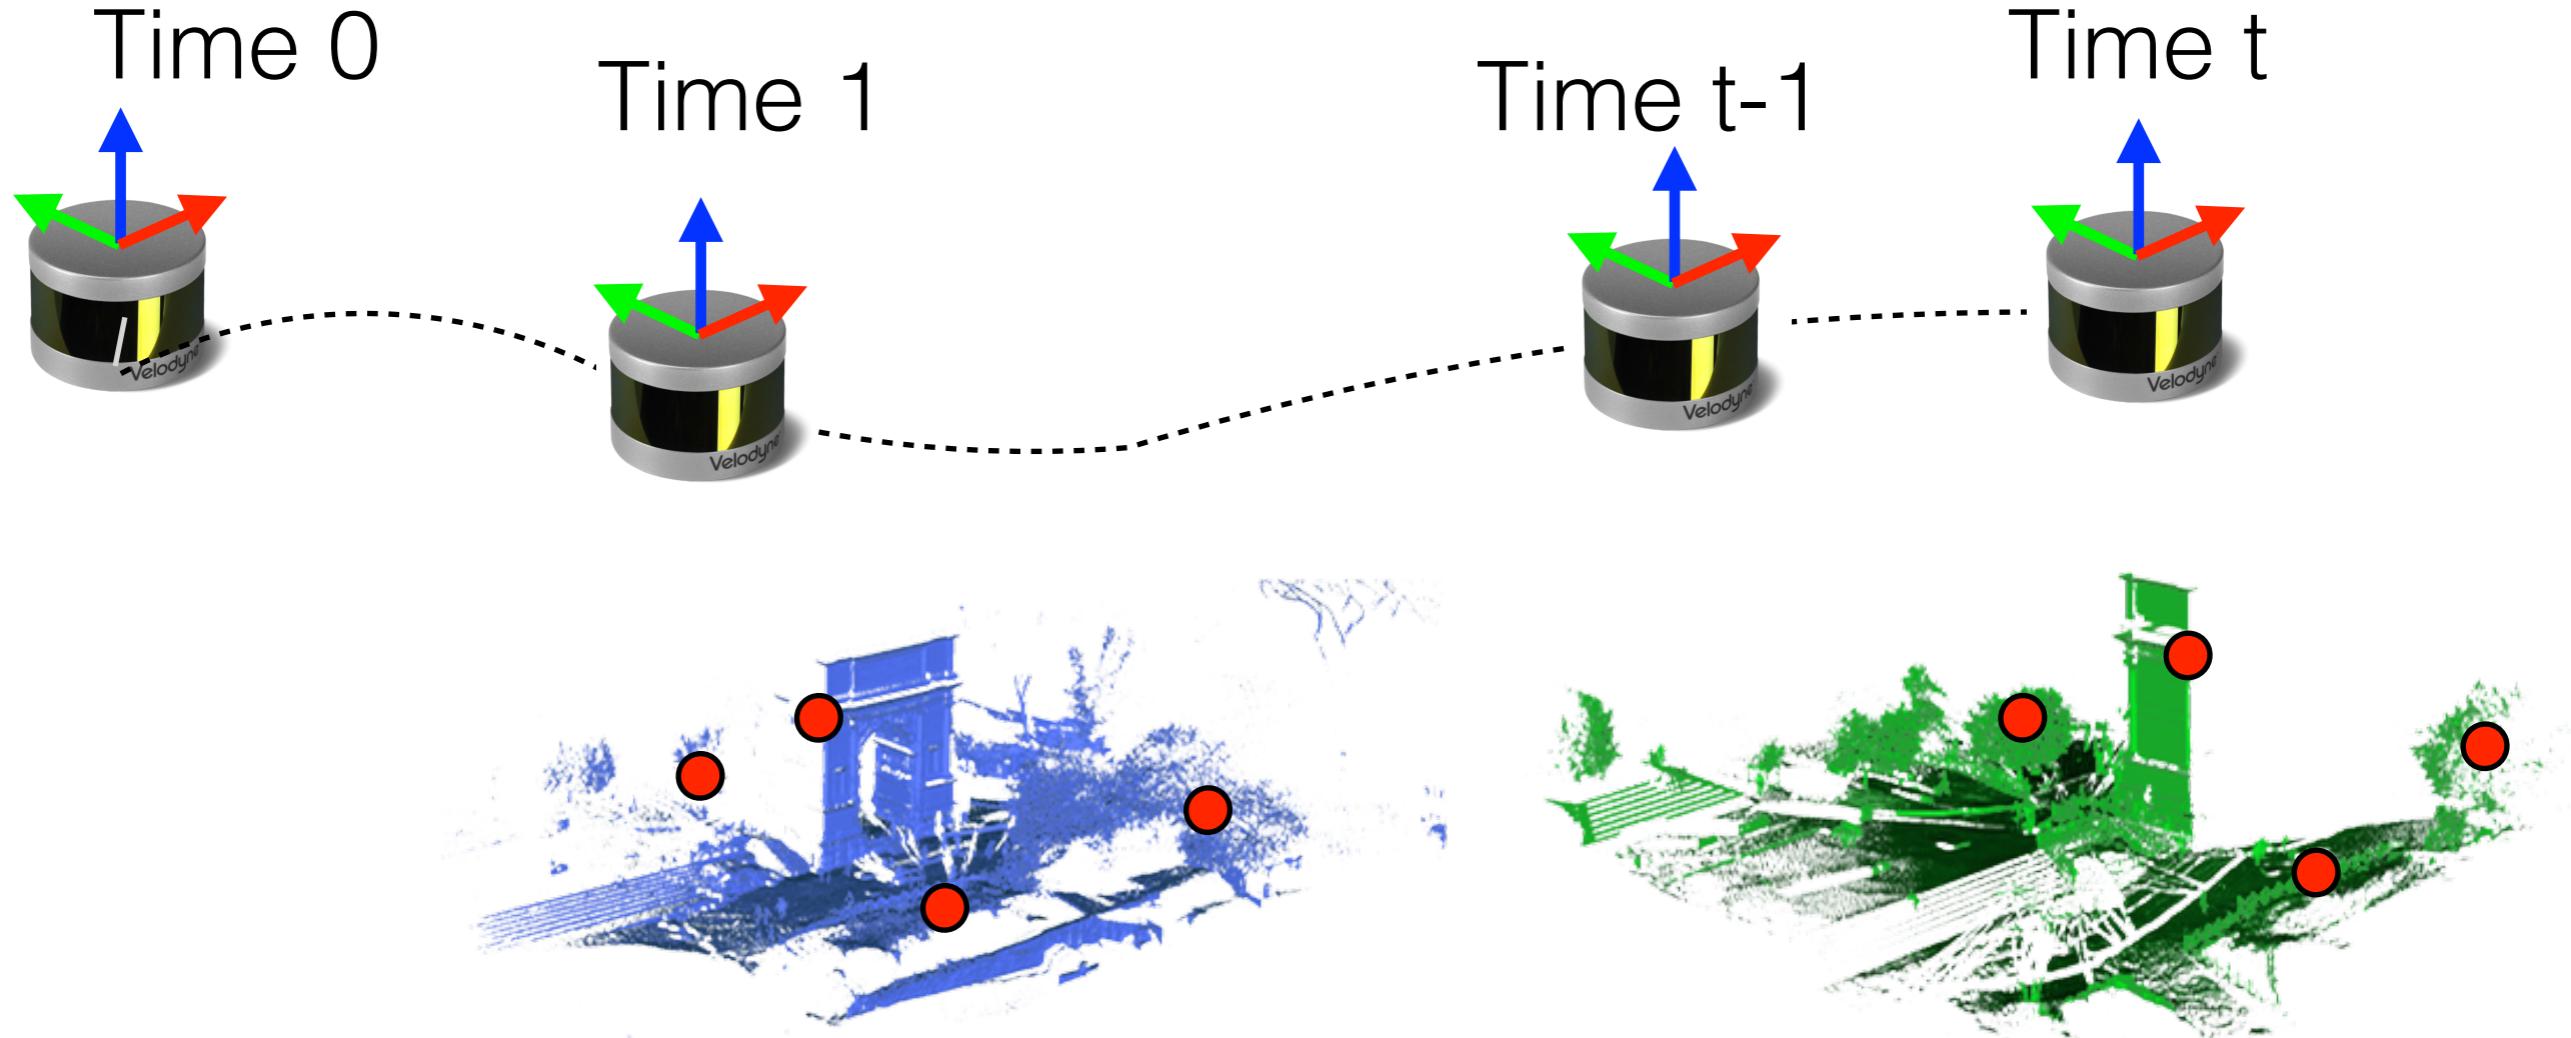


DARPA Subterranean Challenge, in collaboration with JPL

Feature-based Lidar Odometry



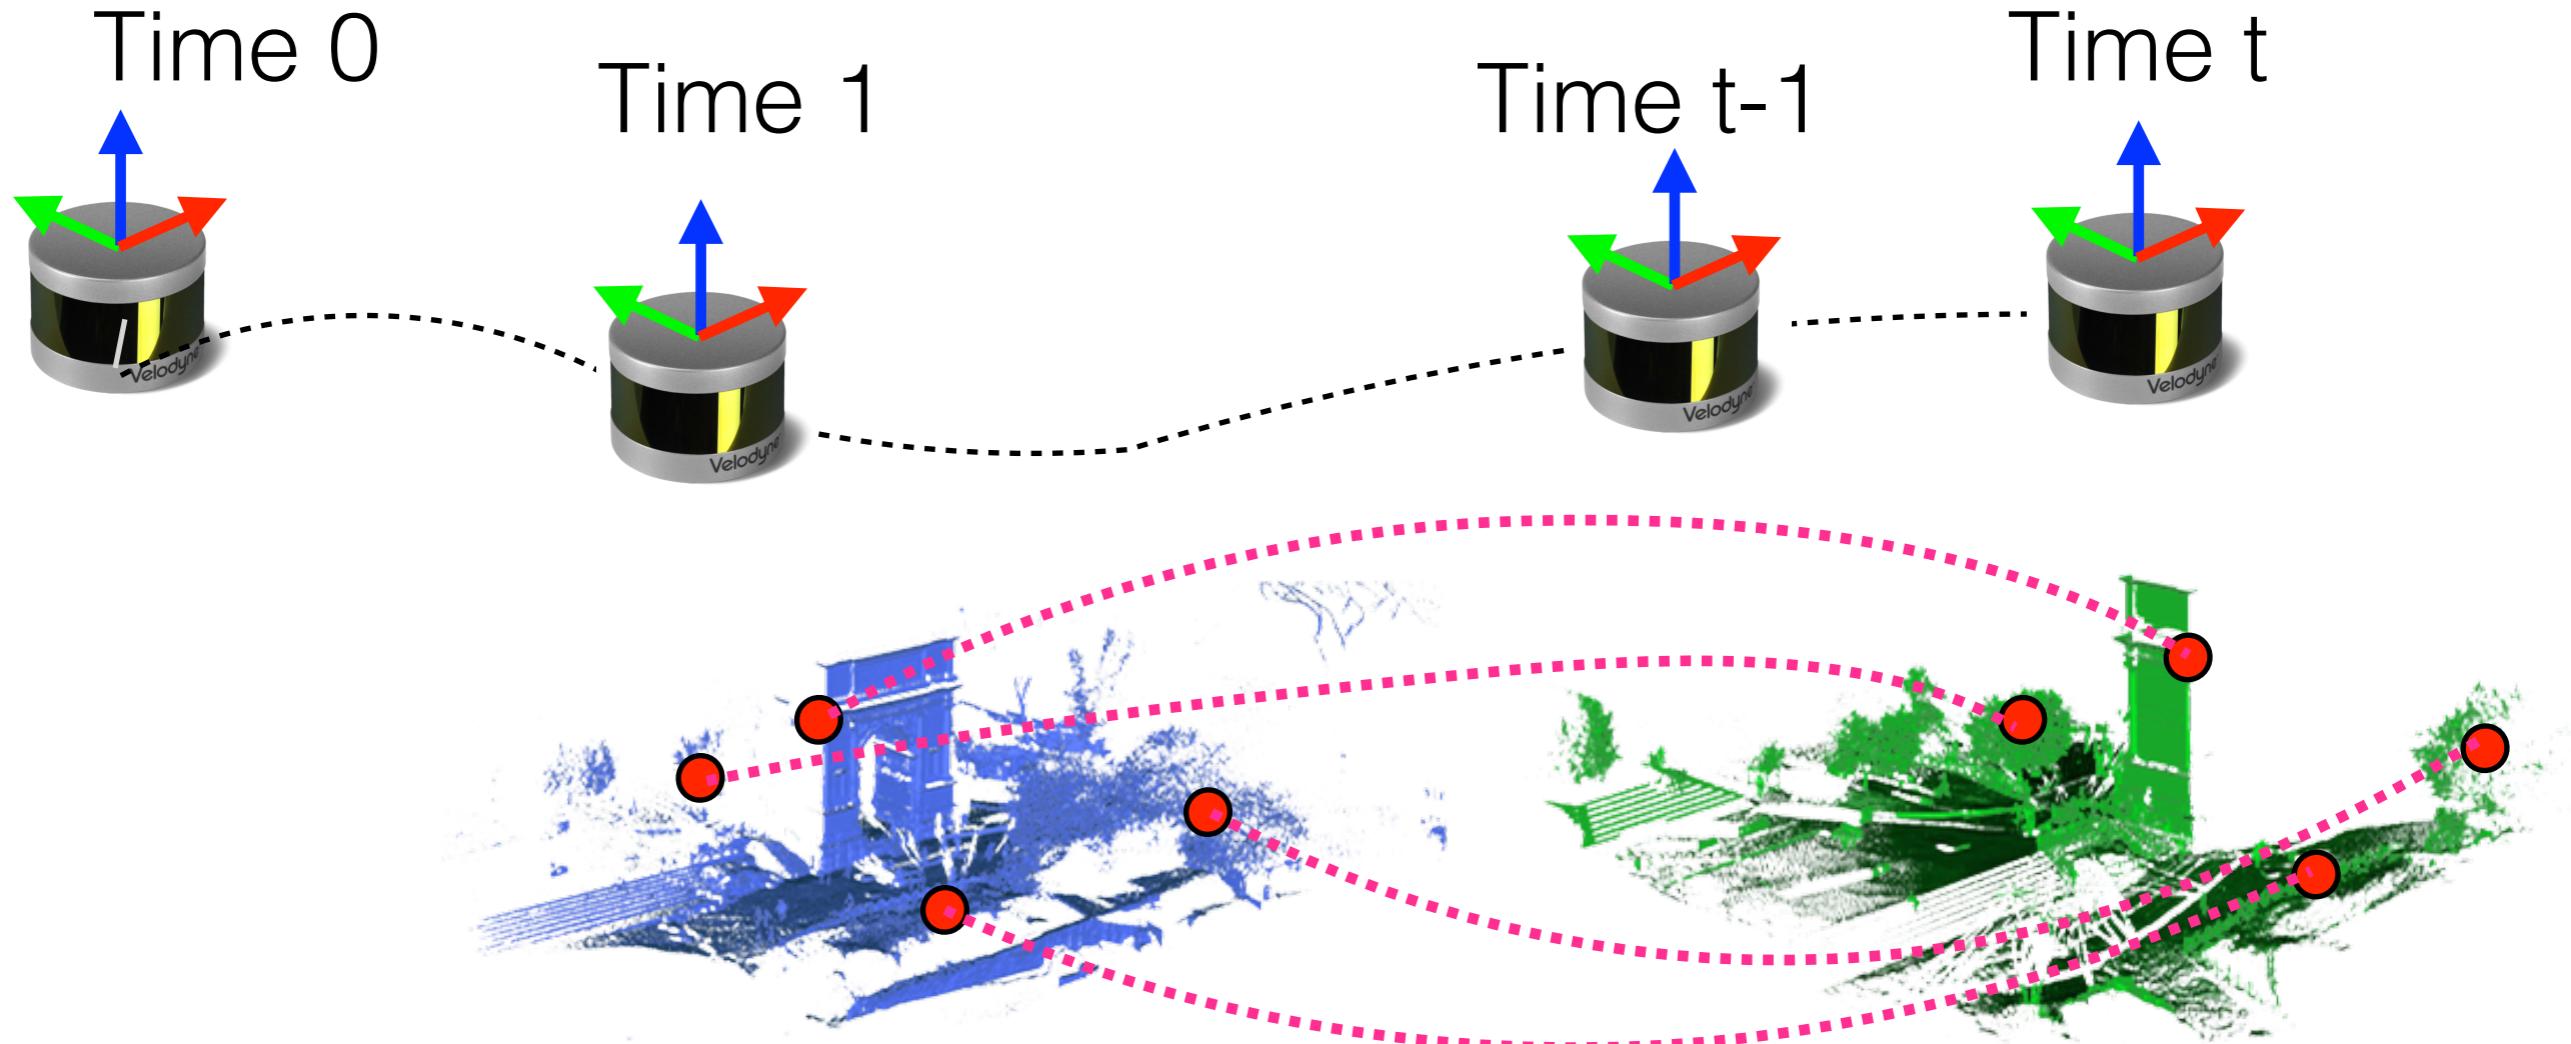
Feature-based Lidar Odometry



Registration: compute relative pose between scans:

- extract features & descriptors
- use descriptors for matching
- compute relative pose

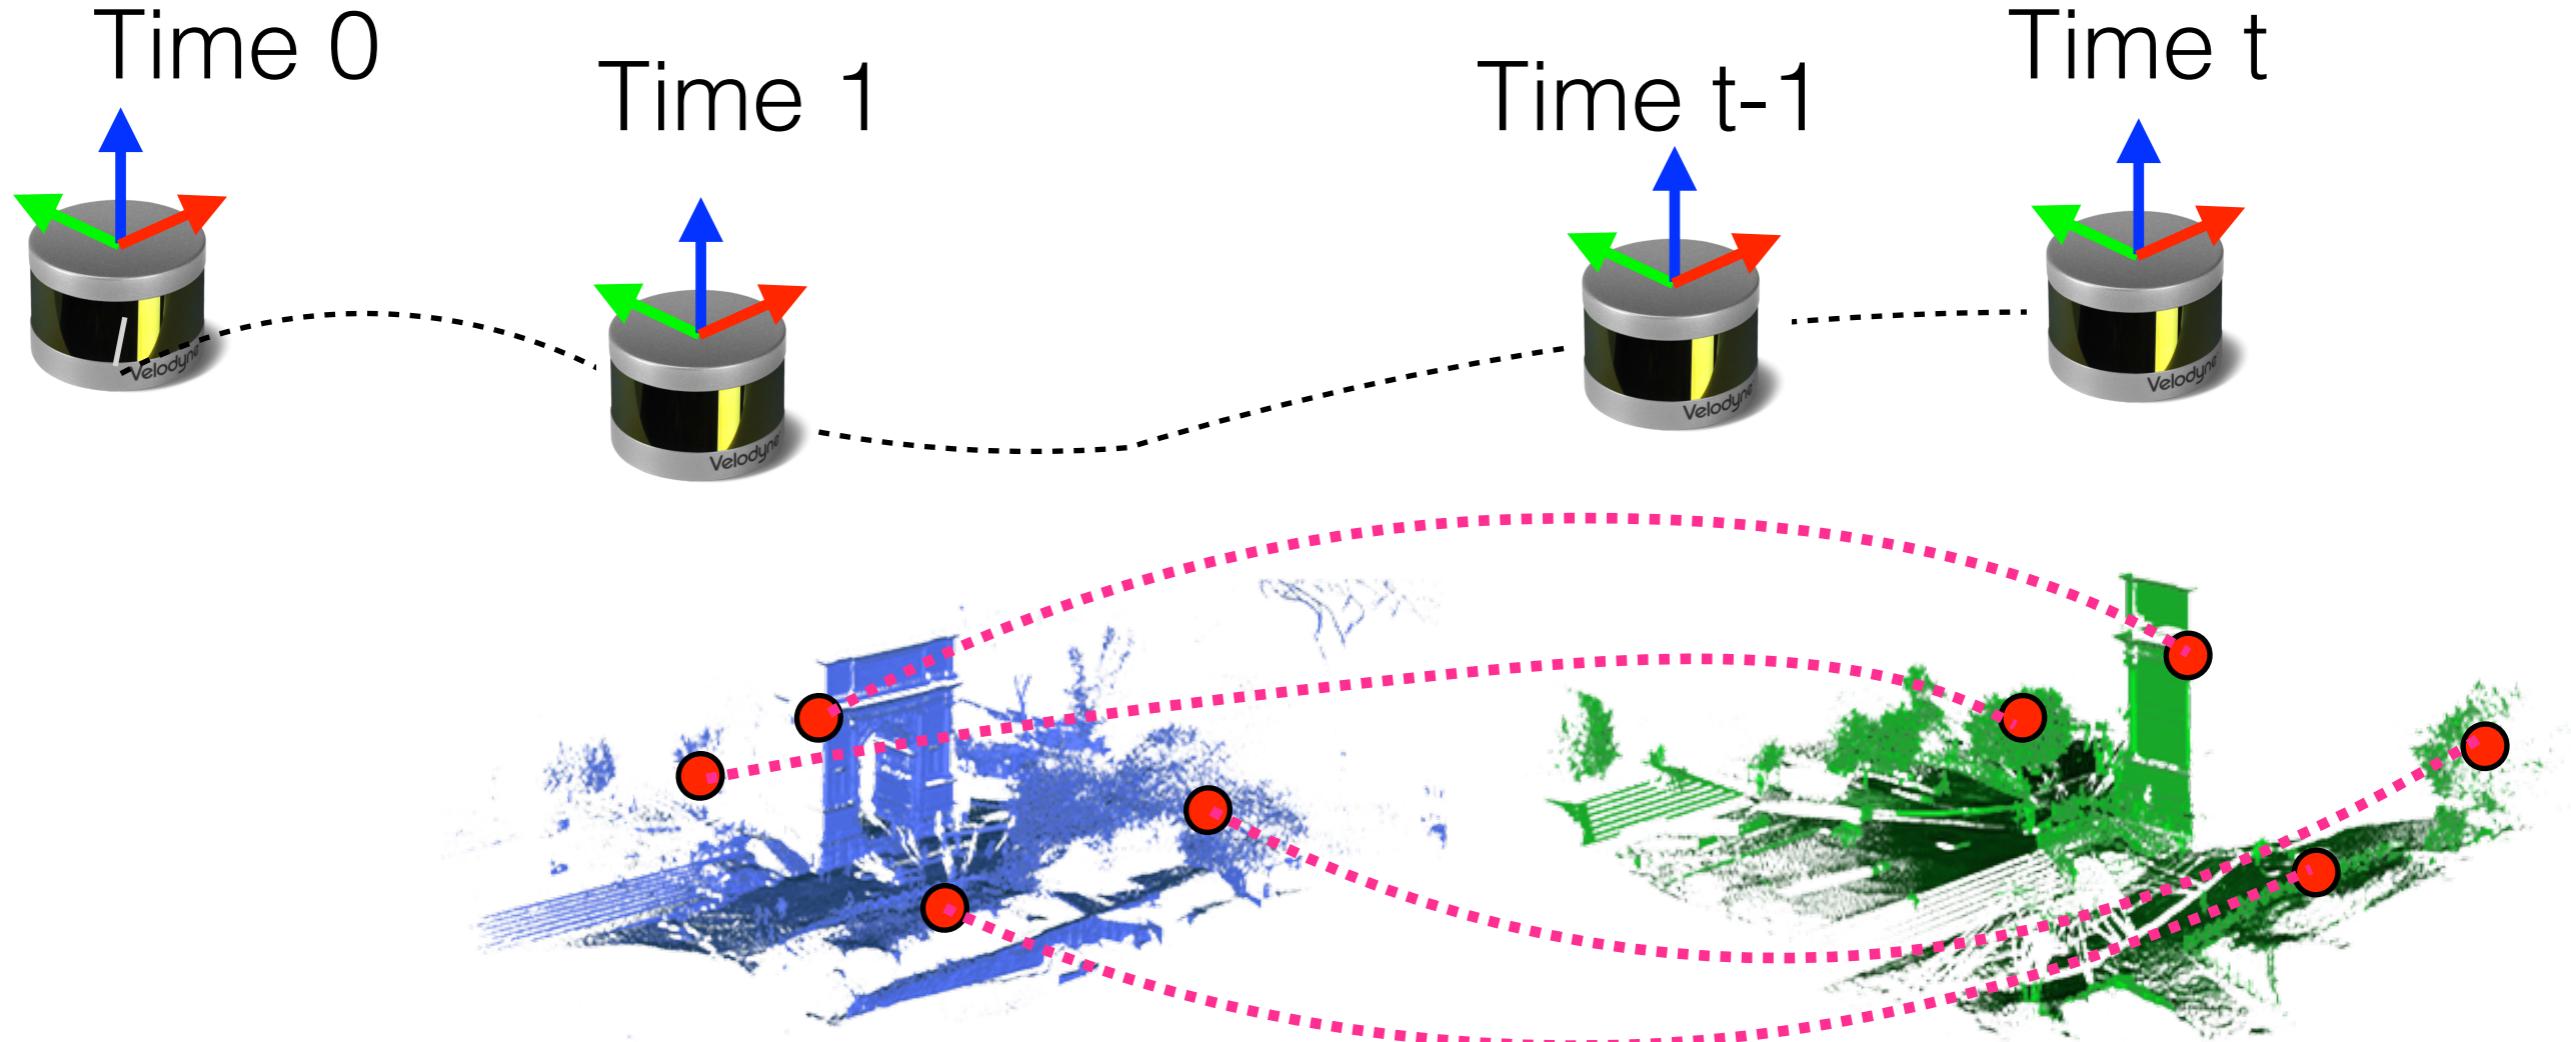
Feature-based Lidar Odometry



Registration: compute relative pose between scans:

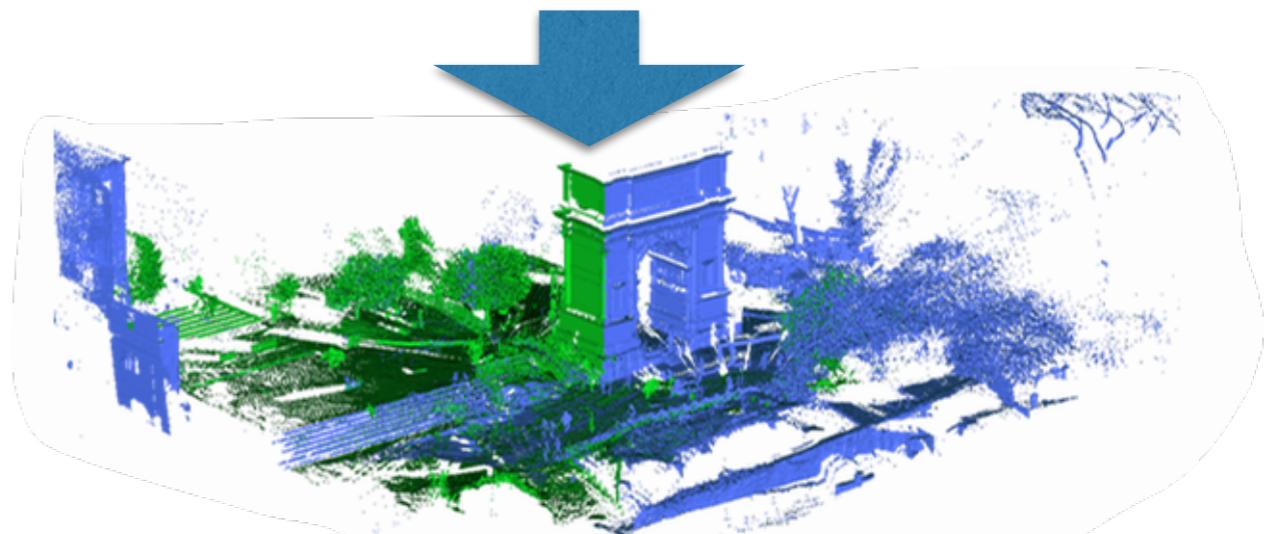
- extract features & descriptors
- use descriptors for matching
- compute relative pose

Feature-based Lidar Odometry

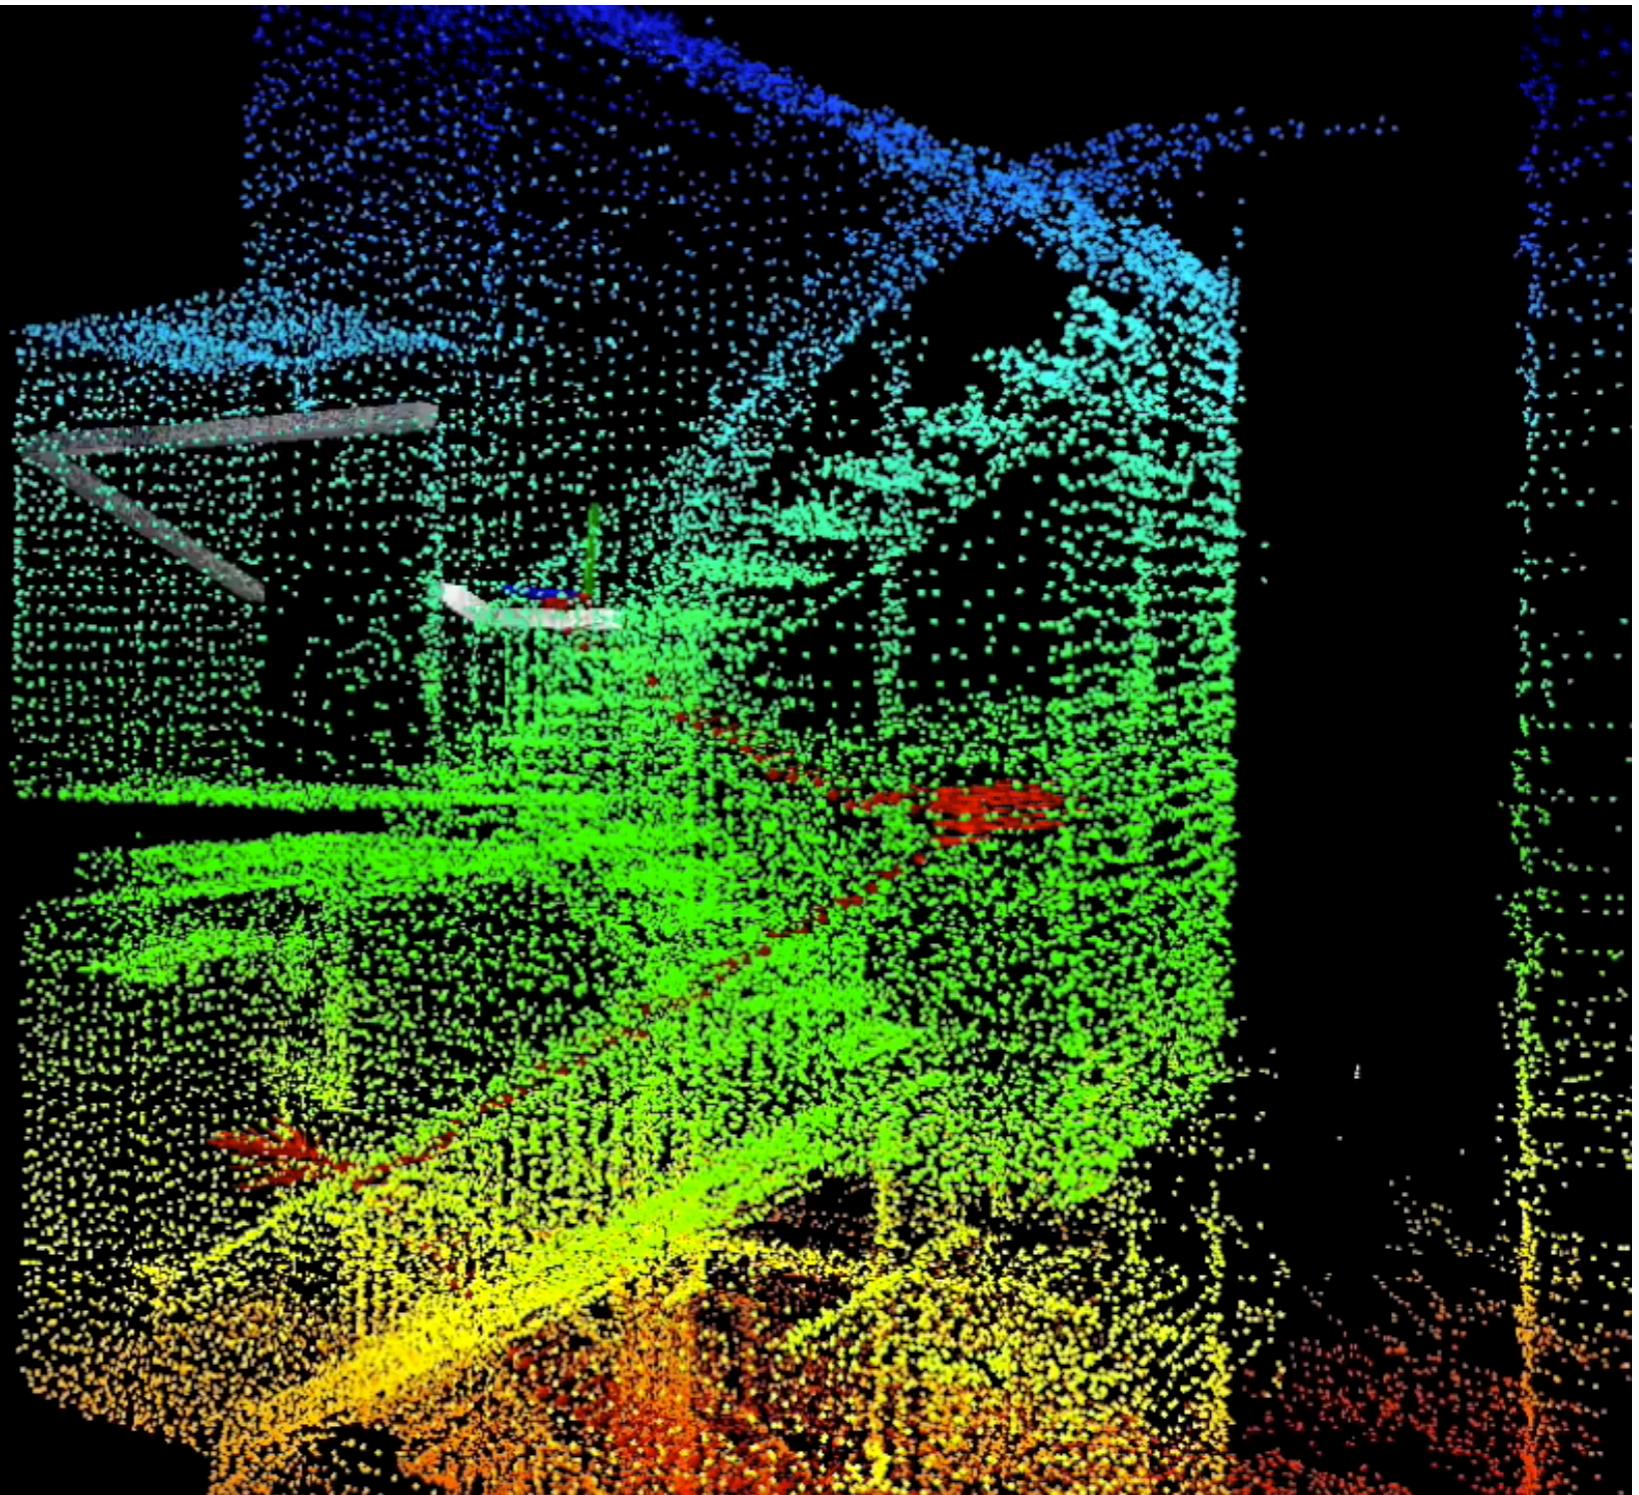


Registration: compute relative pose between scans:

- extract features & descriptors
- use descriptors for matching
- **compute relative pose**



Feature-based Lidar Odometry

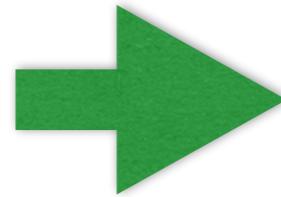
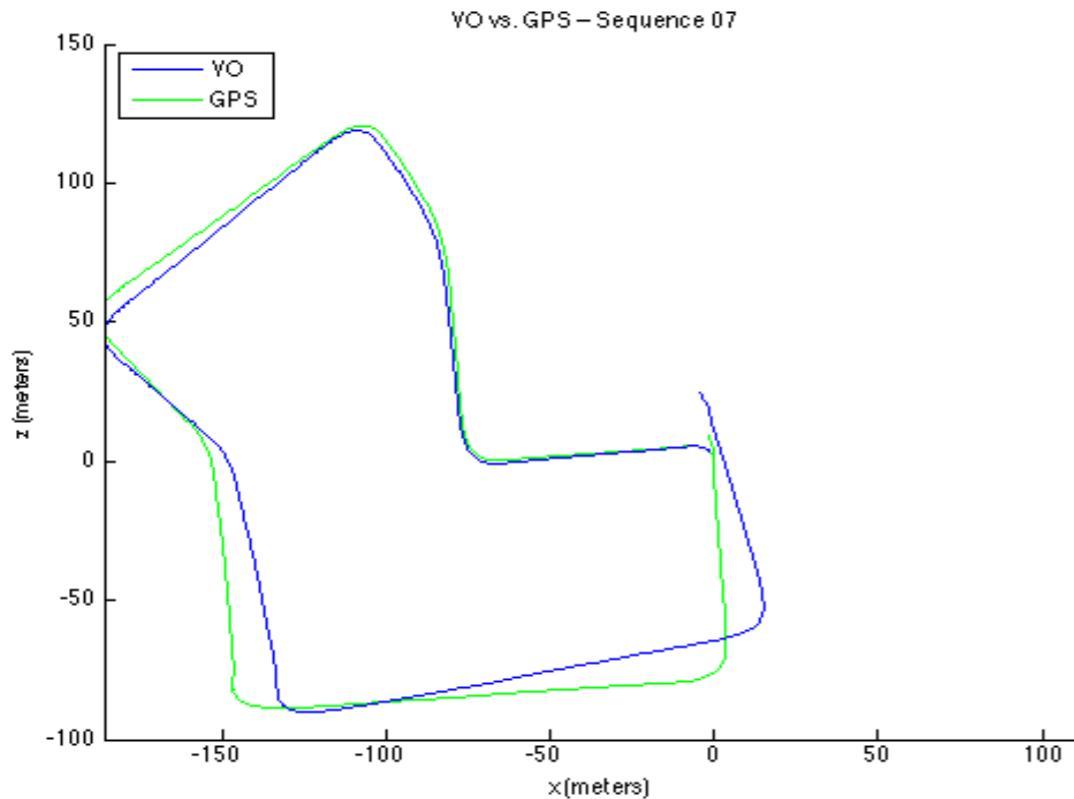


[Zhang and Singh: LOAM: Lidar Odometry and Mapping in Real-time, 2014]

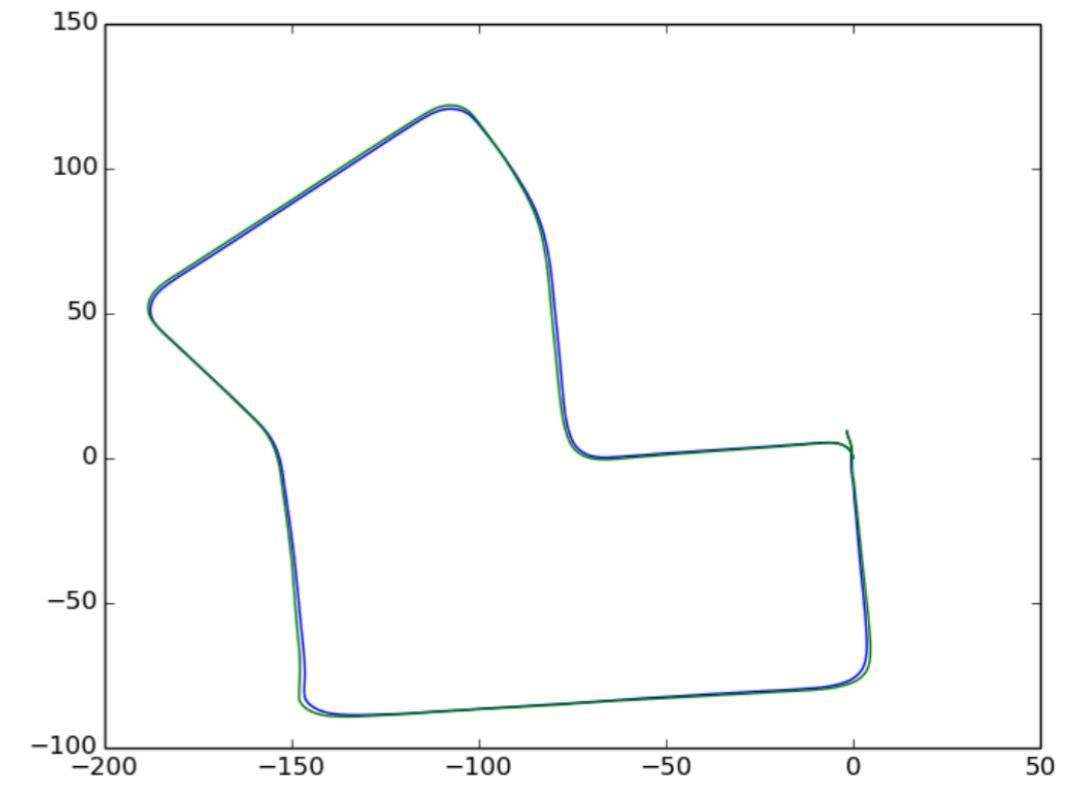
Other approaches: based on Iterative Closest Point (ICP)

Removing Drift via Loop Closure

Visual(-inertial) odometry



SLAM



SLAM requires:

- place recognition (loop closure detection)
- Re-detecting landmarks (e.g., objects)